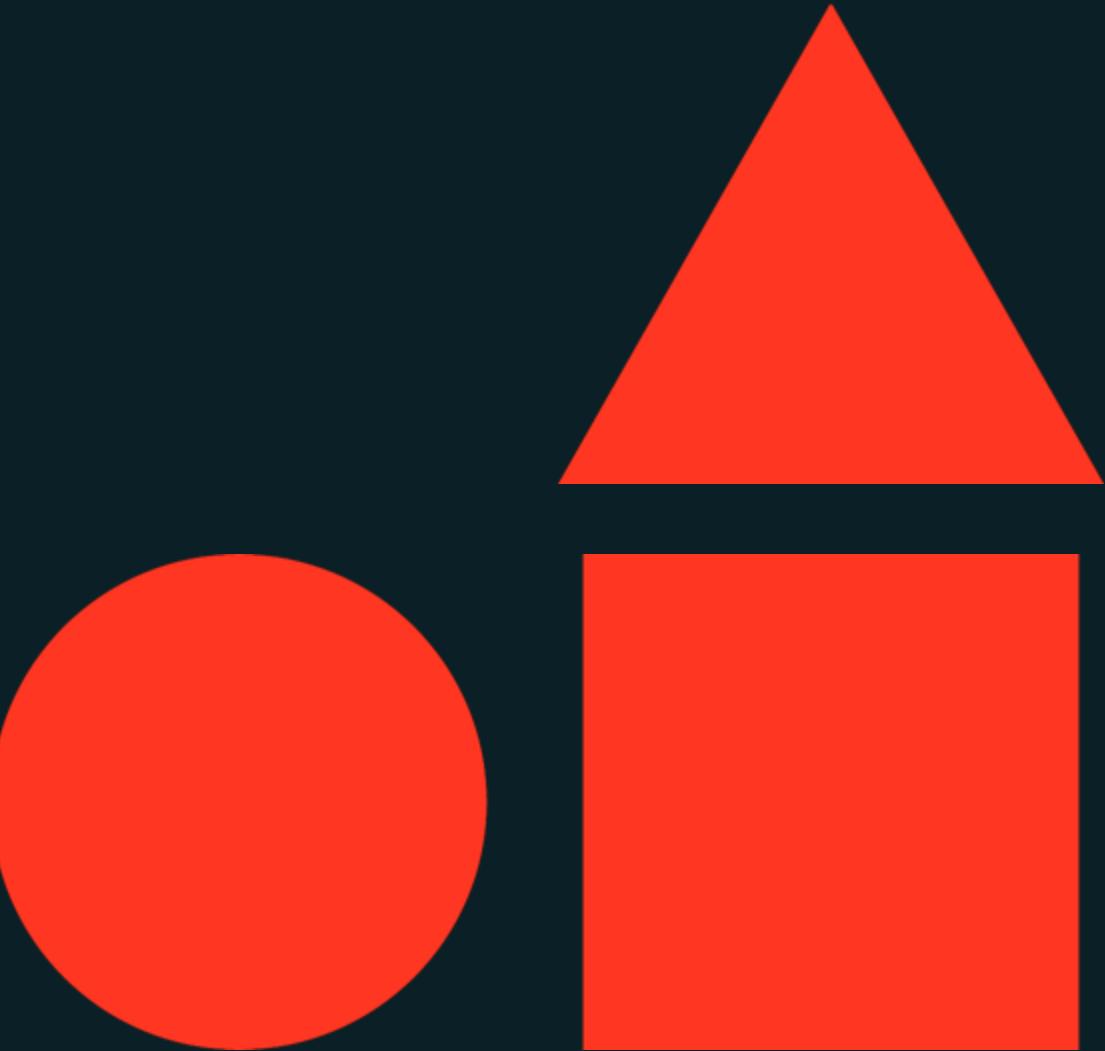




# Data Preparation for Machine Learning

---

Databricks Academy



# Learning goals

Upon completion of this content, you should be able to:

- Describe Databricks Data Intelligence Platform and its features for machine learning.
- Explain data storage and governance features on Databricks.
- Perform exploratory data analysis and feature engineering using Spark and integrated visualization tools.
- Perform data pre-processing for missing data handling, data encoding and data standardization.
- Utilize Feature Store for storing and retrieving features.



# Prerequisites/Technical Considerations

Things to keep in mind before you work through this course

## Prerequisites

-  1 Familiarity with Databricks workspace and notebooks
-  2 Familiarity with Delta Lake and Lakehouse
-  3 Intermediate level knowledge of Python

## Technical Considerations

-  1 Cluster running on **DBR ML 16.3**
-  2 Unity Catalog enabled workspace



# AGENDA

1. Managing and Exploring Data	DEMO	LAB
<b>Managing and Exploring Data in the Lakehouse</b>	✓	✓
02. Data Preparation and Feature Engineering		
<b>Fundamentals of Data Preparation and Feature Engineering</b>		
<b>Data Imputation</b>		
<b>Data Encoding</b>	✓	✓
<b>Data Standardization</b>		
03. Feature Store		
<b>Introduction to Feature Store</b>	✓	✓

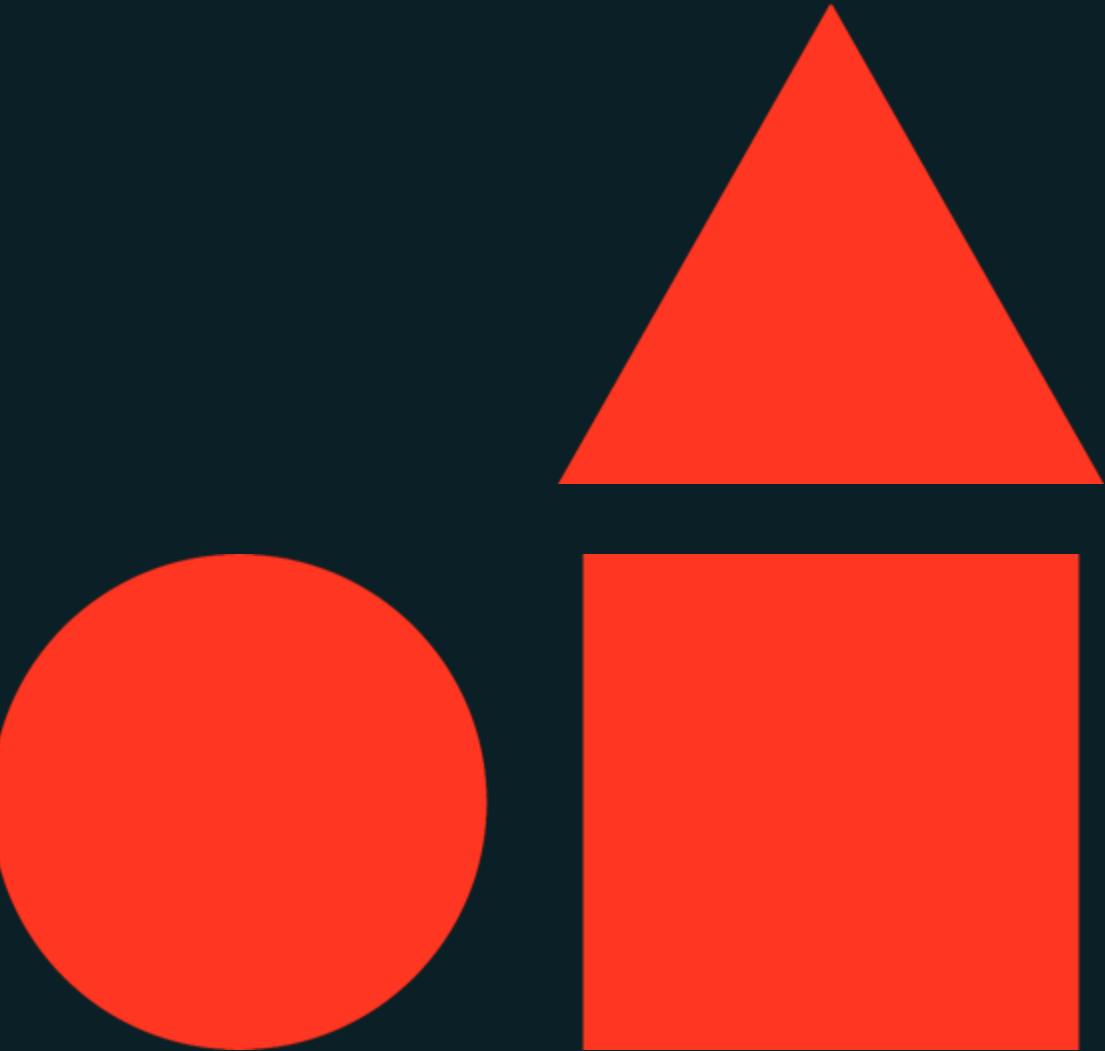




# Managing and Exploring Data

---

Data Preparation for Machine Learning



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Learning objectives

Things you'll be able to do after completing this module

- Understand Databricks Machine Learning as a data-centric platform.
- Explain the benefits of Delta Lake for machine learning, including version control and data integrity.
- Describe the benefits of Unity Catalog for machine learning, including data discovery and collaboration.
- Load and write data from/to Delta tables.
- Use data profiling and visualization to explore and analyze machine learning data.





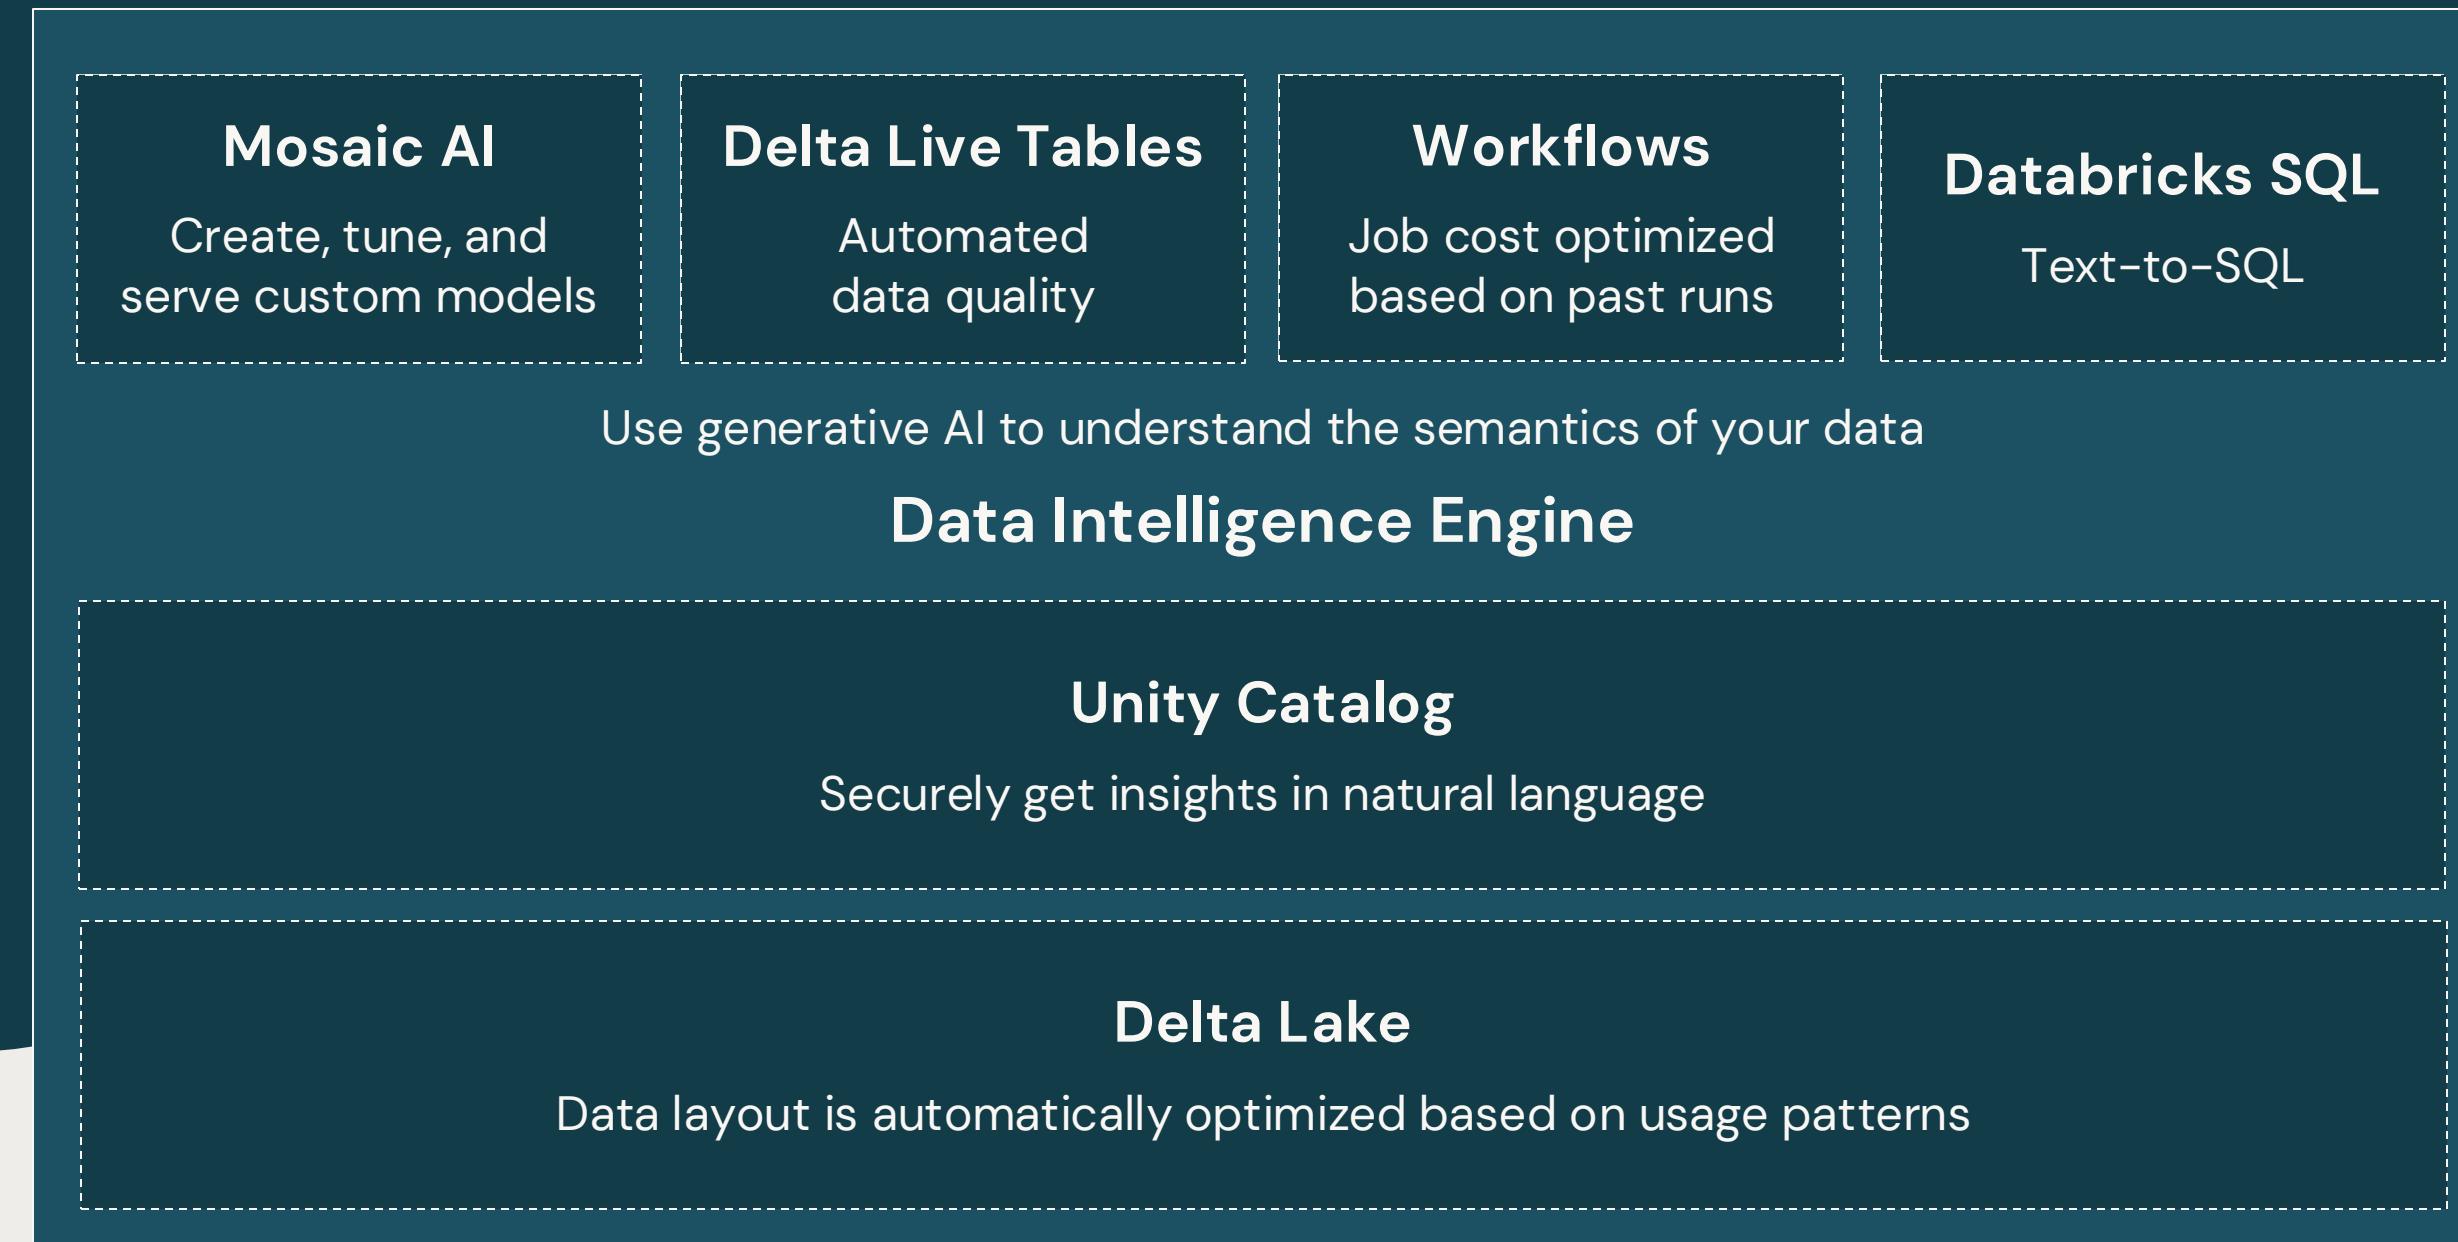
Managing and Exploring Data

LECTURE

# Managing and Exploring Data in the LakeHouse



# Databricks Data Intelligence Platform



## Open Data Lake

All Raw Data  
(Logs, Texts, Audio, Video, Images)

# Data Science & AI on Databricks

## Mosaic AI

### End-to-end AI

- MLOps (MLflow)
- AutoML
- Model Serving
- Monitoring
- Governance

### Gen AI

- Custom models
- Model serving
- RAG

Data Science  
& AI

Mosaic AI

ETL &  
Real-time Analytics

Delta Live Tables

Orchestration

Workflows

Data  
Warehousing  
Databricks SQL

Use generative AI to understand the semantics of your data

## Data Intelligence Engine

## Unity Catalog

Securely get insights in natural language

## Delta Lake

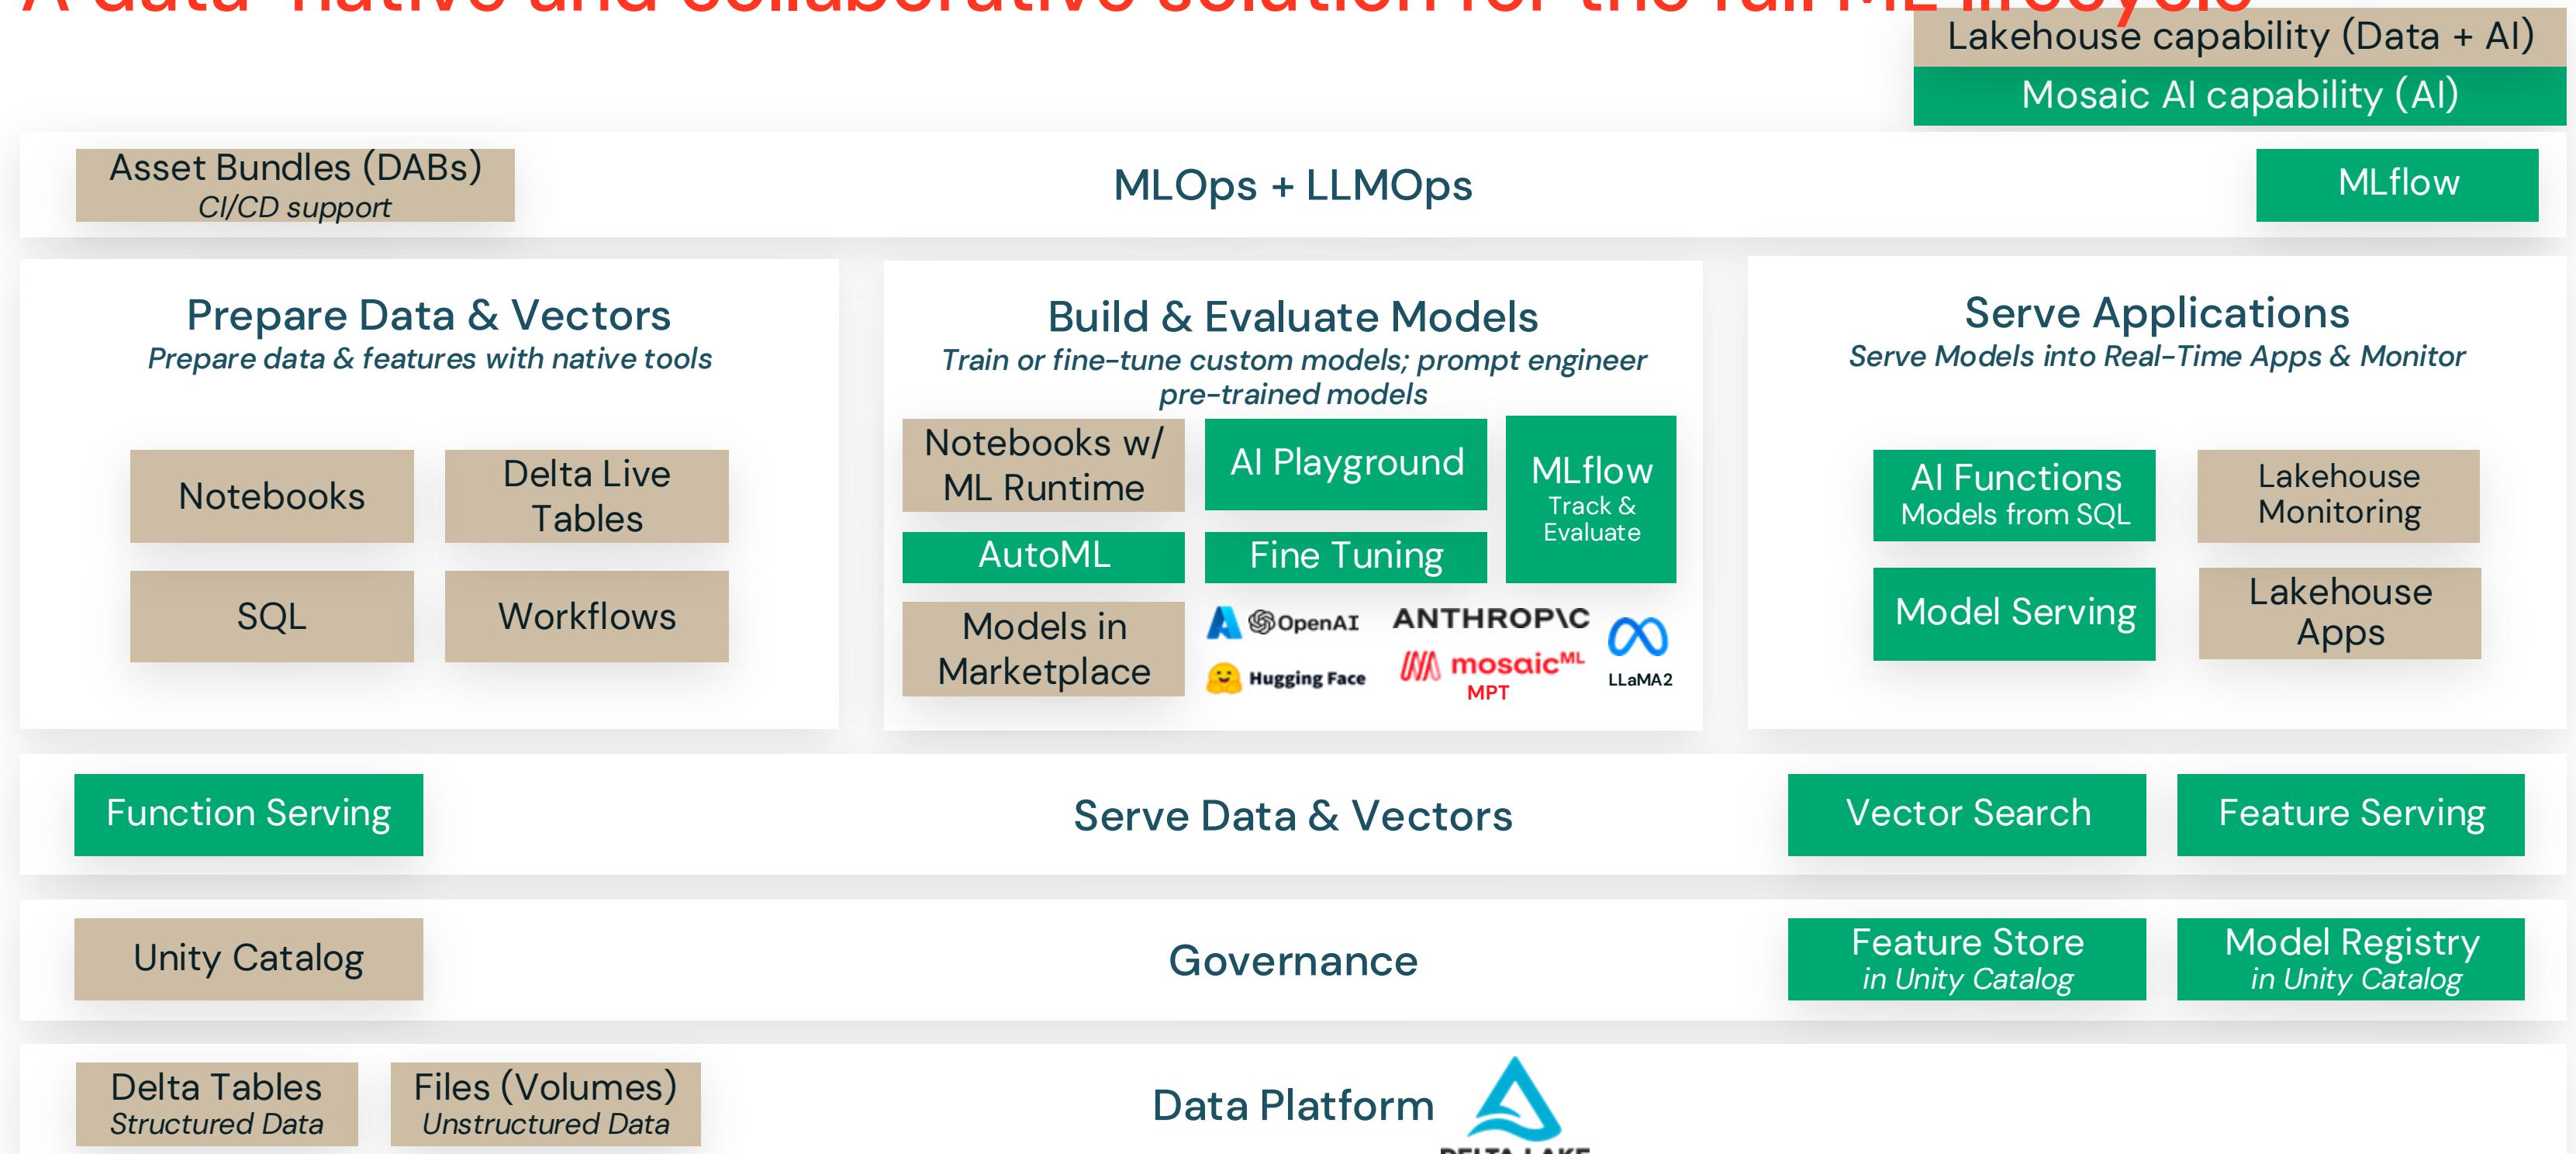
Data layout is automatically optimized based on usage patterns

## Open Data Lake

All Raw Data  
(Logs, Texts, Audio, Video, Images)

# Databricks for Machine Learning

A data-native and collaborative solution for the full ML lifecycle



# Features of Databricks for Machine Learning

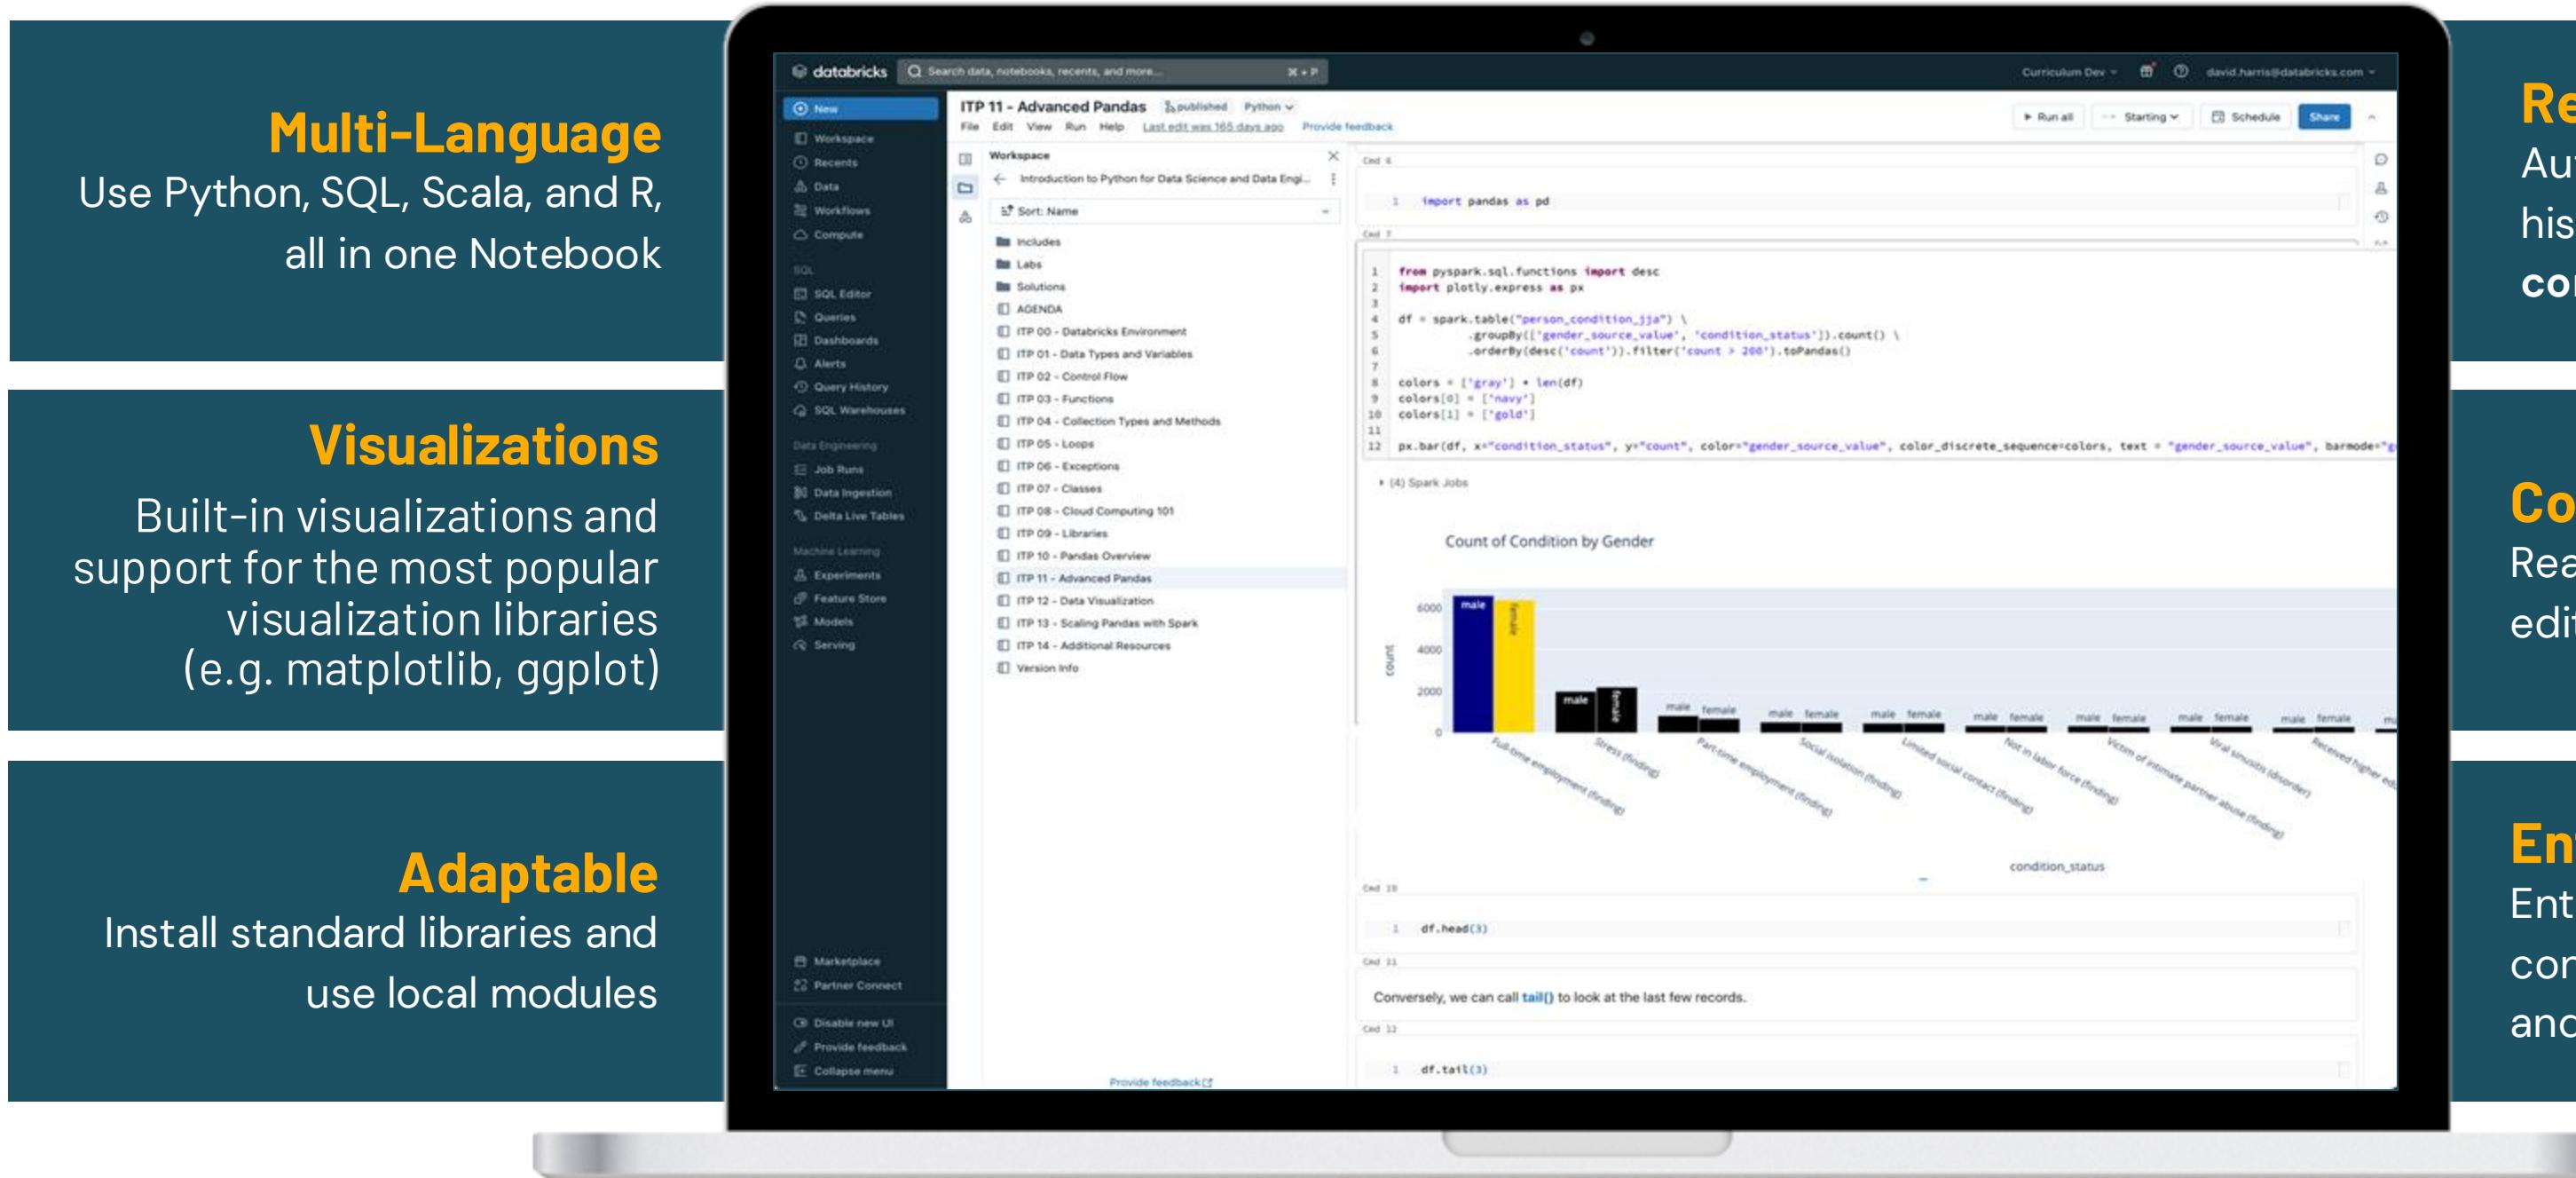
Non exhaustive list of features that will be used throughout this module

- Collaborative notebooks
- ML Runtime
- Governance of Data & Models (*via Unity Catalog*)
- Feature Store
- Managed MLflow
- Model Serving
- AutoML



# Collaborative Multi-Language Notebooks

# Collaborative, reproducible, and enterprise ready



# Multi-Language

# Use Python, SQL, Scala, and R, all in one Notebook

# Visualizations

Built-in visualizations and support for the most popular visualization libraries (e.g. matplotlib, ggplot)

# Adaptable

# Install standard libraries and use local modules

# Reproducible

Automatically track version history, and use **git version control** with Repos

# Collaborative

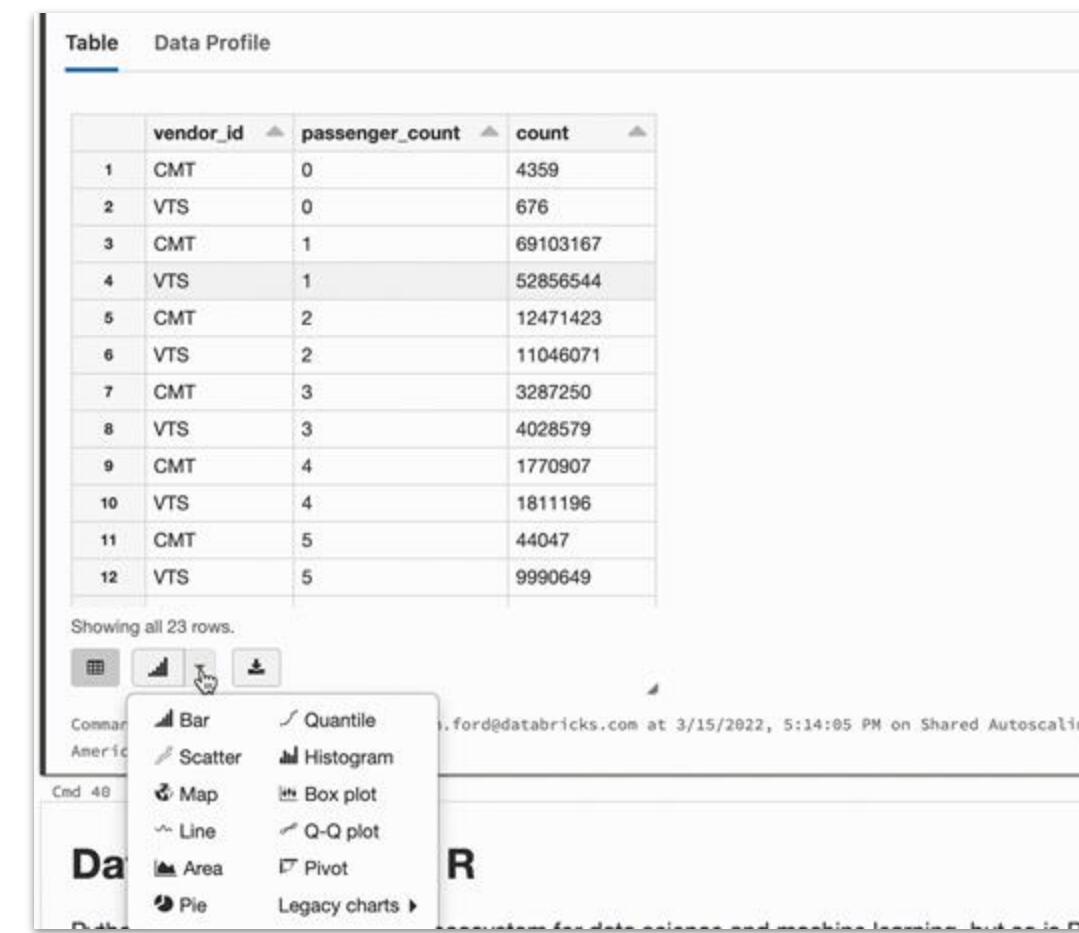
## Real-time co-presence, co-editing, and commenting

# Enterprise Ready

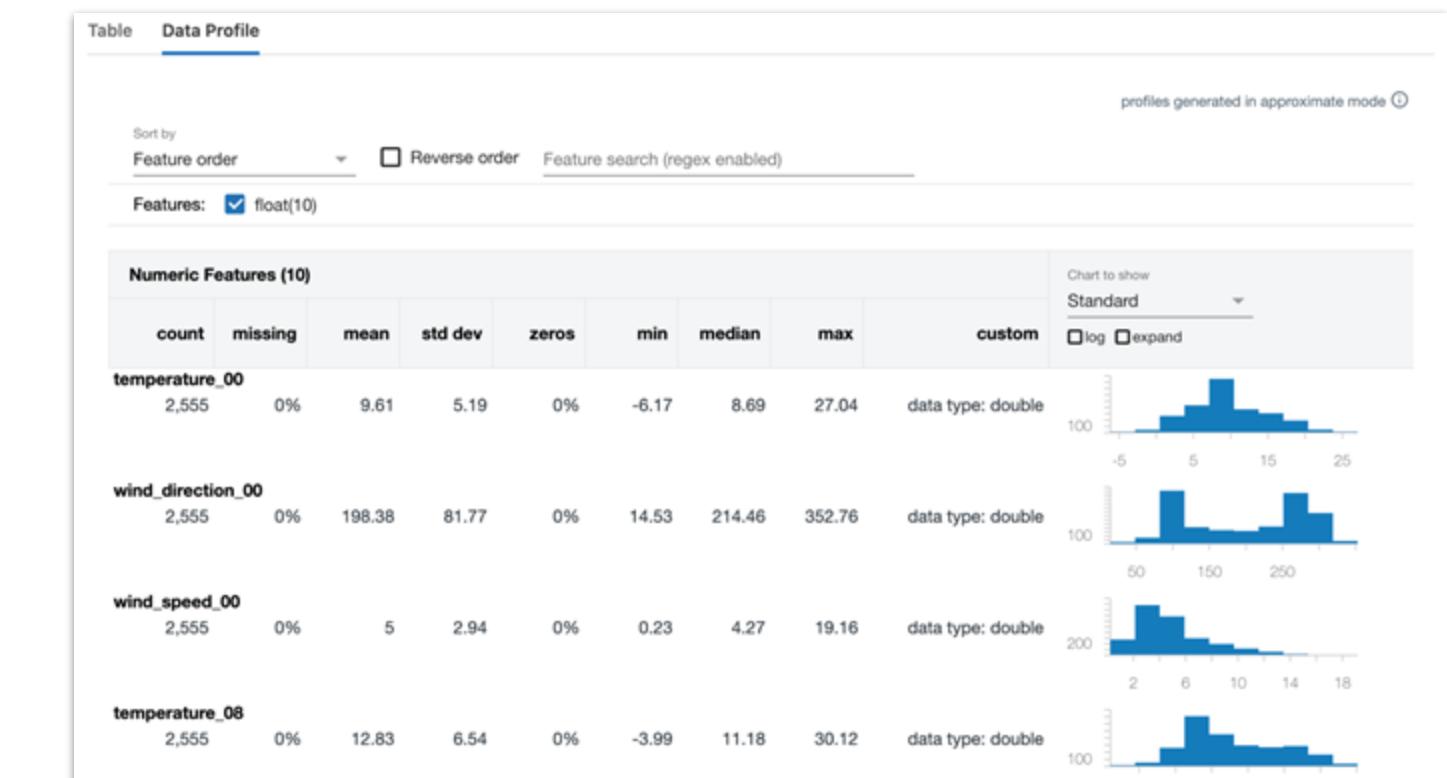
Enterprise-grade access controls, identity management, and auditability

# Quick Exploratory Data Analysis

Native tools for visualizing and understanding data in ML workflow



Create **interactive charts** to visualize data in the Notebook with only two clicks

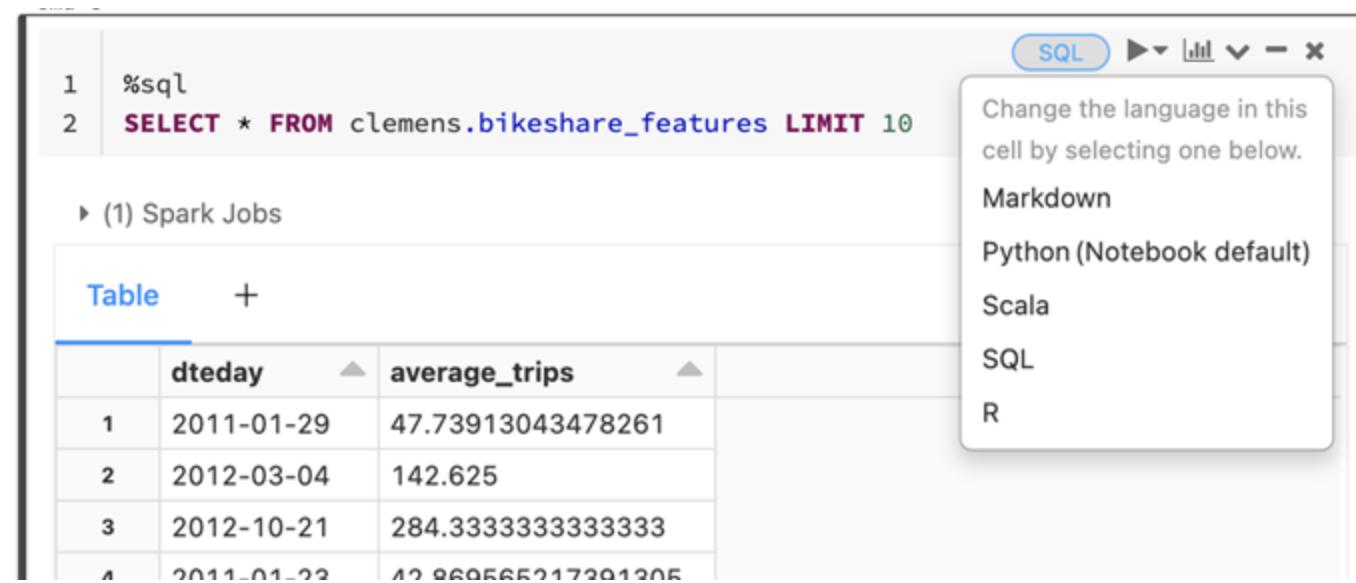


Summarize a data set's essential properties and statistics in a **data profile** with the push of a button



# Tools for Quick ML Model Development

Multi-language support, use standard libraries and custom modules



A screenshot of a Jupyter Notebook cell. The cell contains the following SQL code:

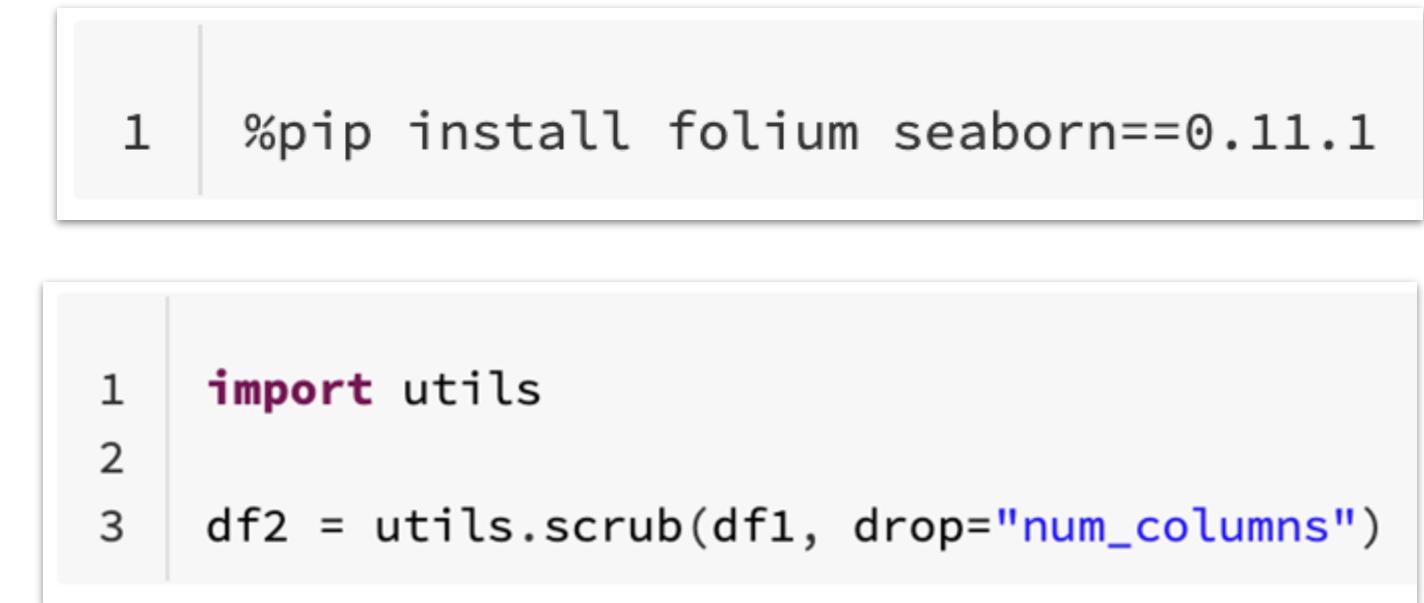
```
1 %sql
2 SELECT * FROM clemens.bikeshare_features LIMIT 10
```

The cell has a "SQL" button at the top. Below the code, there is a table output:

	dteday	average_trips
1	2011-01-29	47.73913043478261
2	2012-03-04	142.625
3	2012-10-21	284.3333333333333
4	2011-01-29	42.96955917291205

To the right of the table, a language selection dropdown menu is open, showing options: Markdown, Python (Notebook default), Scala, SQL, and R.

**Mix and match languages** based on use case and preferred workflow, choosing from **Python, SQL, Scala, and R**



A screenshot of a Jupyter Notebook cell. The cell contains the following commands:

```
1 %pip install folium seaborn==0.11.1
```

```
1 import utils
2
3 df2 = utils.scrub(df1, drop="num_columns")
```

**Install Python libraries** for a notebook without affecting other users with **%pip**  
**Import local modules** using **arbitrary file support** when working in Repos



# Data Science & AI on Databricks

## Unity Catalog

- Context-aware search
- Auto describe tables and columns
- Automated lineage
- End-to-end observability and monitoring
- Sharing AI models

Data Science  
& AI

Mosaic AI

ETL &  
Real-time Analytics

Delta Live Tables

Orchestration

Workflows

Data  
Warehousing  
Databricks SQL

Use generative AI to understand the semantics of your data

## Data Intelligence Engine

### Unity Catalog

Securely get insights in natural language

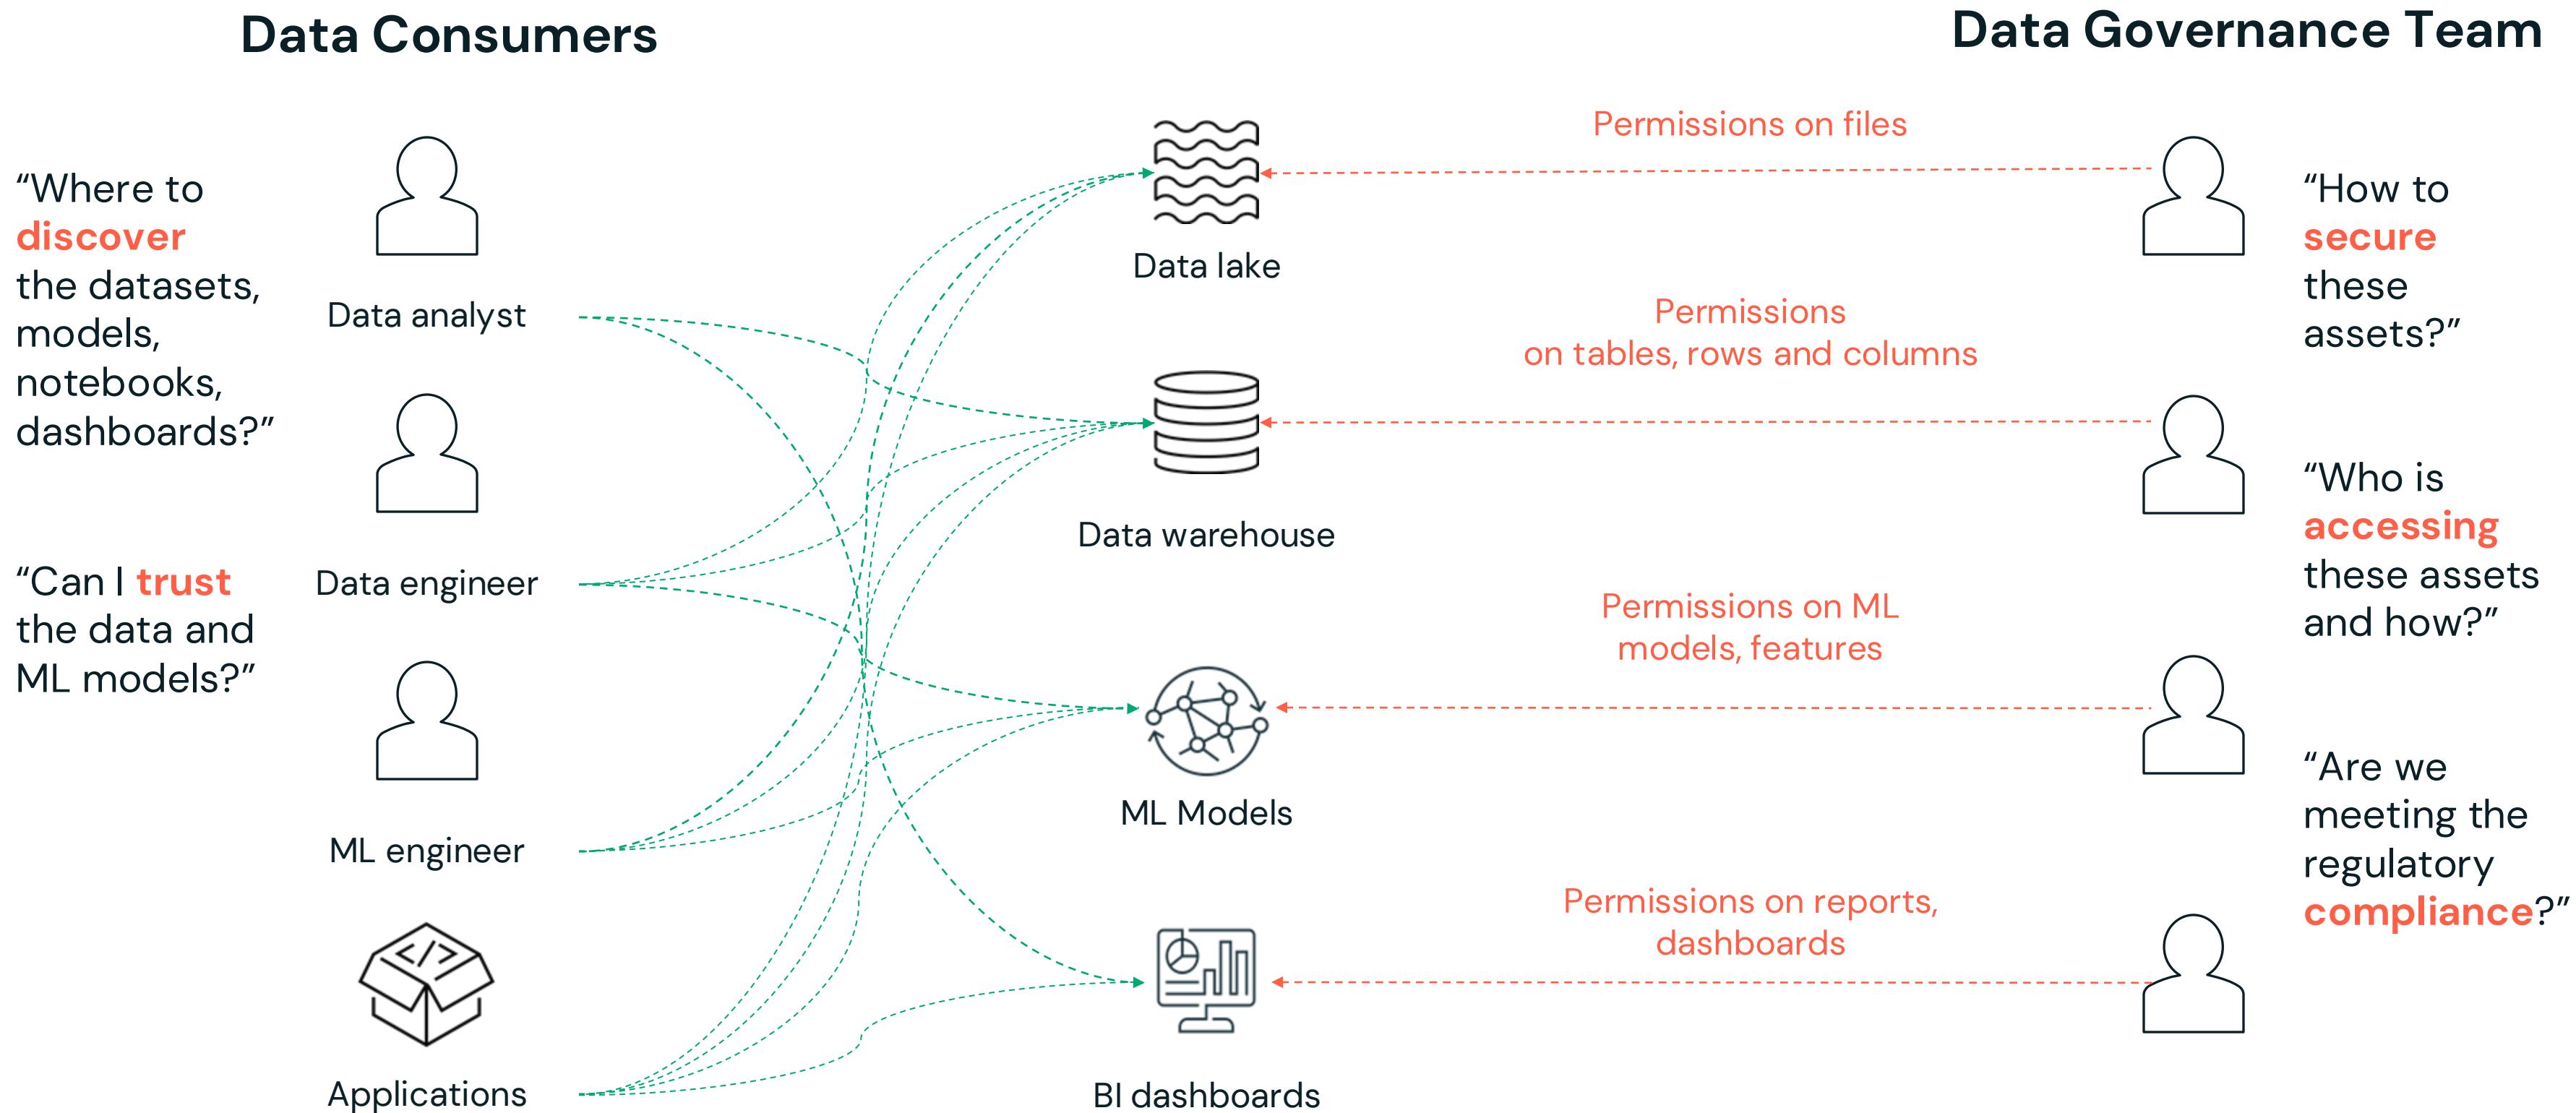
### Delta Lake

Data layout is automatically optimized based on usage patterns

## Open Data Lake

All Raw Data  
(Logs, Texts, Audio, Video, Images)

# Today, data and AI governance is complex



# Unity Catalog (UC)

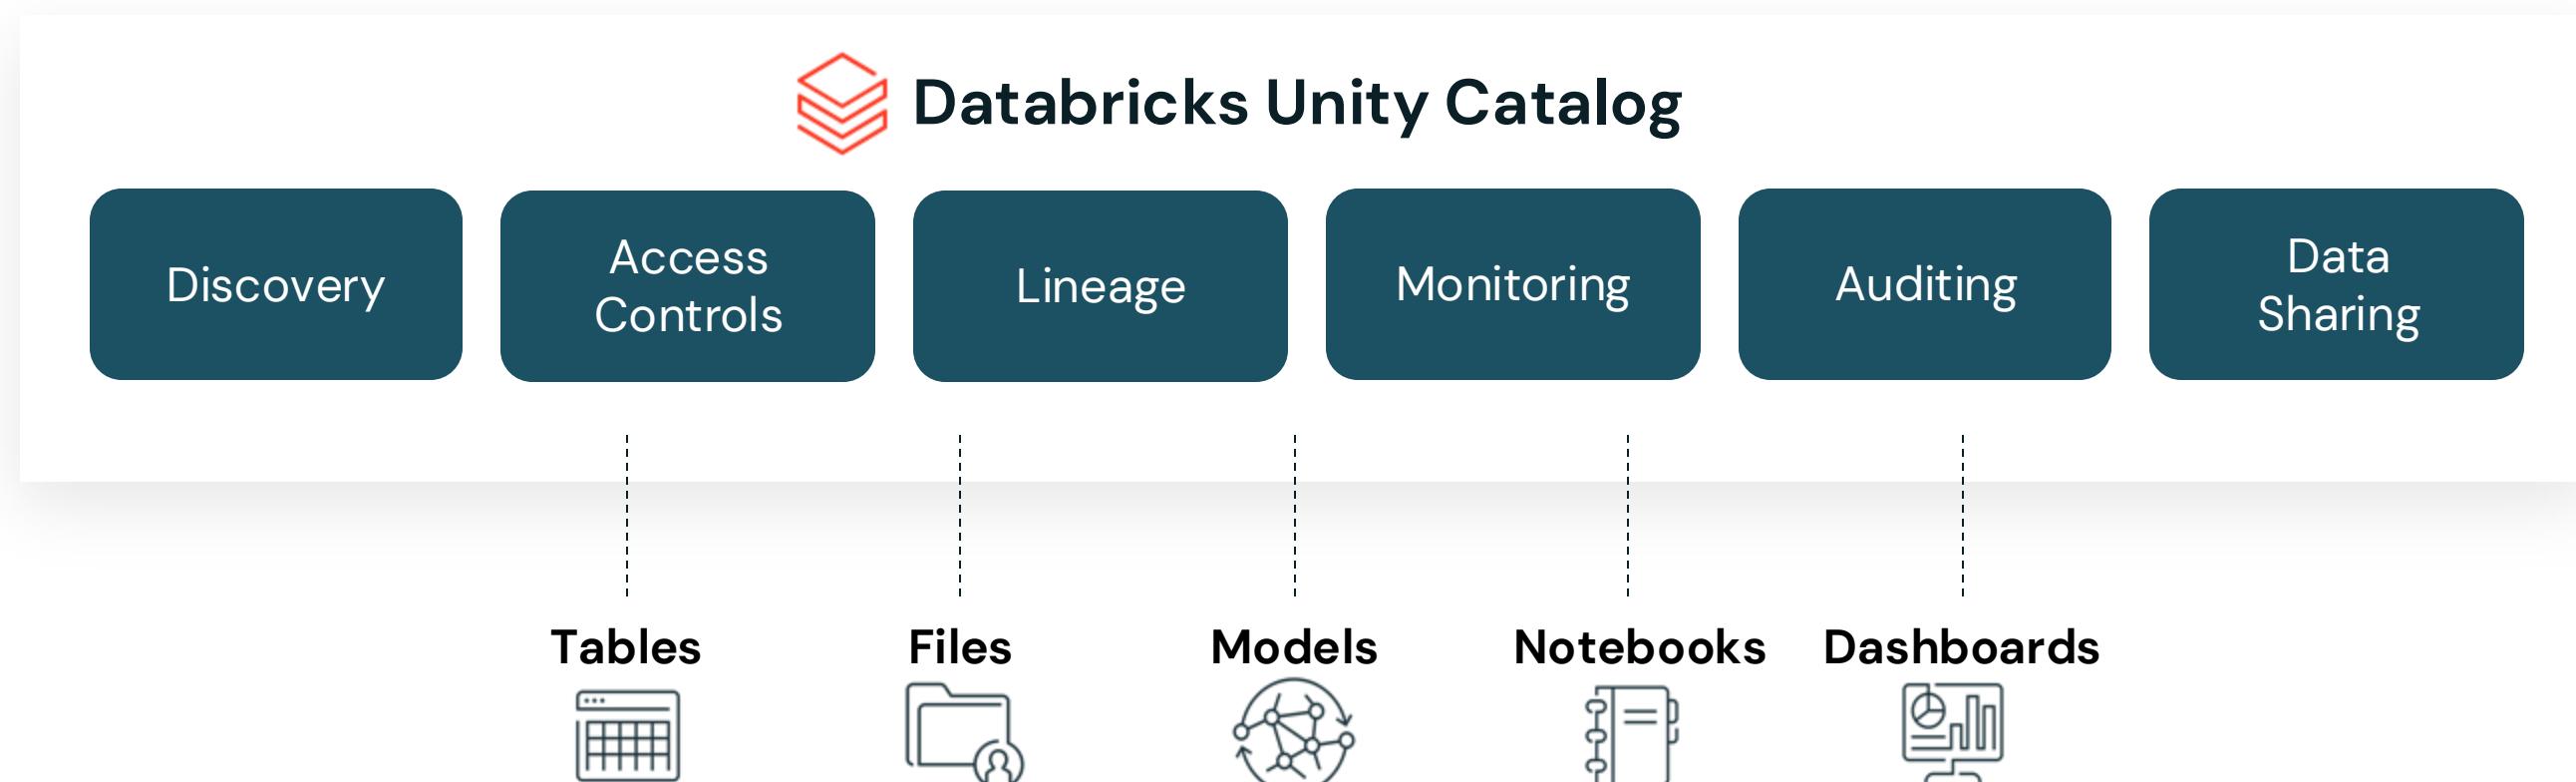
Unified governance for AI; data, code, models

Unified visibility into data and AI

Single permission model for data and AI

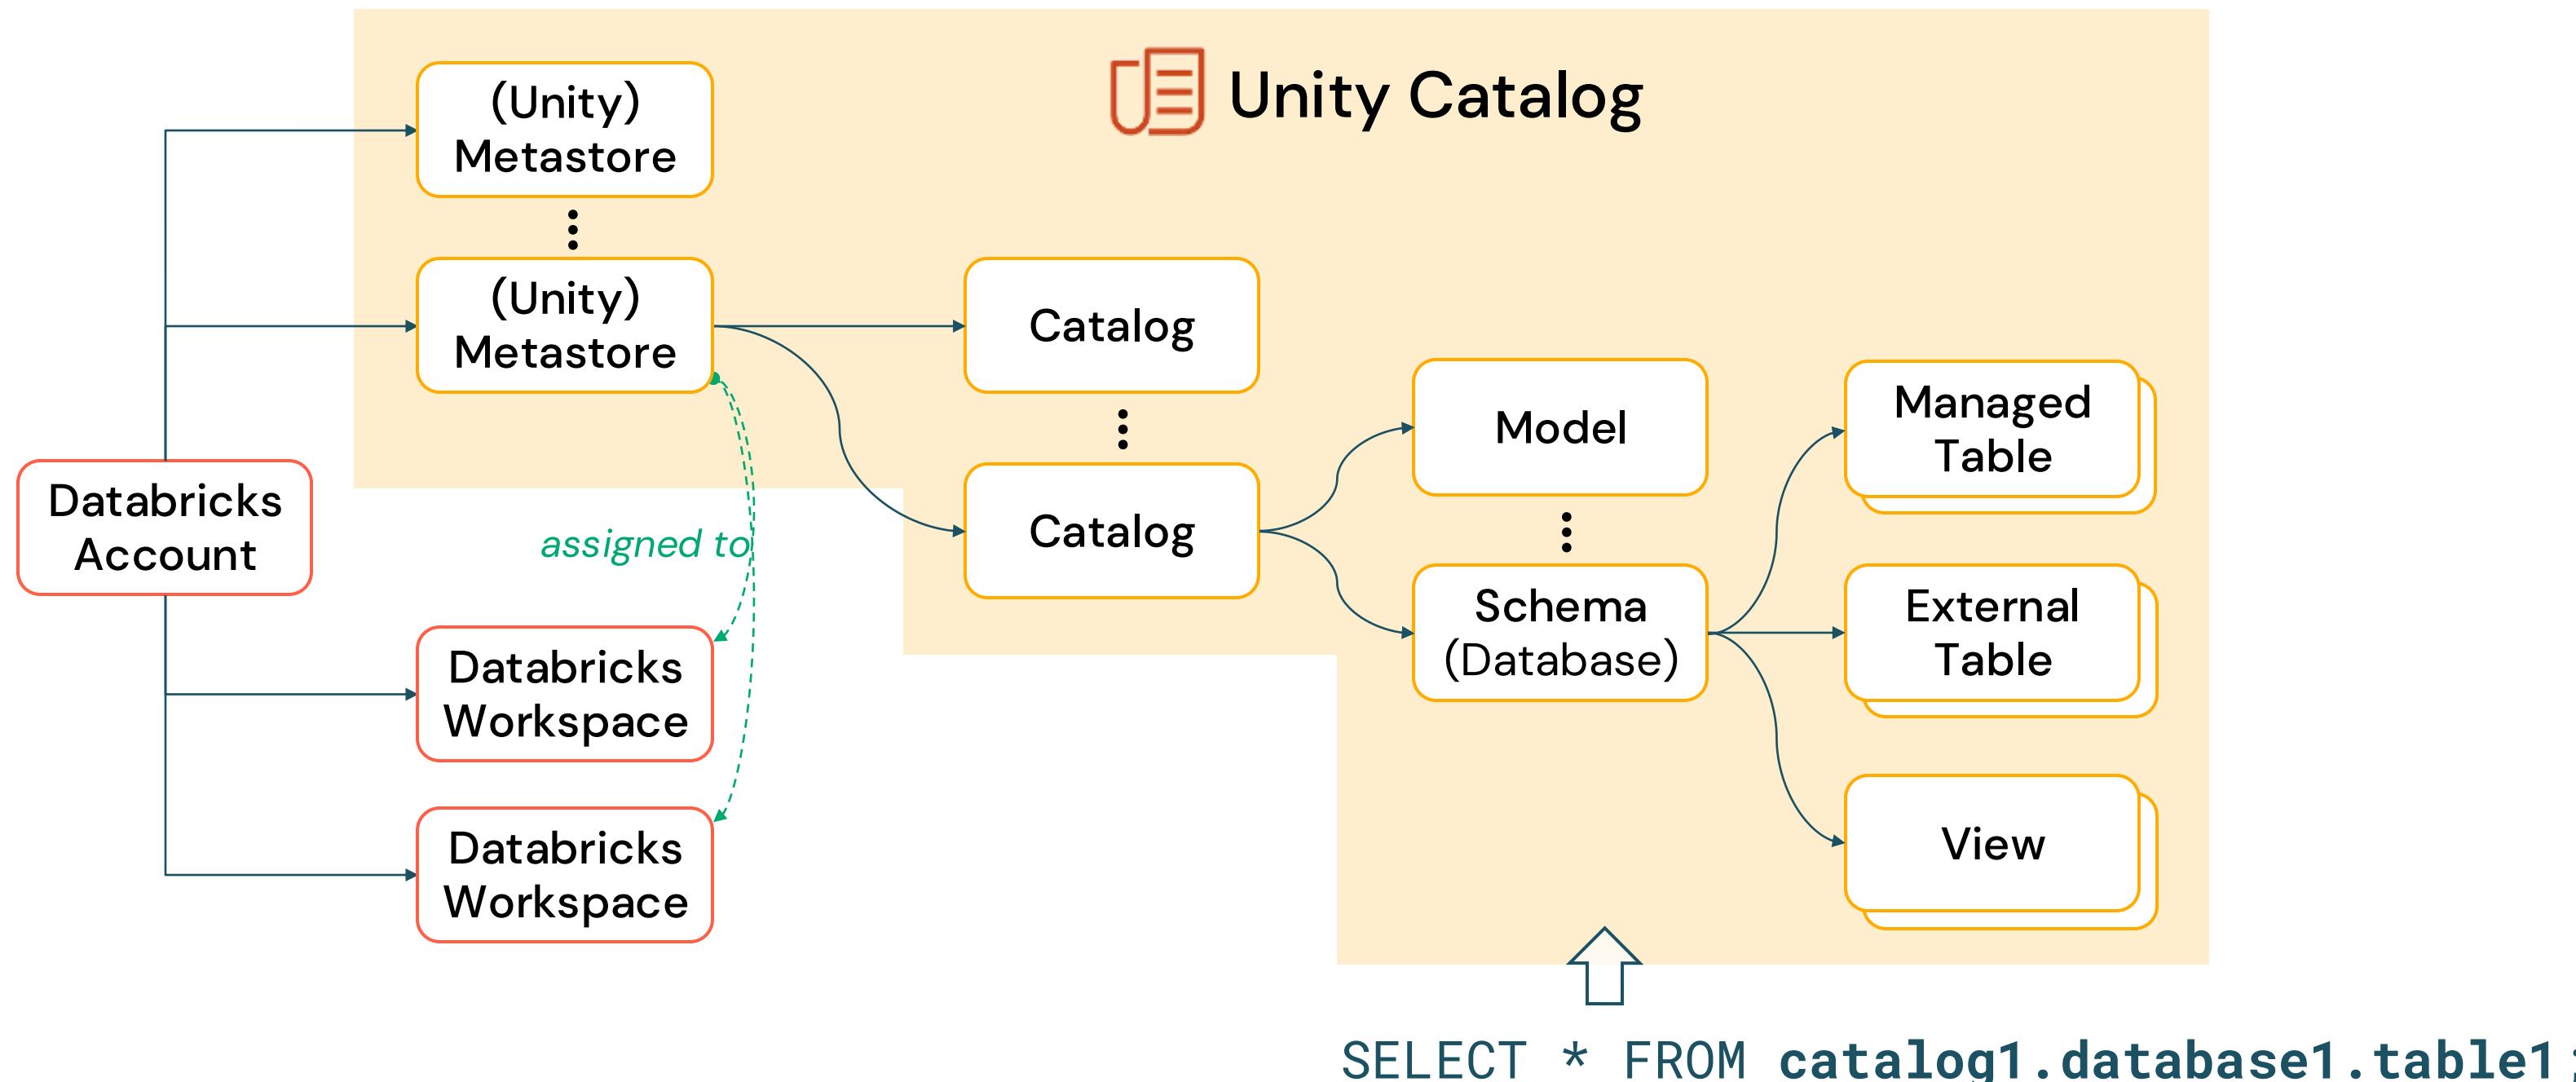
AI-powered monitoring and observability

Open data sharing



# The 3-level Namespace of UC

How to use UC



# Data Science & AI on Databricks

## Delta Lake

- Open-source
- Unified data management layer
- Reliable and fast
- Optimization features for storing data in the cloud

Data Science & AI

Mosaic AI

ETL & Real-time Analytics

Delta Live Tables

Orchestration

Workflows

Data Warehousing  
Databricks SQL

Use generative AI to understand the semantics of your data

## Data Intelligence Engine

## Unity Catalog

Securely get insights in natural language

## Delta Lake

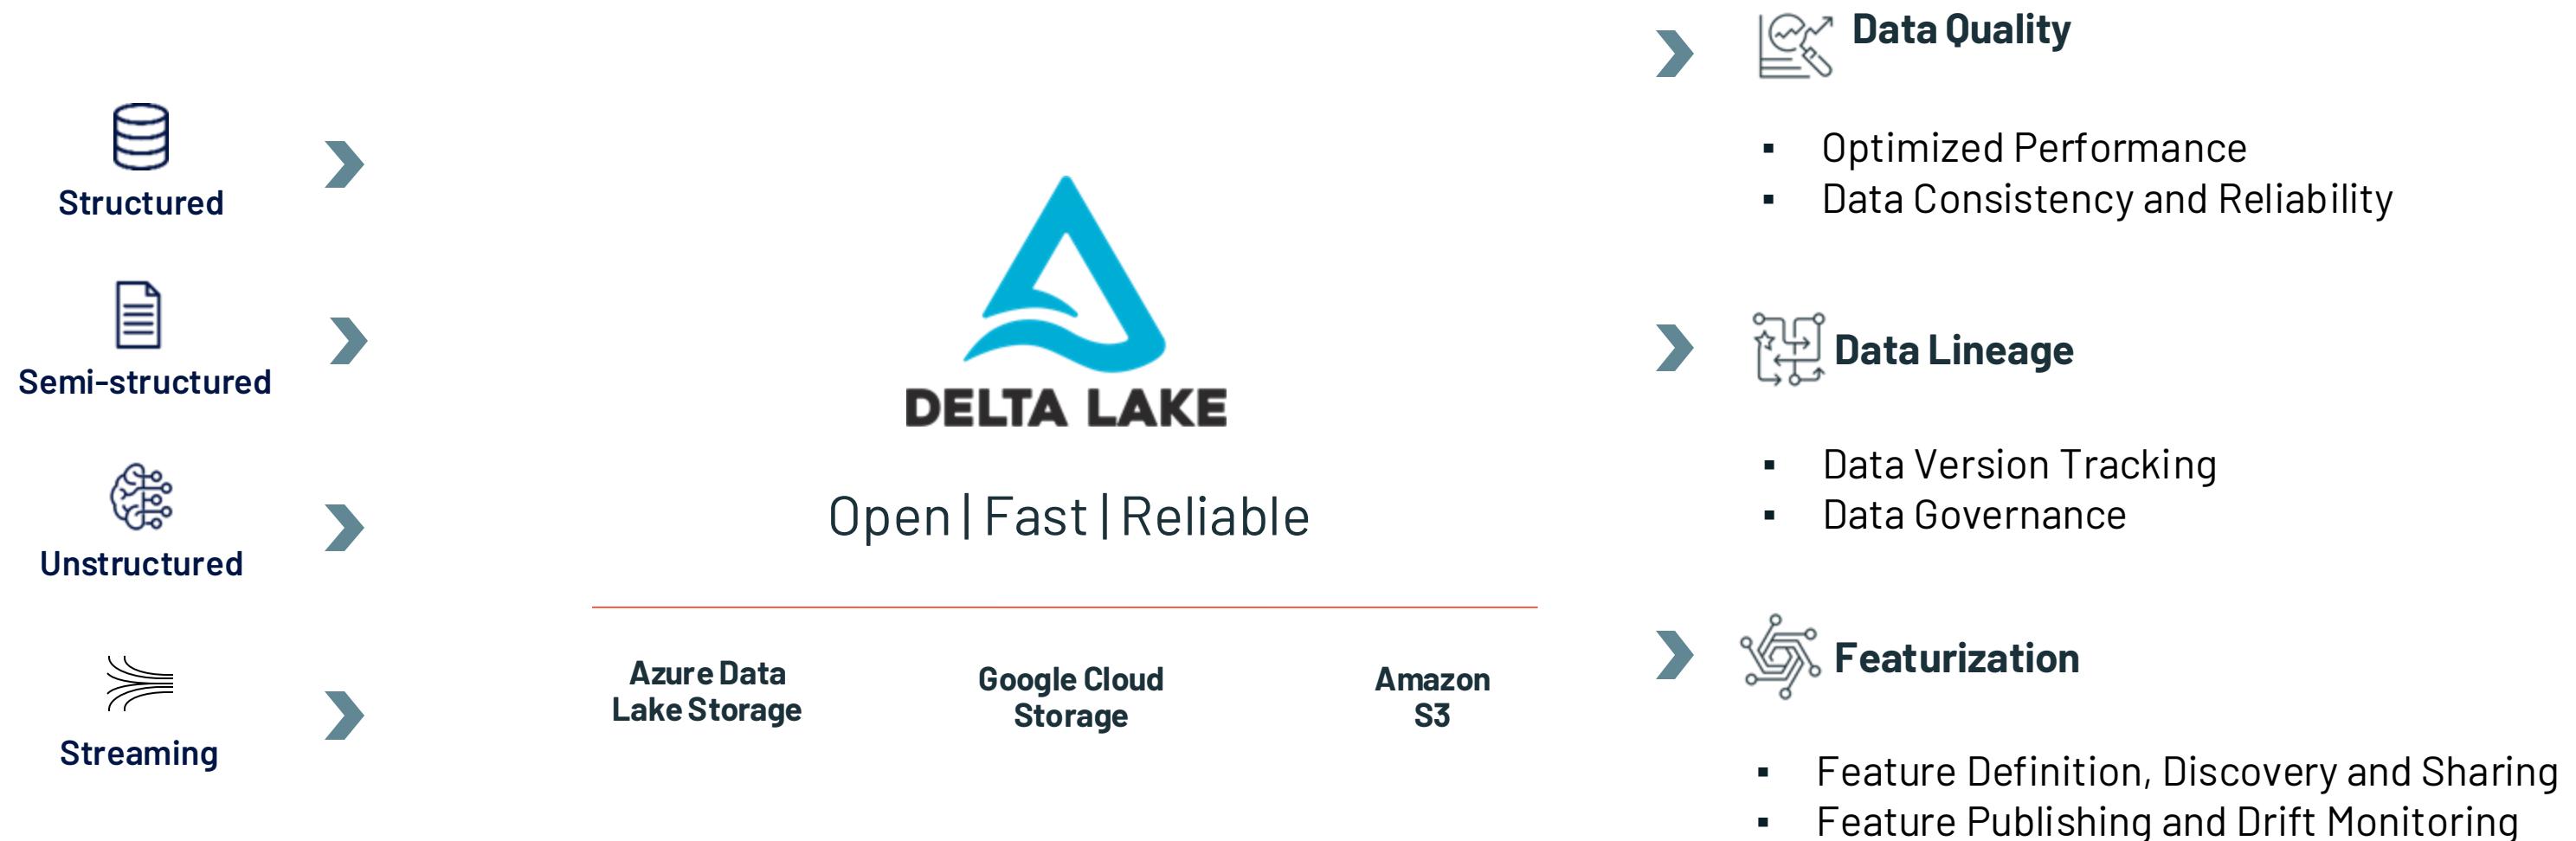
Data layout is automatically optimized based on usage patterns

## Open Data Lake

All Raw Data  
(Logs, Texts, Audio, Video, Images)

# What is Delta Lake?

- A **unified data management layer** that brings data reliability and fast analytics to cloud data lakes. It is the optimized storage layer that provides the foundation for storing data and tables in the Databricks DI Platform.



# Delta Lake and Its Features

Open-source, default storage format on Databricks

- Delta Lake is an **open-source** project.
- It is the **default format** for the tables created in Databricks.
- Delta Lake **optimizes performance** with large datasets, providing **ACID transactions** and **scalable metadata handling**.
- Designed to improve data reliability, quality, and performance in data lakes.



# Delta Lake features

## Key features

- Unified **batch and streaming**
- Automatic **schema validation**
- Support upserts using the merge operation
- Update your table schema without rewriting data.
- Track row-level changes with Change Data Feed
- **Time-travel**; querying previous versions of a table based on version number of timestamp
- **Performance optimization** with data skipping and liquid clustering
- Supports multiple programming languages like Python, Scala, and SQL.

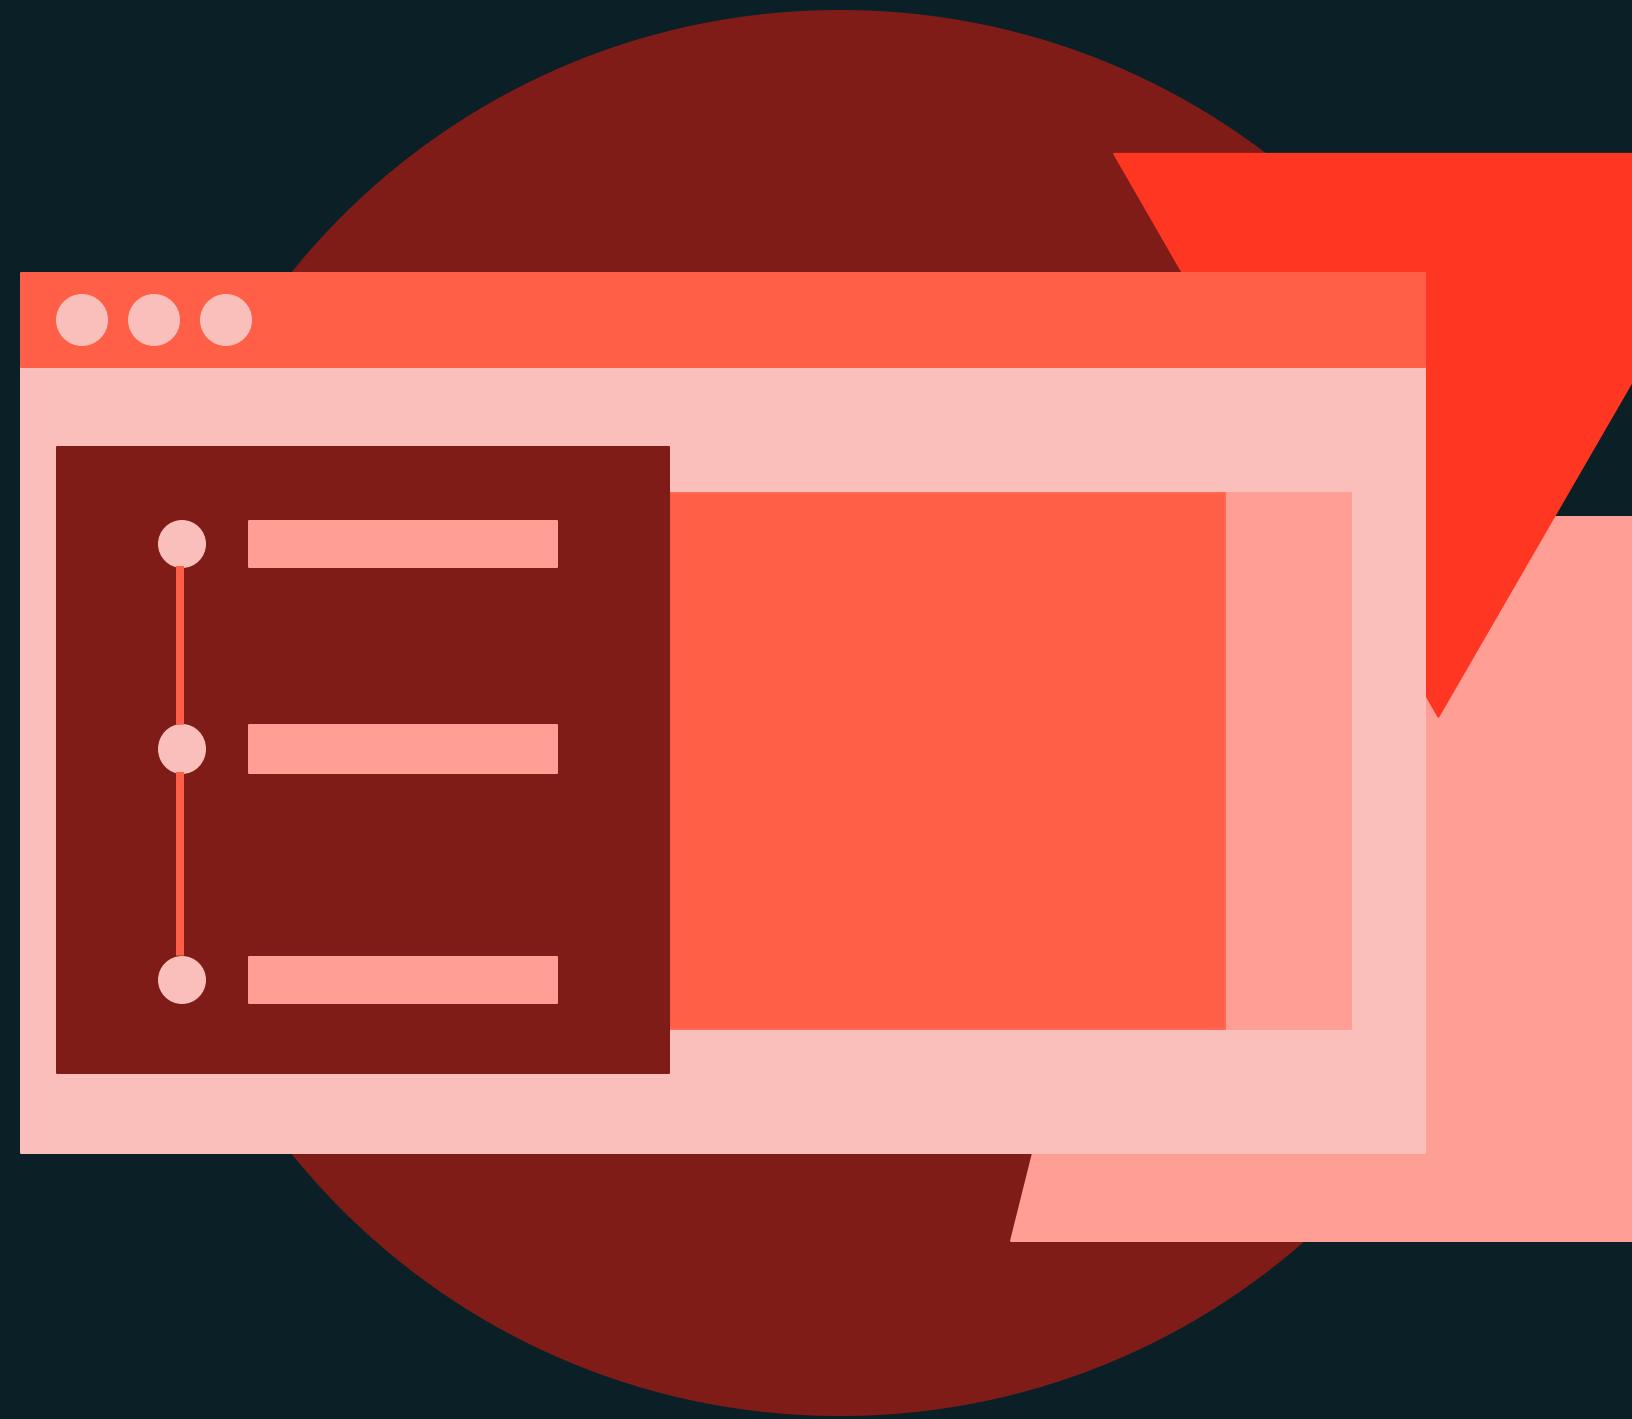




Managing and Exploring Data

**DEMONSTRATION**

# Load and Explore Data



# Demo

## Outline

### What we'll cover:

- Explore data with summary statistics
- Data visualization
  - Integrated visualization tools
  - External visualization tools
- Time-travel with Delta





Managing and Exploring Data

### LAB EXERCISE

# Load and Explore Data



# Lab

## Outline

### What you'll do:

- Read data from Delta table
- Manage data permissions
- Show summary statistics
- Use Data Profiler to explore data
- Time-travel to older versions of data
- Revert to previous versions of the Delta table

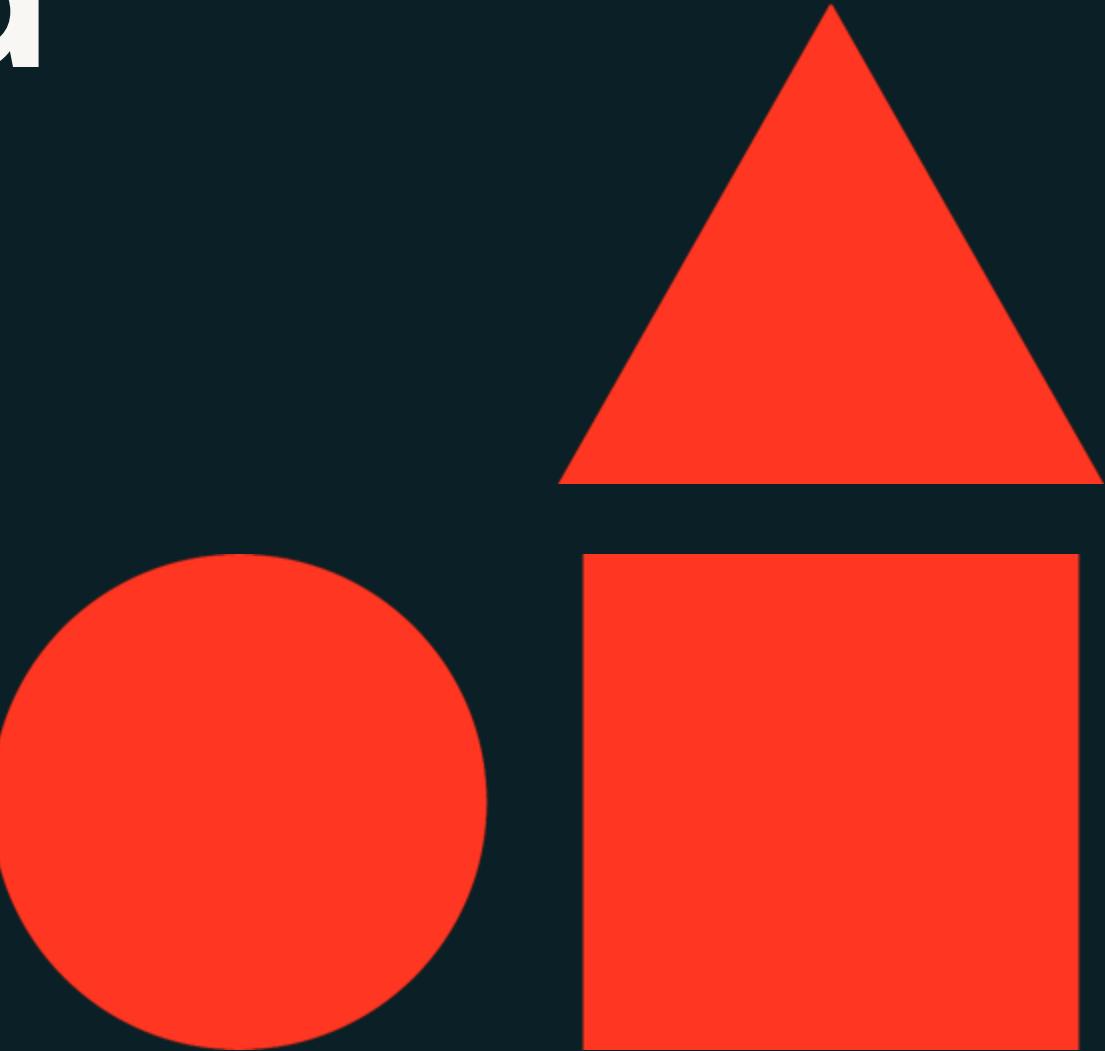




# Data Preparation and Feature Engineering

---

Data Preparation for Machine Learning



# Learning objectives

Things you'll be able to do after completing this module

- Explain data preparation and data splitting for ML, including train-test-holdout and cross-validation.
- Effectively handle missing values and understand the importance of indicator variables.
- Manage categorical features using encoding methods, considering efficiency and associated risks.
- Describe the benefits of feature standardization and recognize challenges in result interpretation.



# Learning objectives

Things you'll be able to do after completing this module

- Build a data imputation and transformation pipeline for diverse datasets.
- Develop an advanced feature engineering pipeline, along with the process of saving it for future use.





Data Preparation and Feature Engineering

LECTURE

# Fundamentals of Data Preparation and Feature Engineering

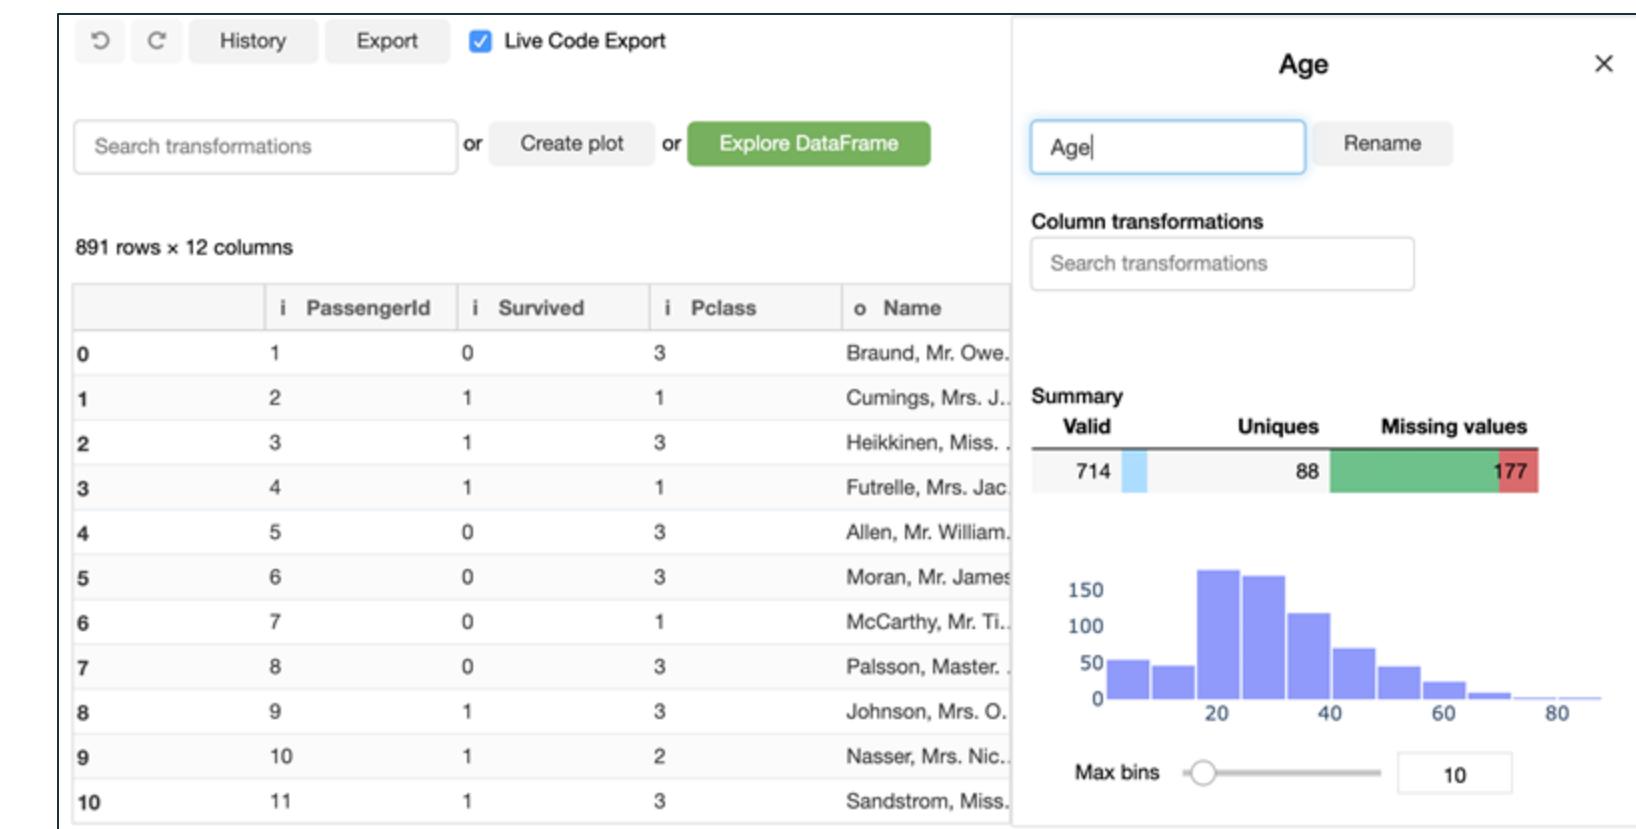


# Data Preparation for ML projects

Goal: Optimize input quality for accurate model predictions

Data preparation includes the following tasks:

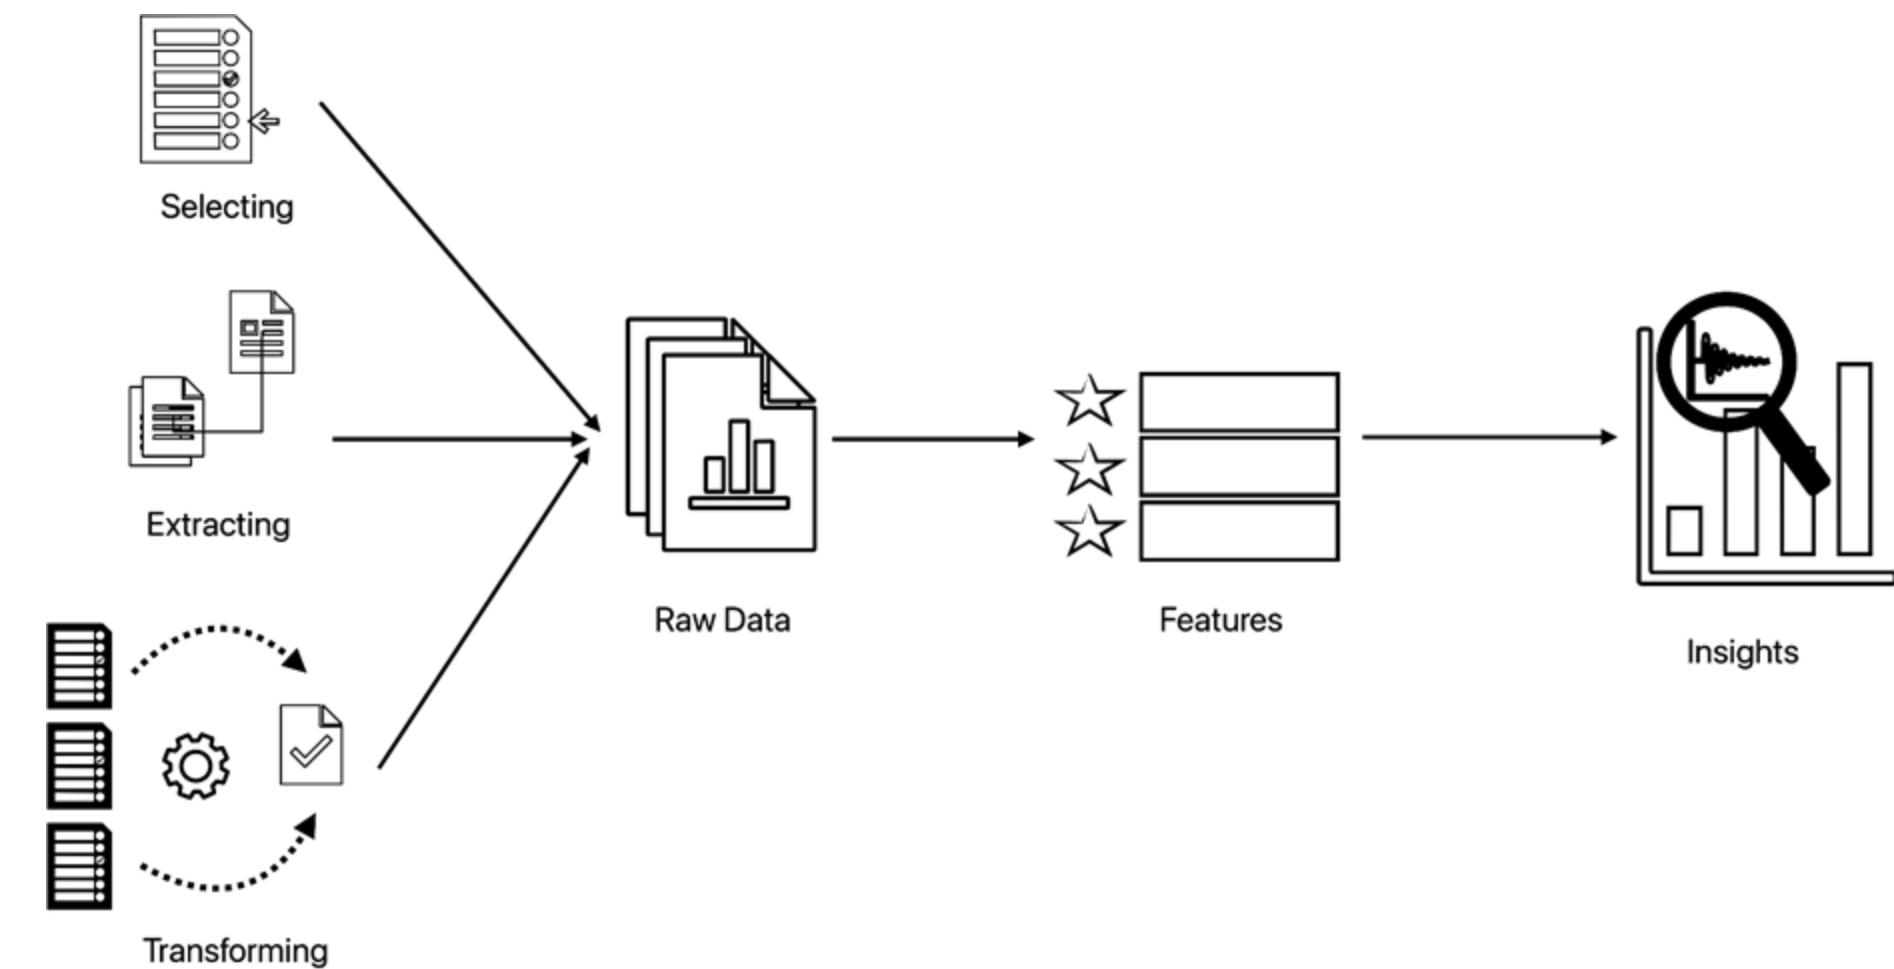
- **Cleaning and formatting data:** This includes tasks such as **handling missing values** or outliers, ensuring data is in the correct format, and removing unneeded columns.
- **Feature Engineering:** This includes tasks like **numerical transformations, aggregating data, encoding text or image data, and creating new features.**



# Feature Engineering

Transforming raw data into model-friendly features

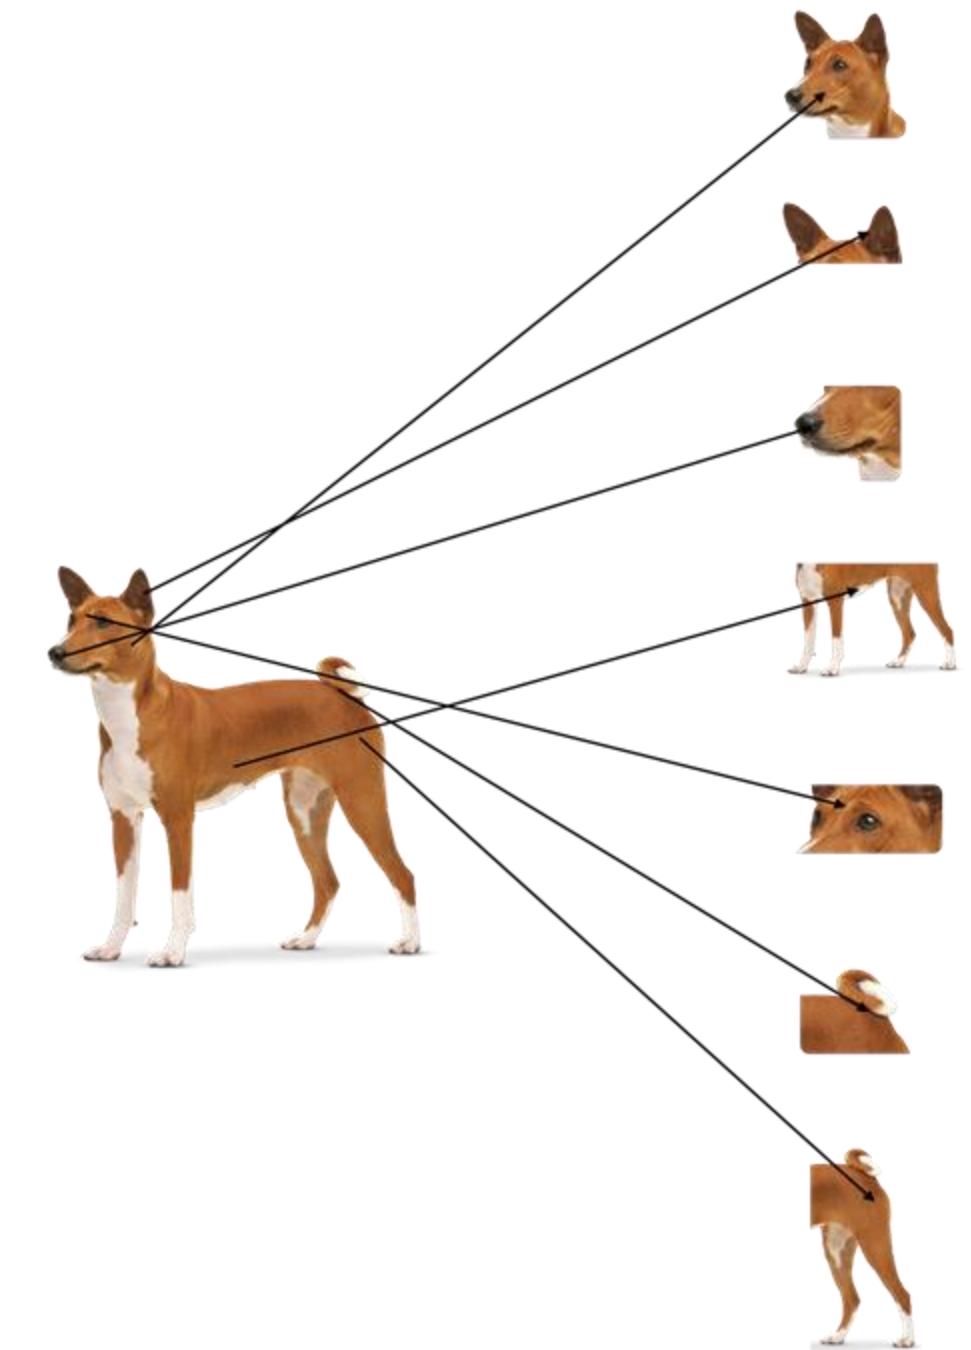
- It is a critical step in preparing data for machine learning models.
- It transforms raw data into a format that allows ML models to extract meaningful patterns and make accurate predictions.



# Feature Extraction

Transforming raw data into a set of features that better represent the underlying patterns

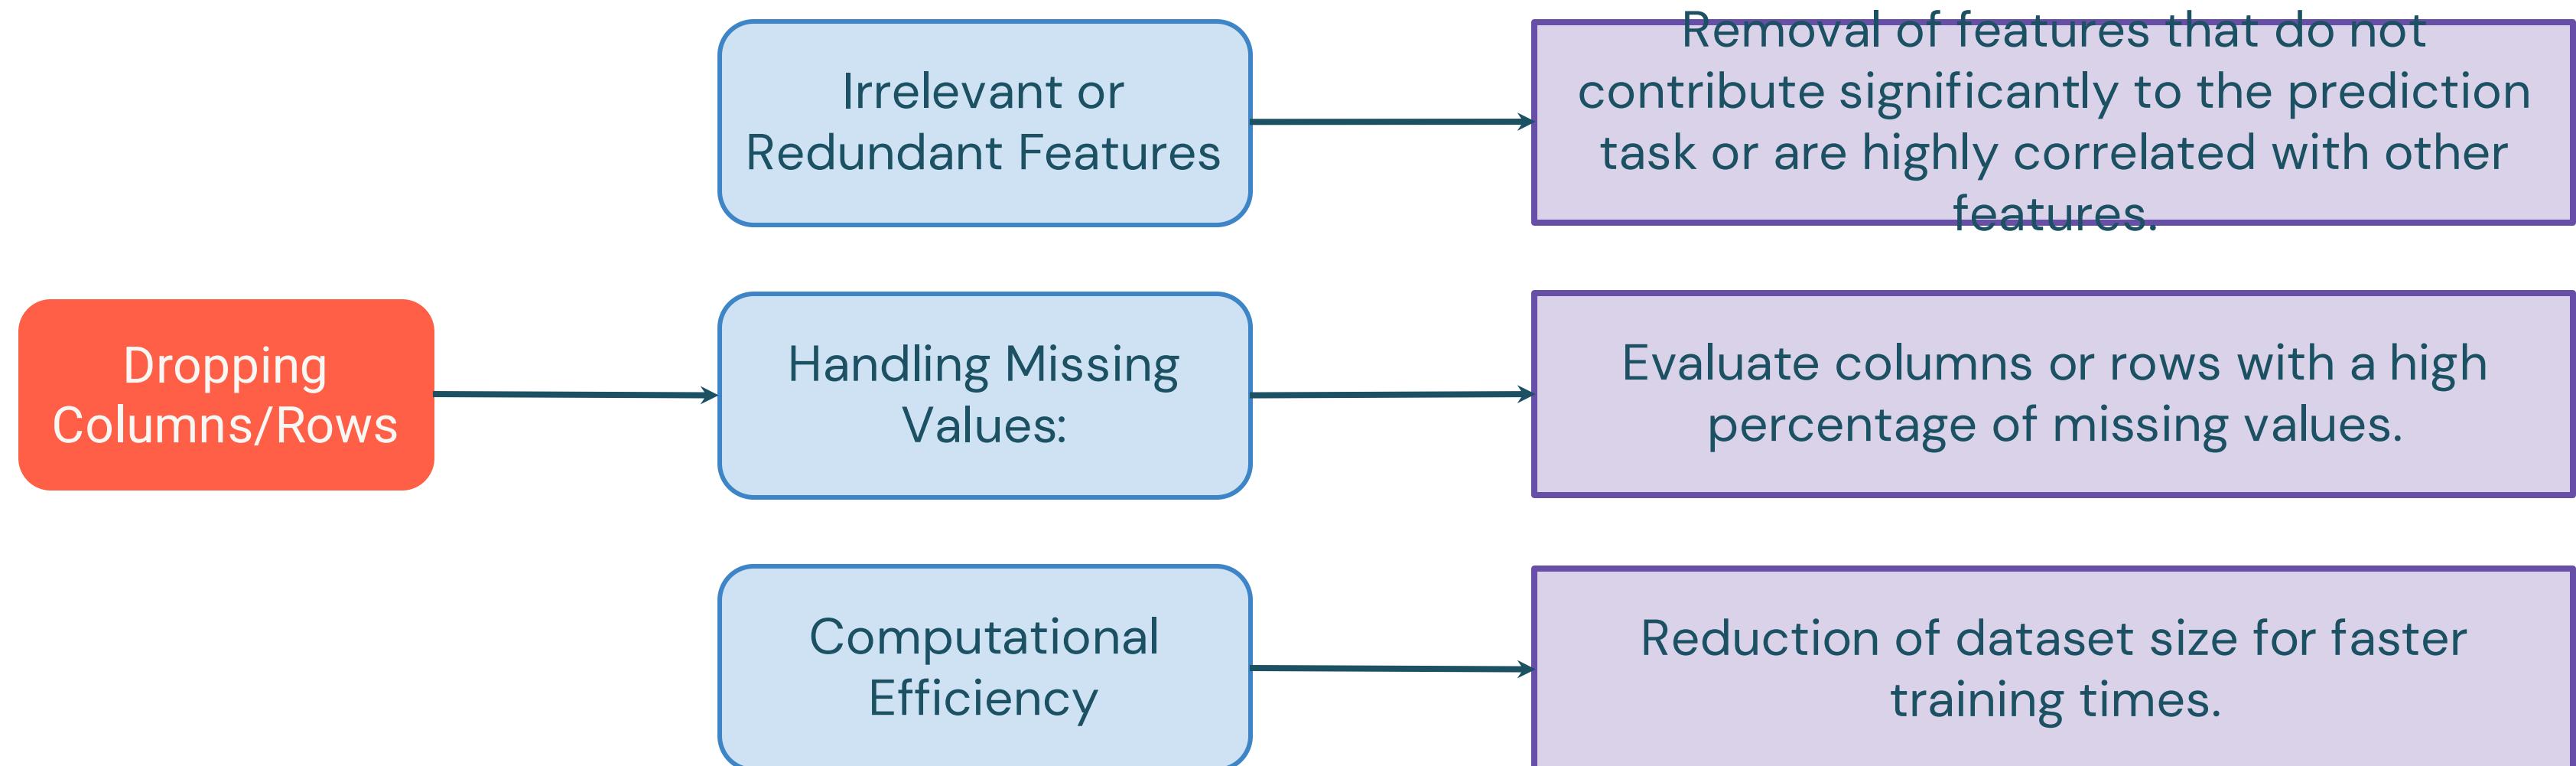
- Transforming Raw Data for Enhanced Modeling
- Dimensionality Reduction for Improved Performance
- Simplifying Feature Engineering



# Feature Selection: Dropping Columns/Rows

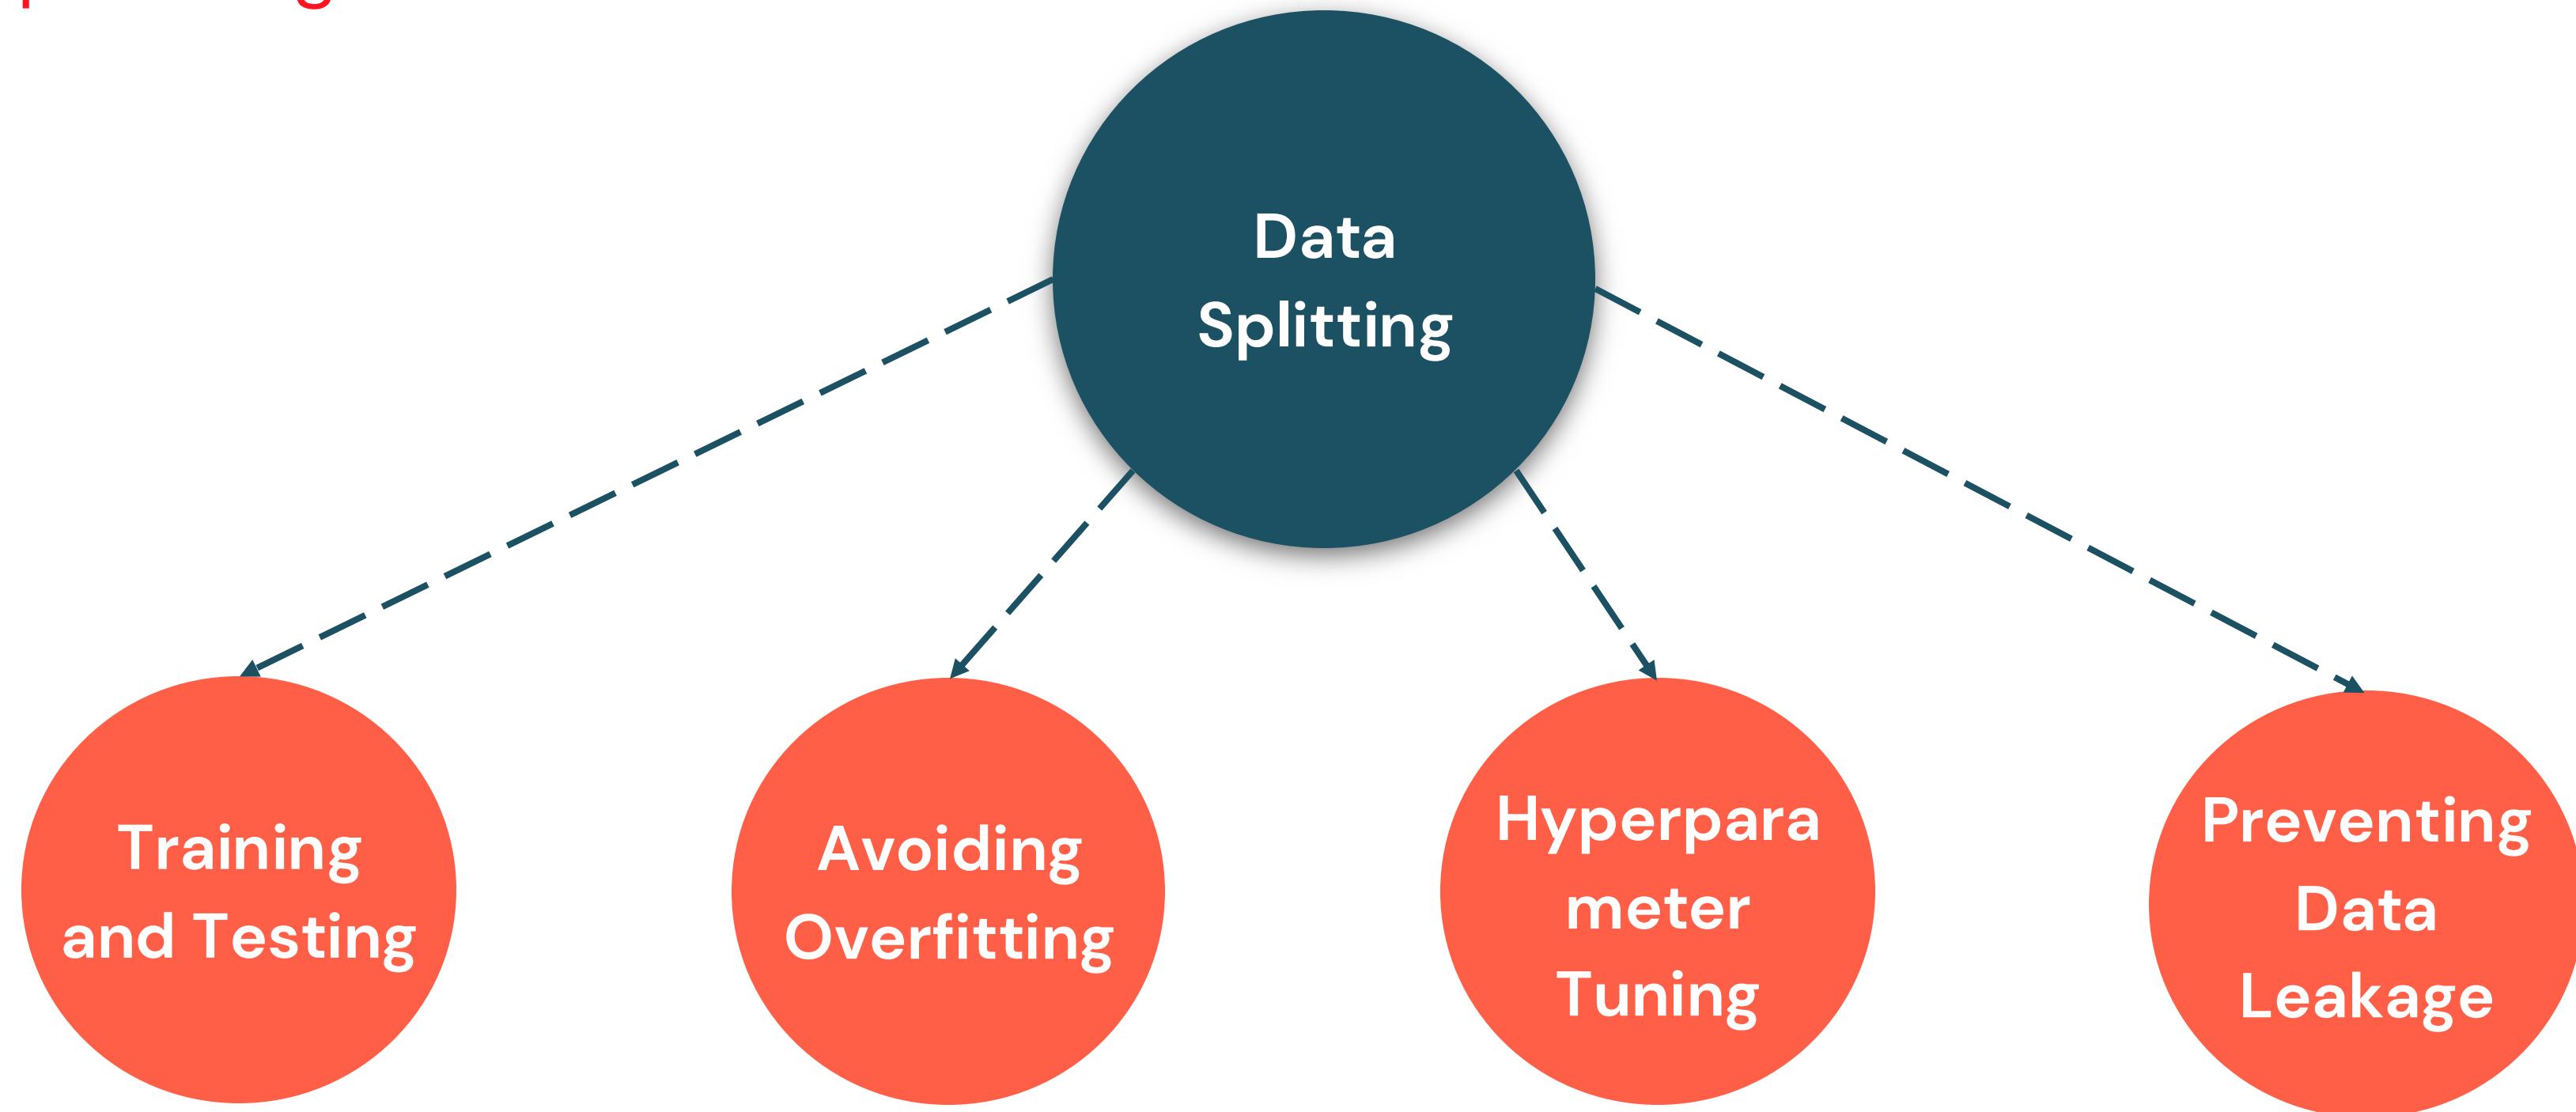
## Enhancing Model Efficiency

Streamlining the dataset by selectively removing columns (features) or rows (instances) to improve model efficiency and effectiveness.



# Why Do We Need Splitting Data

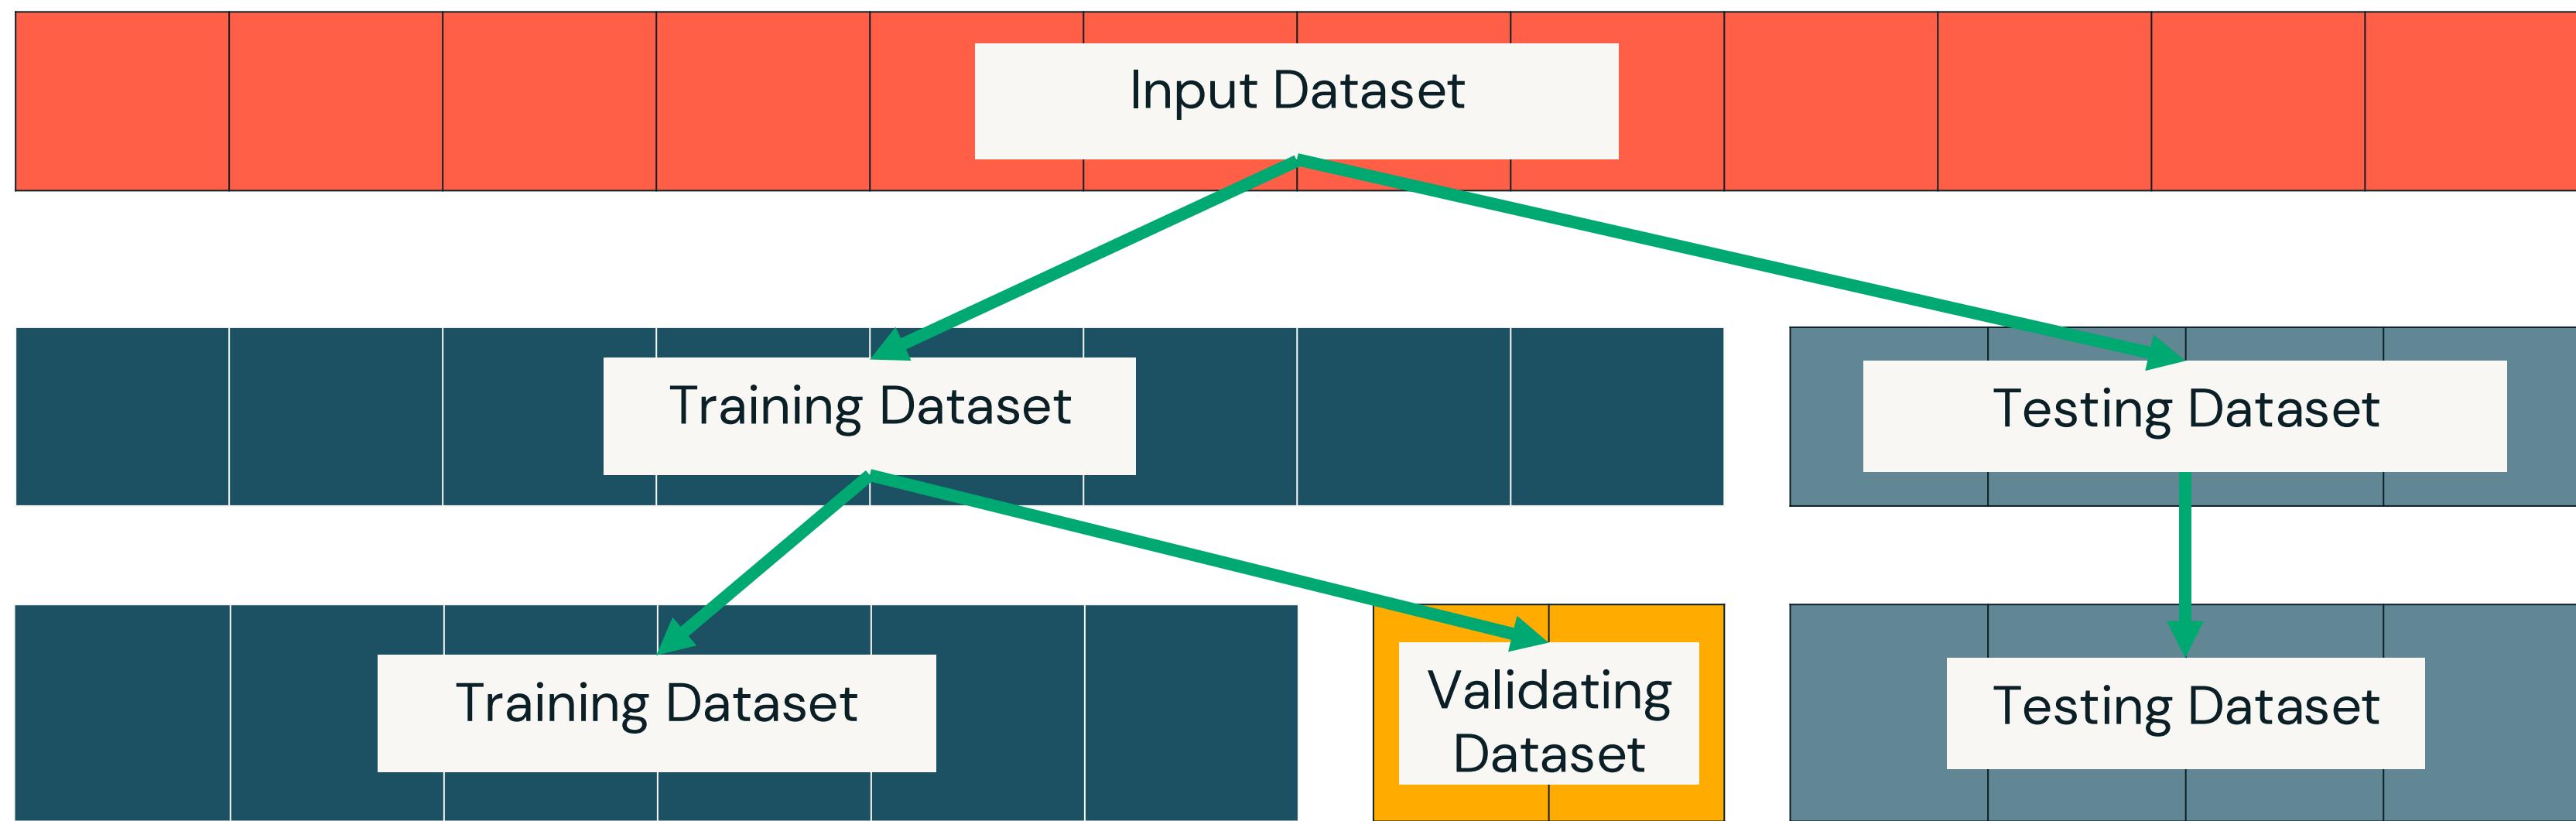
Optimizing Model Evaluation



# Splitting Data into Multiple Sets

## Optimizing Model Training, Validation, and Testing

Splitting the dataset is a fundamental step in data preparation to facilitate effective model training, validation, and testing.



# Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.

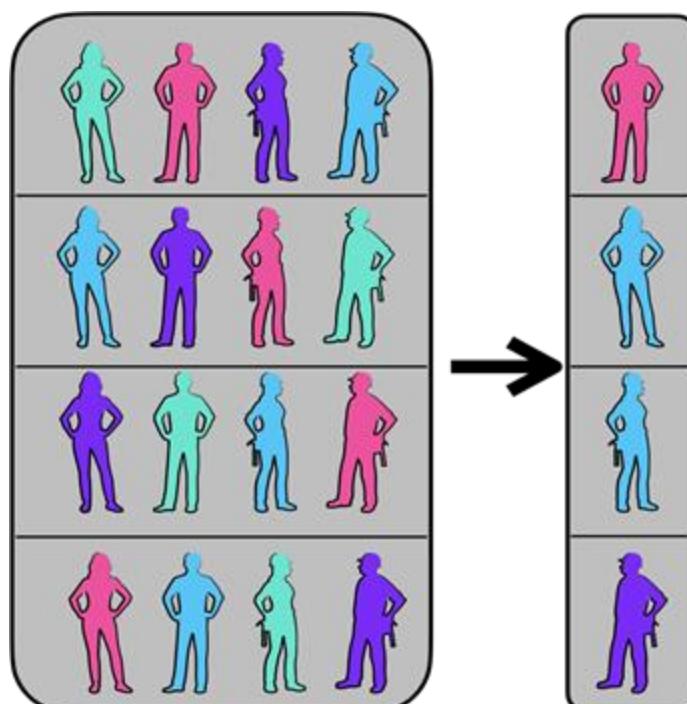


# Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.

Here's an overview of common sampling methods:

## Random Sampling:



This is the most common and basic sampling technique in machine learning. It randomly selects data points from a dataset without following any specific pattern, assuming that each data point has an equal probability of being chosen.

### Example with a pandas DataFrame

```
# my_data is a pandas DataFrame and we want to sample 50 data points  
  
my_random_sample = my_data.sample(n=50)
```

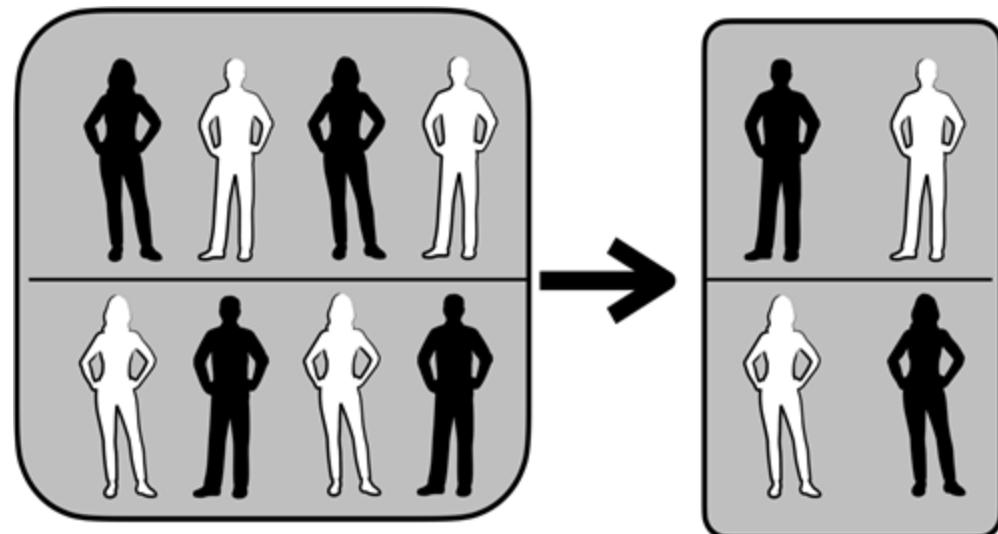


# Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.

Here's an overview of common sampling methods:

## Stratified Sampling:



This sampling method is used when a dataset consists of subgroups, ensuring that samples are taken from each. The population is divided into homogeneous subgroups, known as strata, and an appropriate number of instances is selected from each stratum to ensure the sample accurately represents the entire population.

## Example with a pandas DataFrame & scikit-learn

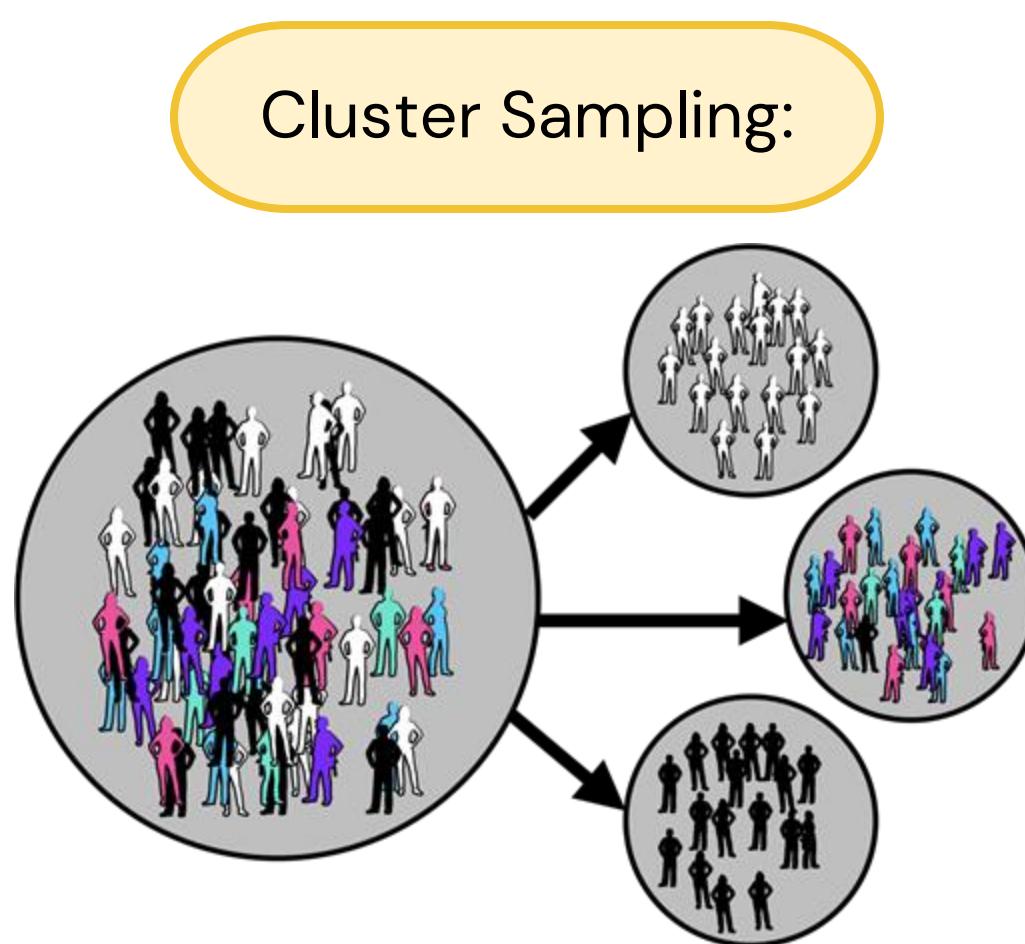
```
# my_data is a pandas DataFrame  
# group_label is the column containing subgroups  
  
from sklearn.model_selection import train_test_split  
  
stratified_train, stratified_test = train_test_split(my_data, test_size =0.2, stratify =  
my_data['group_label'])
```



# Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.

Here's an overview of common sampling methods:



Cluster sampling is a sampling technique where the population is divided into naturally occurring groups, or clusters, and a subset of these clusters is randomly selected for study. Instead of sampling individuals directly, researchers collect data from all individuals within the chosen clusters.

Example with a pandas DataFrame and numpy

```
# my_data is a pandas DataFrame  
# my_col is the column containing subgroup identifiers  
# Randomly select M clusters from the range 1 to N  
  
import pandas as pd  
import numpy as np  
  
clusters = np.random.choice(np.arange(1, N+1), size=M, replace=False)  
cluster_sample = my_data[my_data['my_col'].isin(clusters)]
```

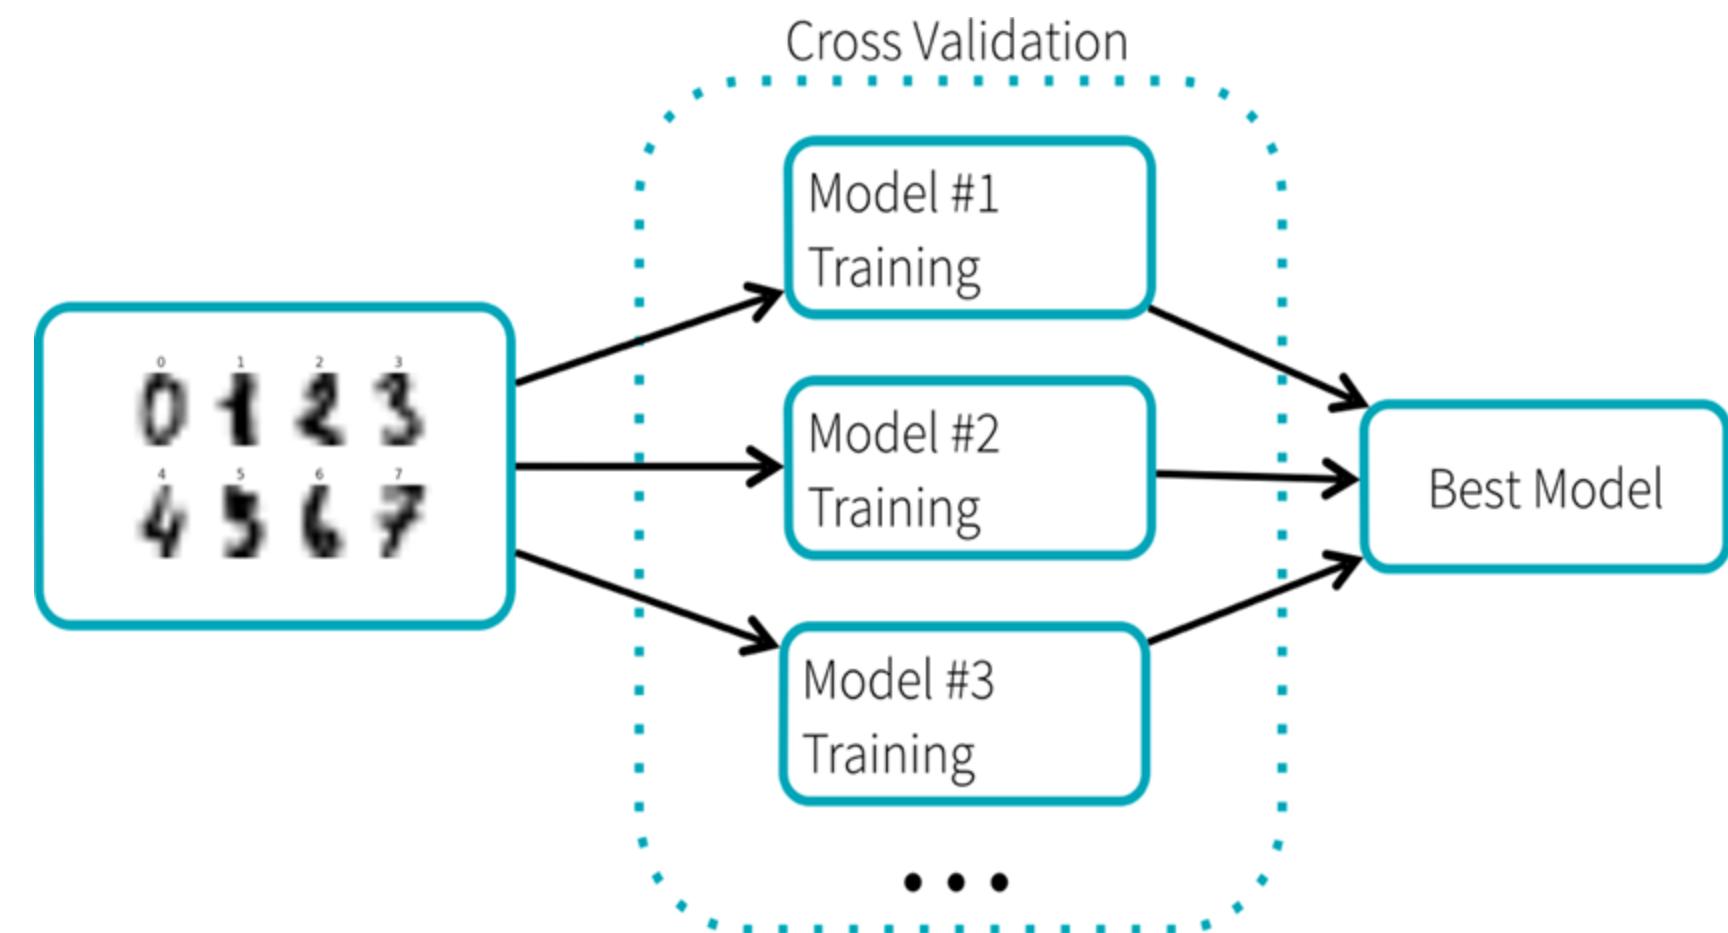


# Cross-Validation

Cross-validation is a resampling technique used to assess the performance of a machine learning model. It involves partitioning the dataset into subsets to train and evaluate the model multiple times, providing a more robust estimate of its generalization performance.

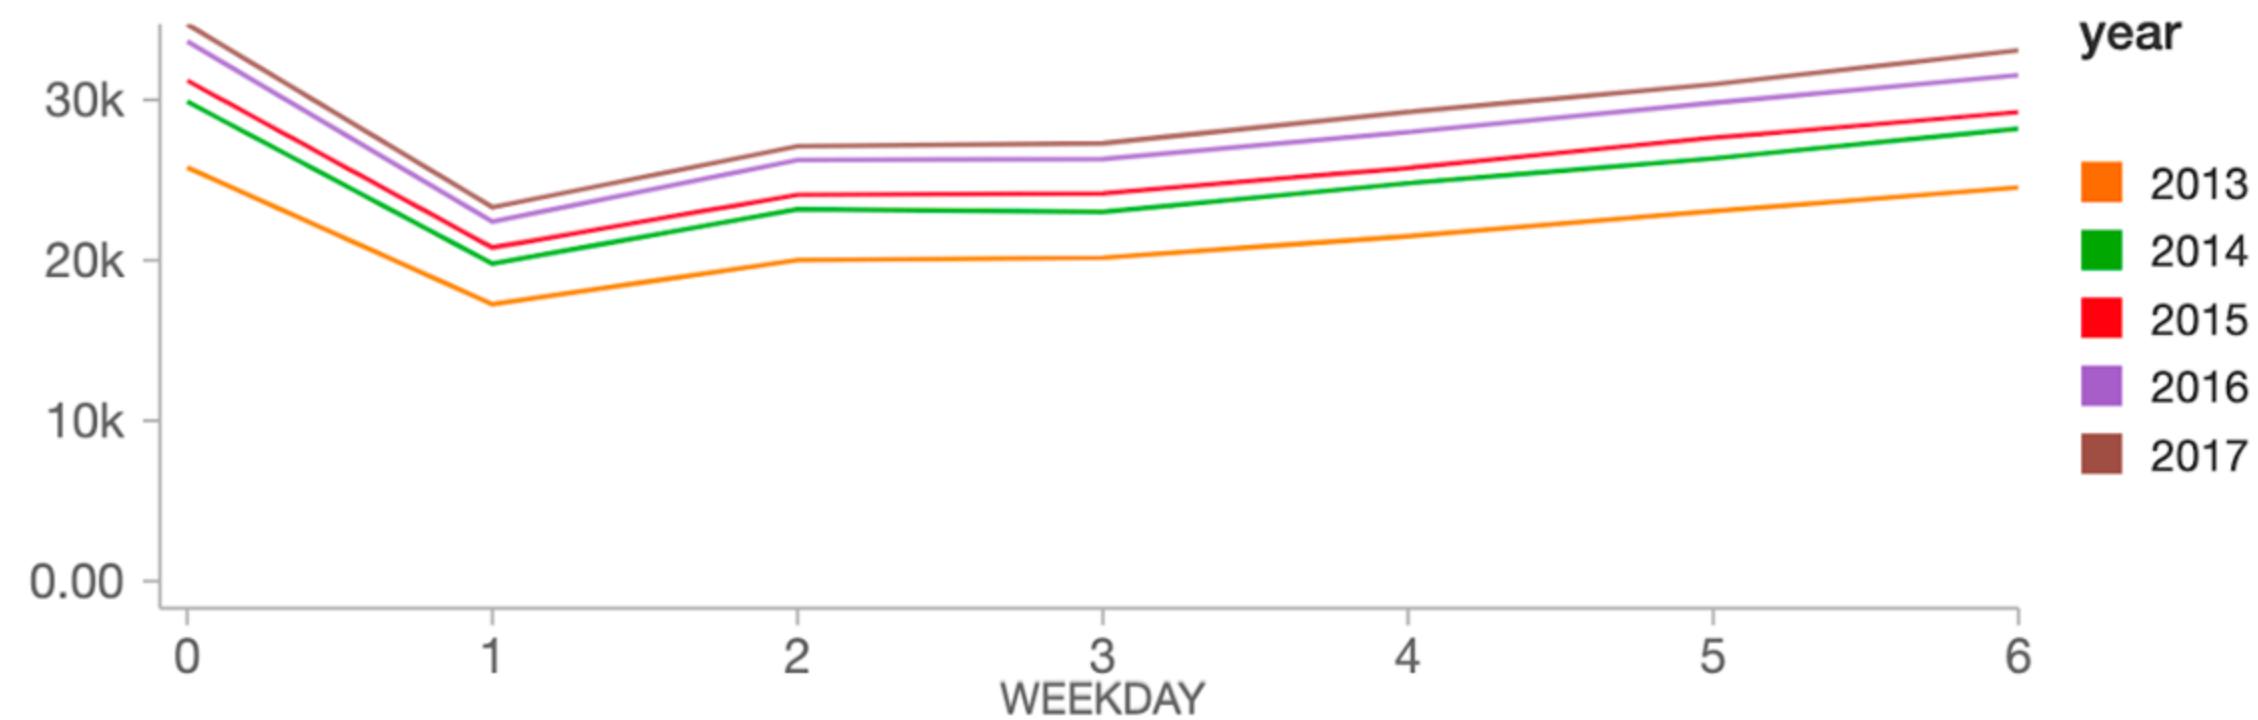
**Here are the key elements of cross-validation:**

- K-Fold Cross-Validation
- Stratified K-Fold Cross-Validation
- Shuffle Split Cross-Validation
- Nested Cross-Validation



# Sampling for Time-Series Data

Sampling for **time-series data requires special consideration** due to the temporal dependencies inherent in the data. **Traditional random sampling or shuffling may not be appropriate**, as the order of events in time is crucial.





Data Preparation and Feature Engineering

LECTURE

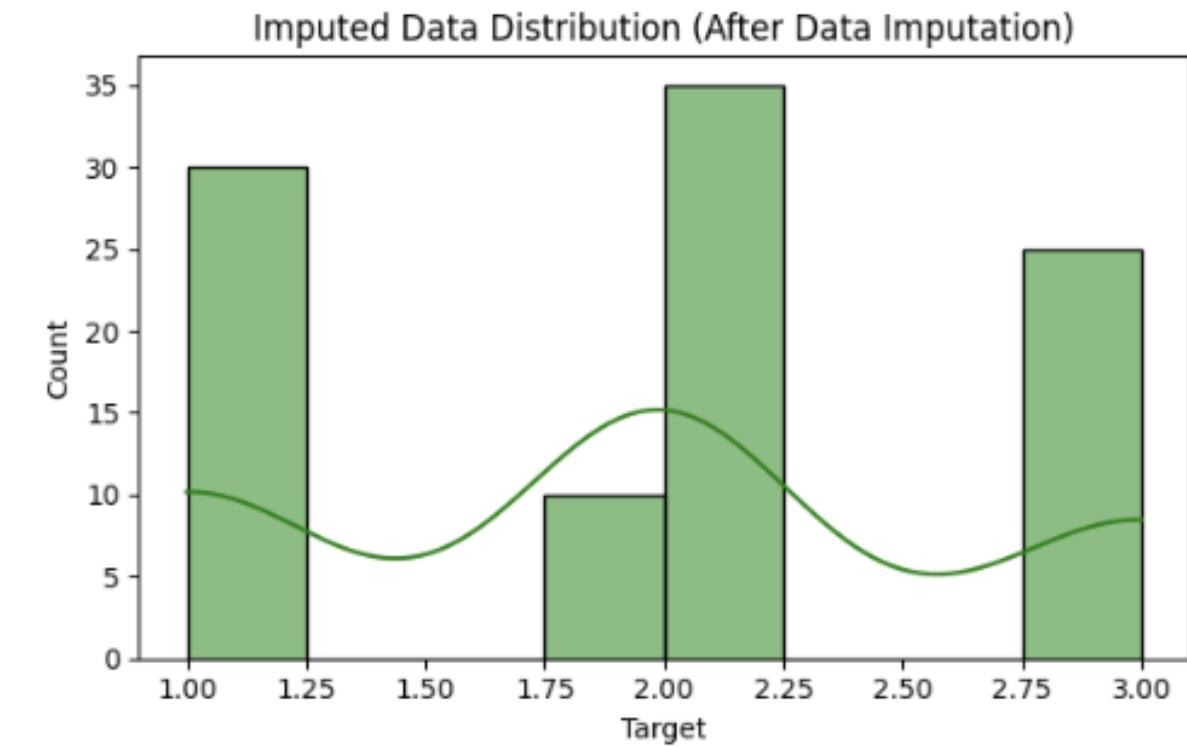
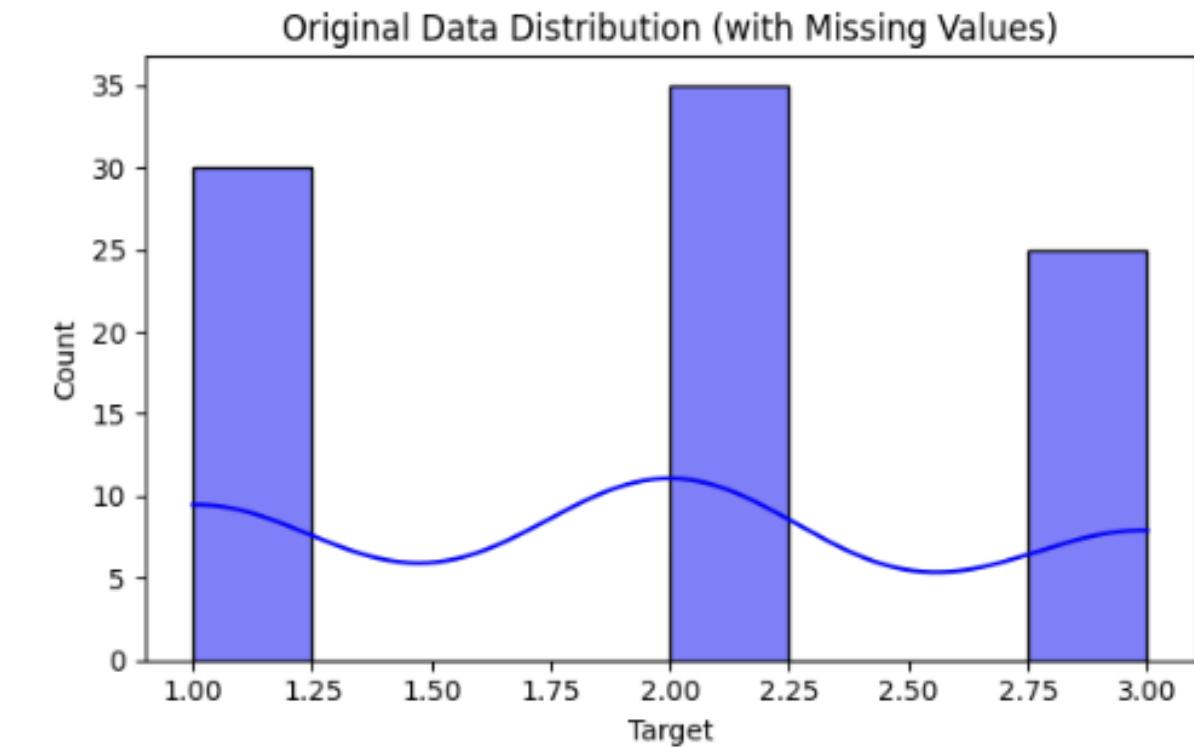
# Data Imputation



# Data Imputation

Data imputation is **the process of filling in missing values** in a dataset with estimated or predicted values.

The goal of data imputation is to enhance the quality and completeness of the dataset, ultimately improving the performance and reliability of the machine learning model.

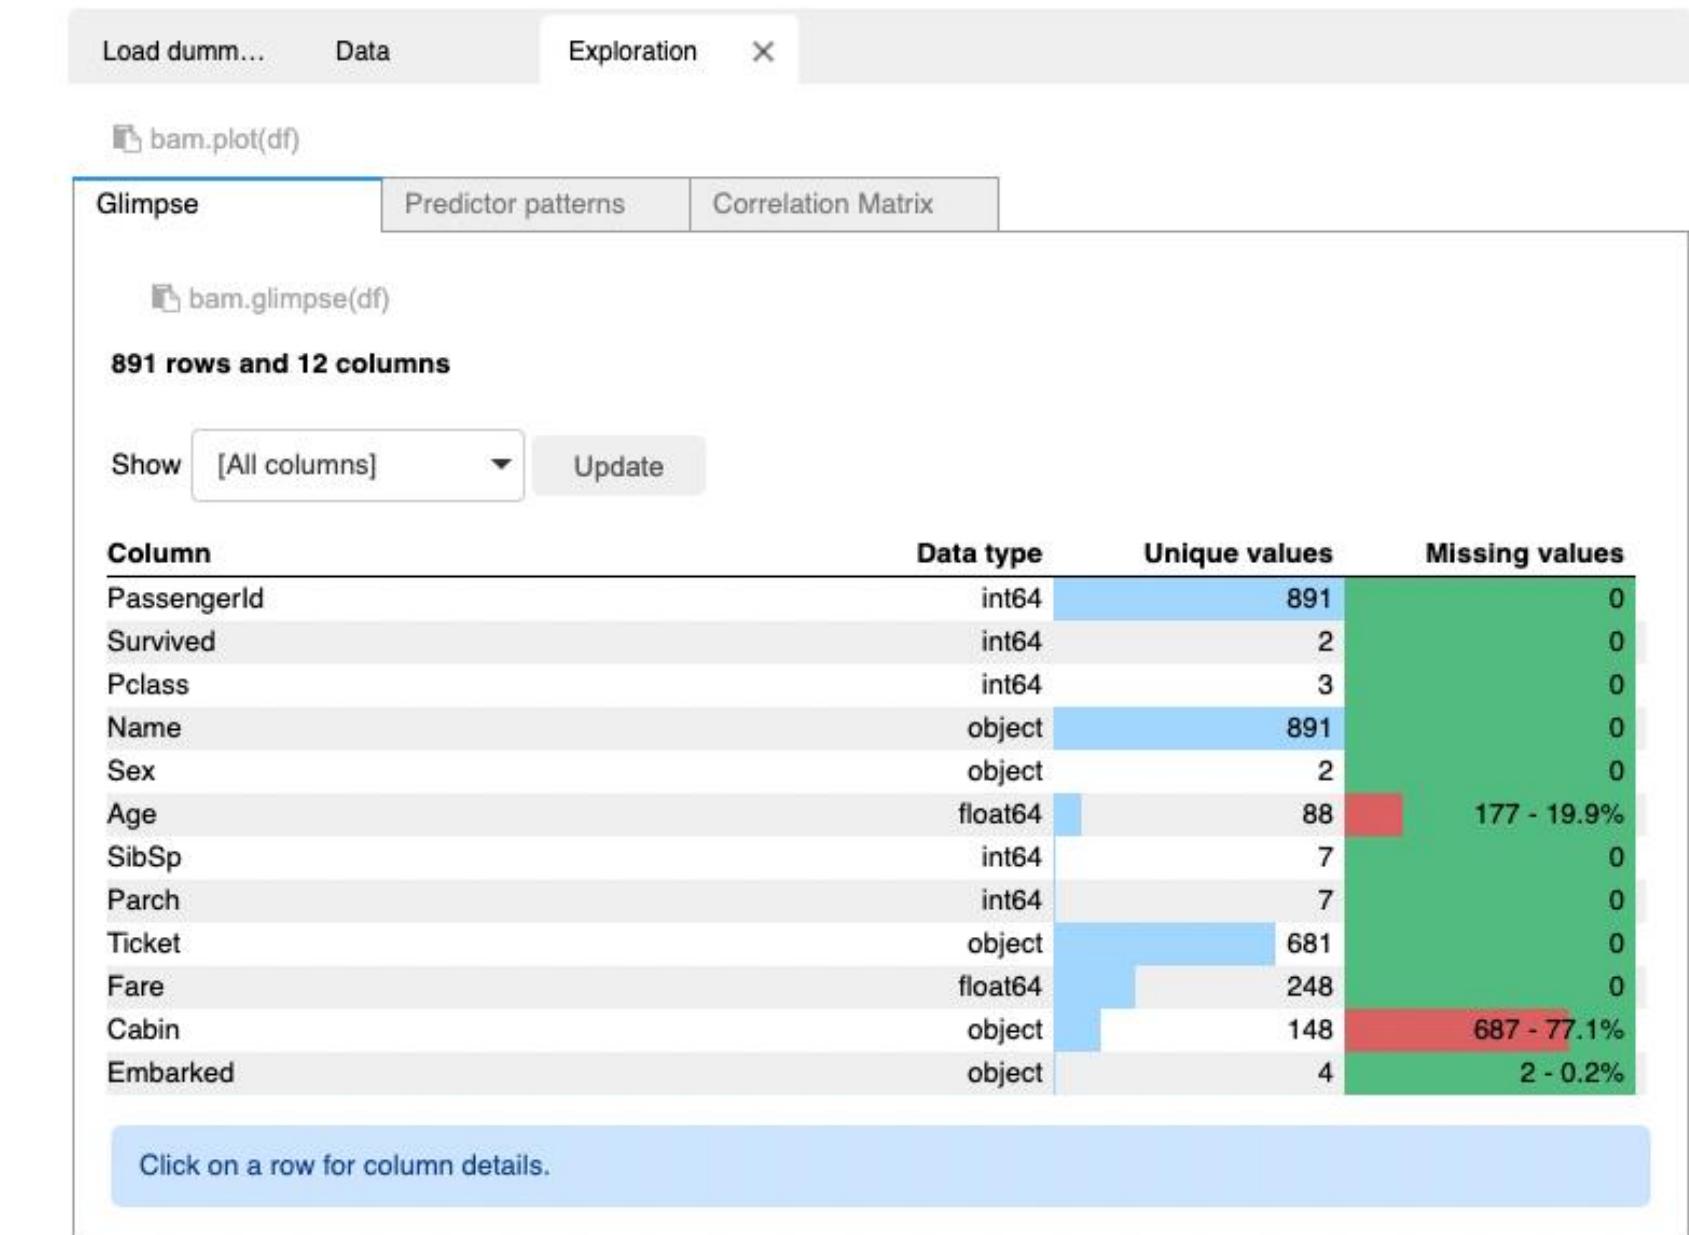


# Problems with Missing Data

Impacting the performance and reliability of ML models

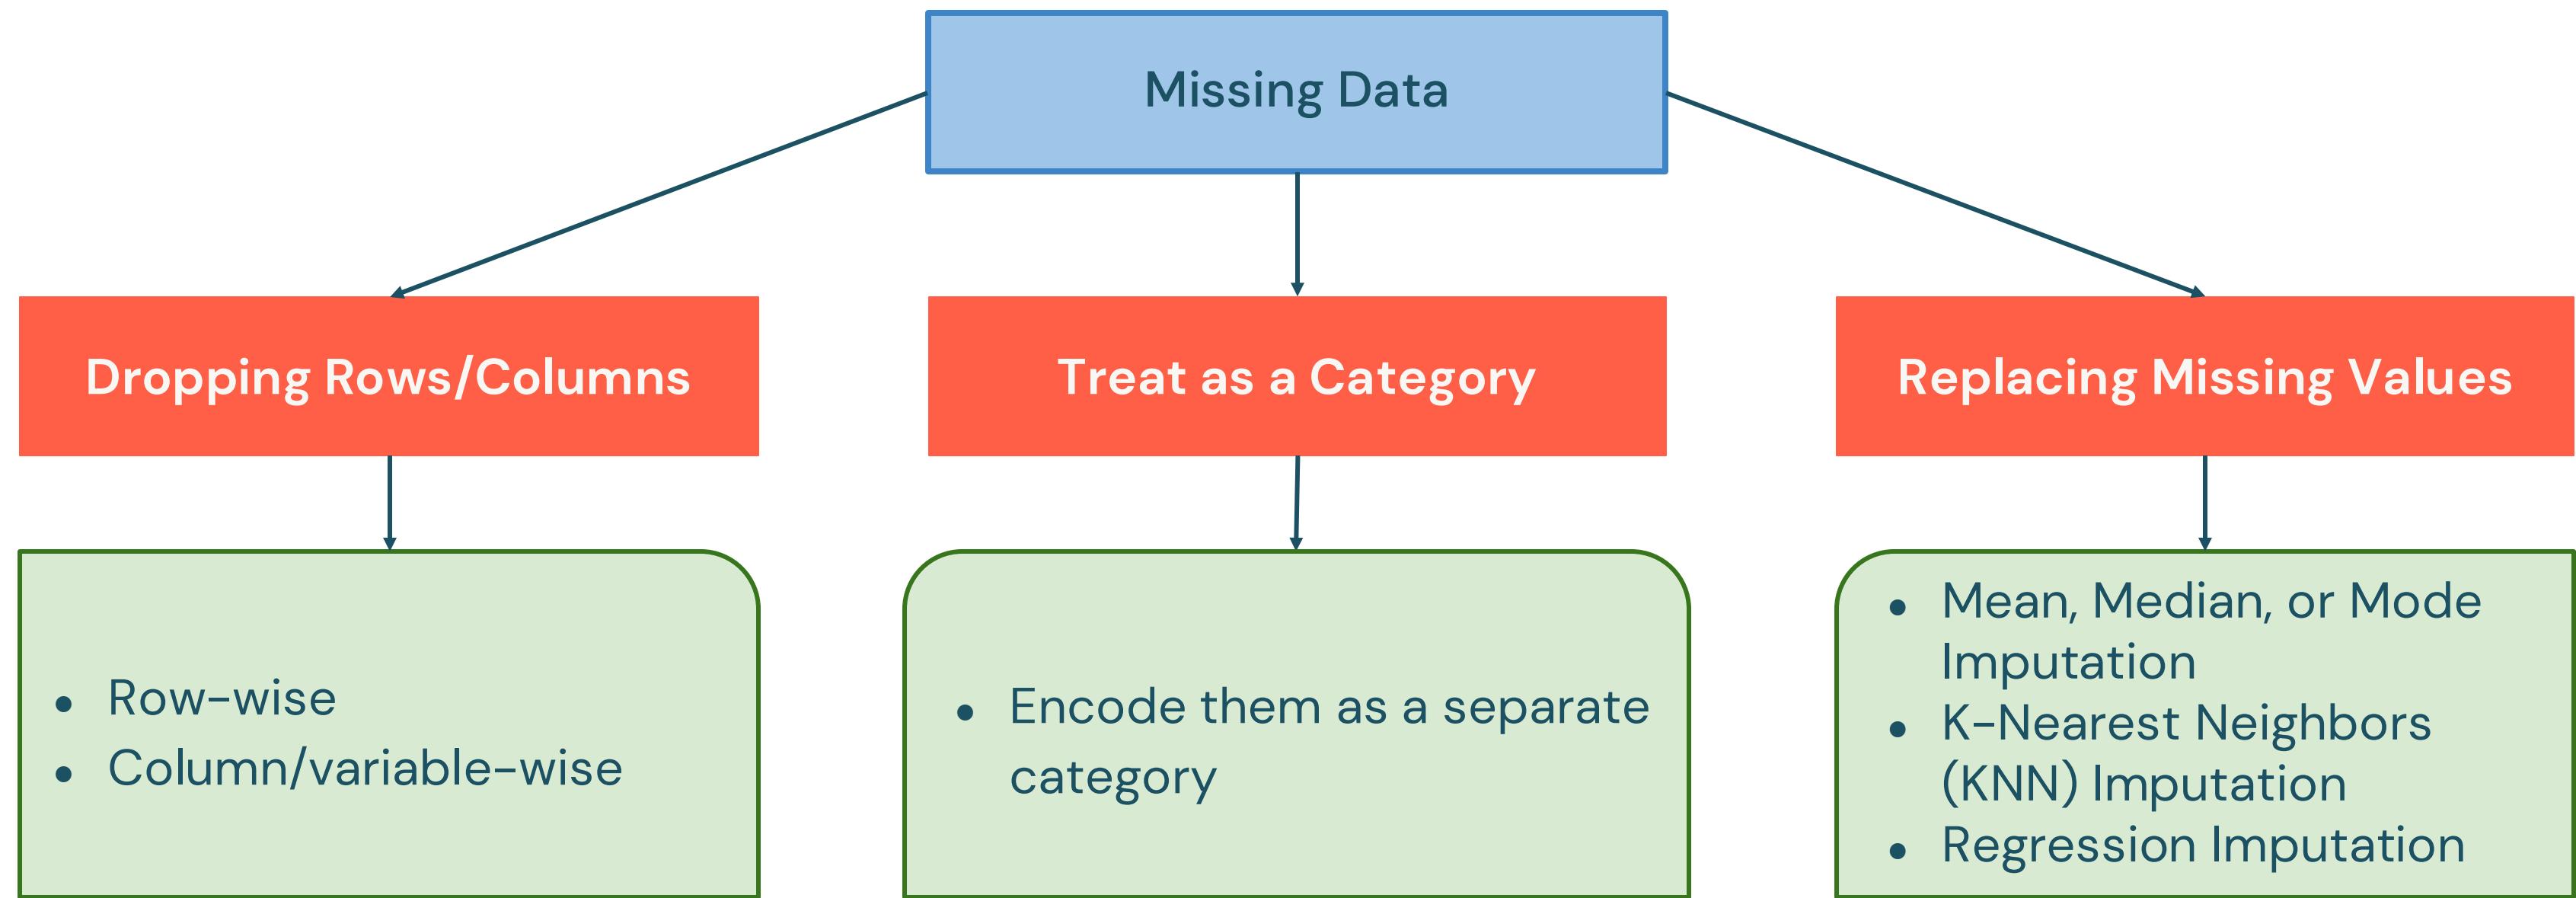
- Reduced Model Performance
- Biased Inferences
- Imbalanced Representations
- Increased Complexity in Model

Handling



# How to Handle Missing Data

## Data imputation methods



# Replacing Missing Values

## Data imputation methods

### Mean - Mode Imputation

Before	After
10	10.0
15	15.0
-	18.3
20	20.0
25	25.0

### K-Nearest Neighbors (KNN with k=2)

Before	After
8	8.0
-	10.0
12	12.0
15	15.0
-	13.0

### Multiple Imputation (Regression)

F1	F2	Before	After
X	X	10	10.0
X	X	15	15.0
X	$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$		Y
X	X	20	20.0
X	X	25	25.0



# Factors Influencing Imputation Method

There isn't a definitive best approach

## Nature of Data

- Is data type is **continuous, ordinal or categorical?**
- Data **distribution**: For example, median imp. Might work better for non-normal distributions

## Amount of Missing Data

- How much of data is missing?
- Will imputing data improve the quality of dataset or will it add more bias?
- Multiple imp. and KNN might better for large missingness.

## Form of Missingness

- Is data is **missing at random** or is there a **missingness pattern**.
- Domain knowledge and how would imputation affect downstream analysis.



# Marking Imputed Data (*Best Practice*)

Keep track of imputed data

Important to mark imputed data, for:

- Model Evaluation
- Data Quality Assessment
- Enabling Transparency of Dataset
- Error Identification

ID	Name	Age	Age_imputed
1	Alice	25.0	0
2	Bob	30.0	0
3	Charlie	26.0	1
4	David	28.0	0
5	Eva	22.0	0





Data Preparation and Feature Engineering

LECTURE

# Data Encoding



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Data Encoding

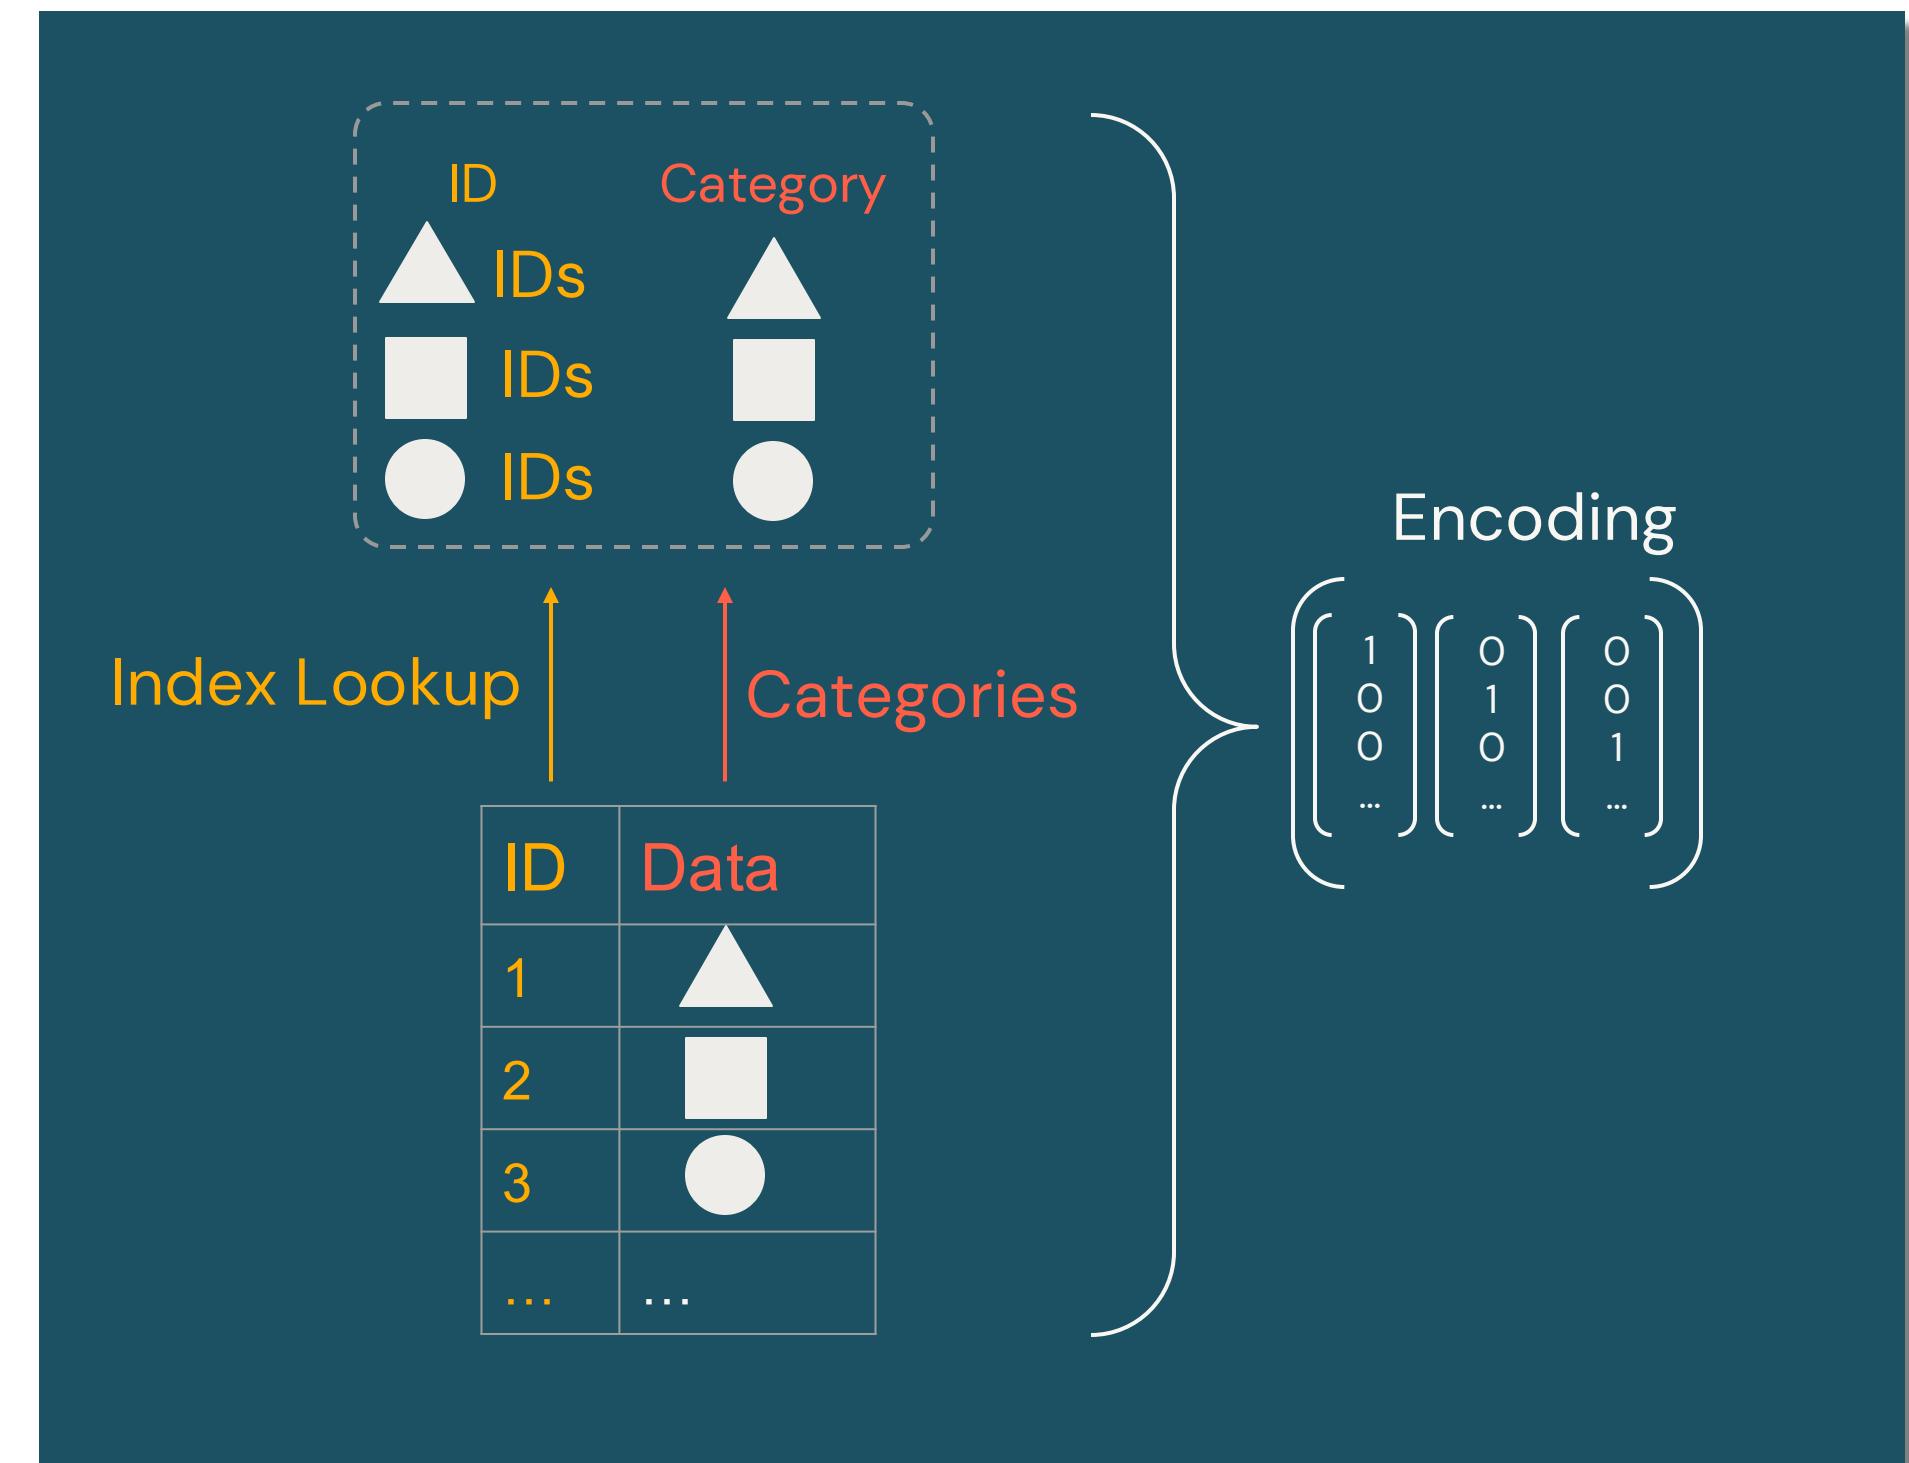
## Why Encoding?

Data encoding is a important pre-processing step in **preparing categorical data for machine learning algorithms**, as vast majority of algorithms accept numerical input exclusively.

## Issues with classic ML

- Handling high-cardinality features
- Introducing unintended relationships
- Overfitting
- Increased computational cost
- Possible lack of interpretability

## Encoding Process



# Working with Categorical Features

## High Cardinality

### Issue: Large set of categories

Many unique values in a categorical feature can lead to a large number of dummy variables, increasing dimensionality and potentially causing issues. This is known as the **high-cardinality problem**.

### Possible Solutions:

#### Group Rare Categories:

ID	Category
0	A
1	B
2	A
3	C
4	D
5	D
6	D
7	D

ID	Category	Cat_Group
0	A	A
1	B	Rare
2	A	A
3	C	Rare
4	D	D
5	D	D
6	D	D
7	D	D

#### Top-N Categories:

ID	Category
0	A
1	B
2	A
3	C
4	D
5	D
6	D
7	D

ID	Category	Cat_Group
0	A	A
1	B	Other
2	A	A
3	C	Other
4	D	D
5	D	D
6	D	D
7	D	D



# Working with Categorical Features

## Missing Values

### Issue: Categorical gaps

Categorical features often have **missing values**, which need to be addressed before model training.

## Possible Solutions:

### Imputation:

ID	Feature A	Feature B
0	1.0	10.0
1	2.0	NaN
2	NaN	30.0
3	4.0	40.0
4	5.0	50.0

ID	Feature A	Feature B
0	1.0	10.0
1	2.0	32.5
2	3.0	30.0
3	4.0	40.0
4	5.0	50.0

### Consider Missing as a Separate Category:

ID	Feature A	Feature B
0	1.0	10.0
1	2.0	NaN
2	NaN	30.0
3	4.0	40.0
4	5.0	50.0

ID	Feature A	Feature B
0	1.0	10.0
1	2.0	-1.0
2	-1.0	30.0
3	4.0	40.0
4	5.0	50.0



# Working with Categorical Features

## Encoding Categories

### Issue: String types

Models require numerical input, and categorical variables need to be **encoded**.

## Possible Solutions:

### One-Hot Encoding

ID	Category
0	A
1	B
2	A
3	C
4	A

ID	Cat_A	Cat_B	Cat_C
0	1	0	0
1	0	1	0
2	1	0	0
3	0	0	1
4	1	0	0

### Label Encoding:

ID	Category	Label
0	A	0
1	B	1
2	A	0
3	C	2
4	A	0

### Ordinal Encoding:

ID	Category	Ordinal
0	A	0.0
1	B	1.0
2	A	0.0
3	C	2.0
4	A	0.0



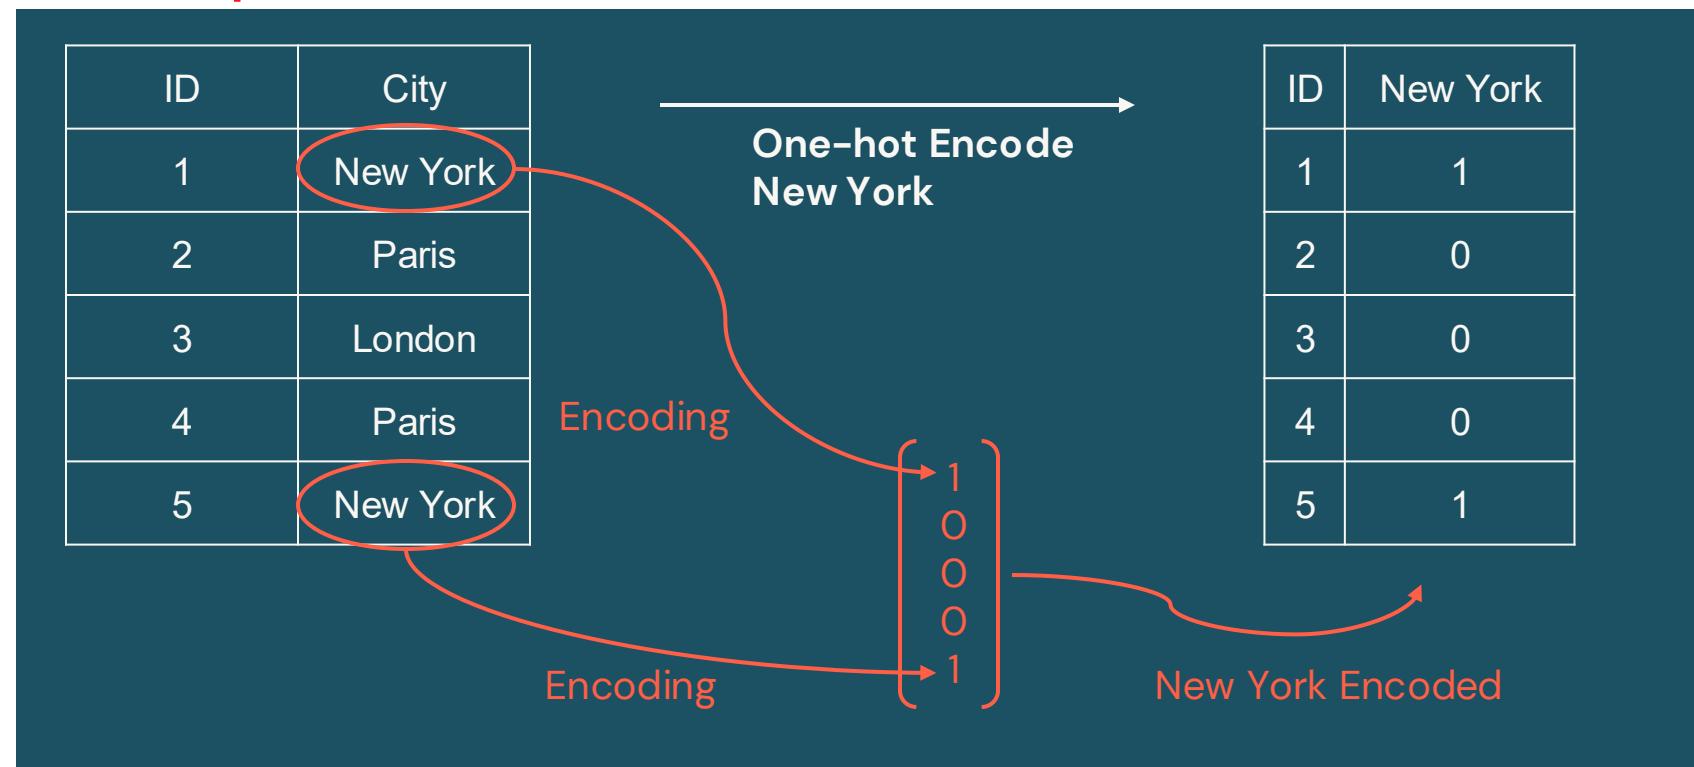
# One-hot Encoding

Represent categorical variables as binary vectors

## Procedure:

1. Add binary columns (0 or 1) for each unique value in categorical columns.
2. Each category is now a column that map to either 0 or 1 in reference to the ID column

## Example:



ID	City
1	New York
2	Paris
3	London
4	Paris
5	New York

↓

ID	New York	Paris	London
1	1	0	0
2	0	1	0
3	0	0	1
4	0	1	0
5	1	0	0



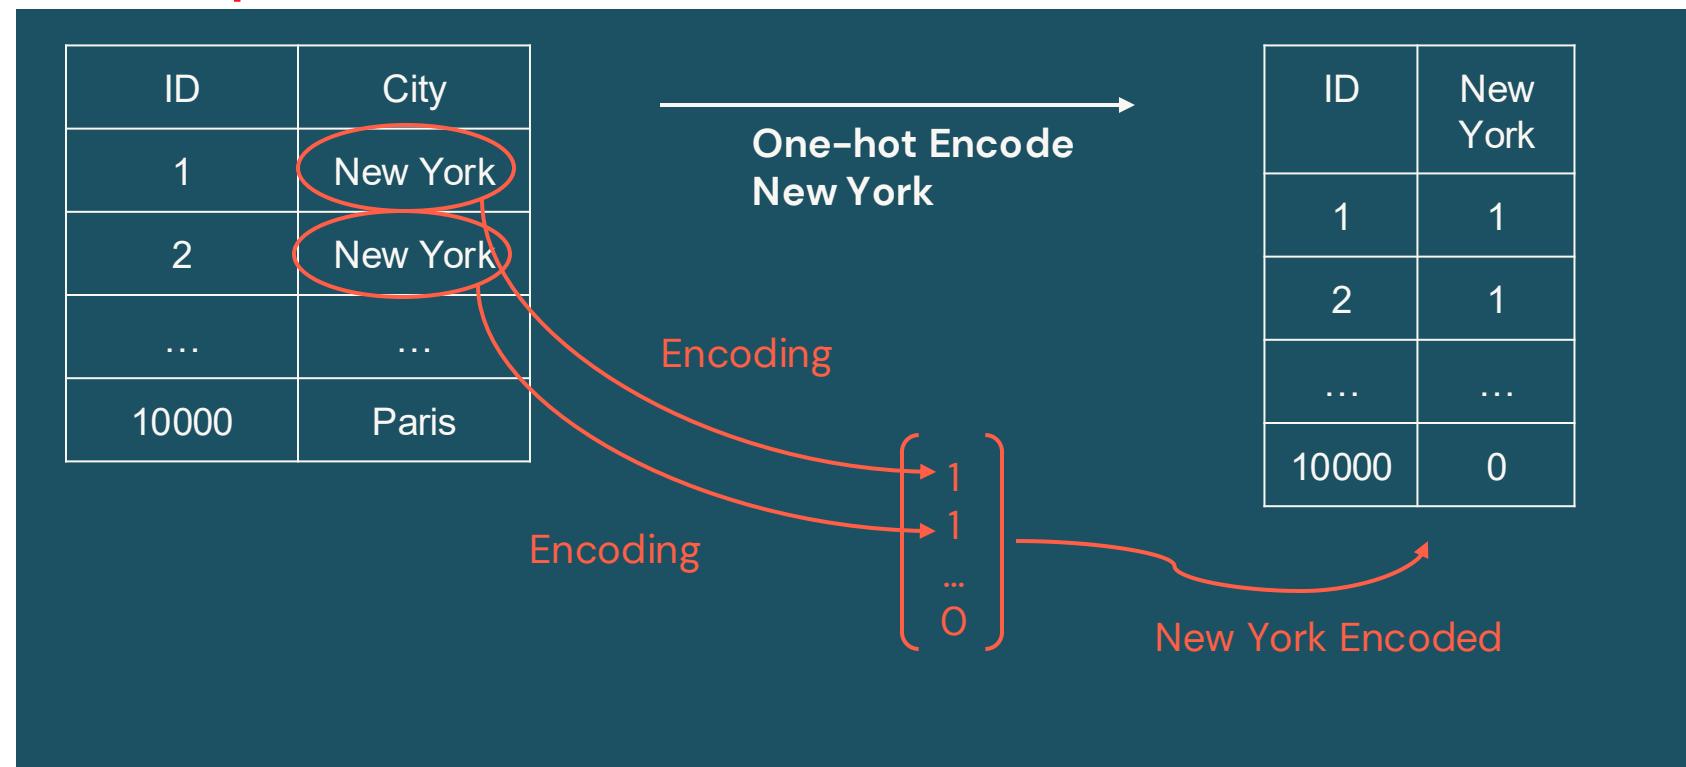
# One-hot Encoding

Represent categorical variables as binary vectors

## Drawbacks:

1. Induces sparsity
2. Limits split options for tree-based models
3. Inefficiency with high-cardinality variables
4. Obscures feature importance
5. Increases computational costs

## Example:



ID	City
1	New York
2	Paris
3	London
4	Paris
5	New York

ID	New York	Paris	London
1	1	0	0
2	0	1	0
3	0	0	1
4	0	1	0
5	1	0	0



# Label Encoding for Ordinal Features

Convert categorical data into numerical labels (*aka “String Indexing”*)

## Procedure:

1. **Assign numeric labels:** Map each category to a numeric value based on its natural order.
2. **Transform the Feature:** Replace each categorical value in the feature column with its corresponding numeric label.

## Example:

String Value	Numeric Value
Freshman	1
Sophomore	2
Junior	3
Senior	4

ID	High School Grade Level	Age
1	Freshman	14
2	Senior	17
3	Junior	16
4	Freshman	15
5	Sophomore	16



ID	High School Grade Level	Age
1	1	14
2	4	17
3	3	16
4	1	15
5	2	16



# Label Encoding for Ordinal Features

Convert categorical data into numerical labels (*aka “String Indexing”*)

## Drawbacks:

1. Arbitrary Assignments Can Mislead Models.
2. Misinterpretation of Relationships.
3. Model Sensitivity.
4. Loss of Information.
5. Not Suitable for all algorithms.

## Example:

String Value	Numeric Value
Freshman	1
Sophomore	2
Junior	3
Senior	4

Less Important  
More Important

ID	High School Grade Level	Age
1	Freshman	14
2	Senior	17
3	Junior	16
4	Freshman	15
5	Sophomore	16



ID	High School Grade Level	Age
1	1	14
2	4	17
3	3	16
4	1	15
5	2	16



# Label Encoding for Ordinal Features

Convert categorical data into numerical labels (*aka “String Indexing”*)

## Drawbacks:

1. Arbitrary Assignments Can Mislead Models.
2. Misinterpretation of Relationships.
3. Model Sensitivity.
4. Loss of Information.
5. Not Suitable for all algorithms.

## Example:

String Value	Numeric Value
Freshman	Equal Importance 1
Sophomore	Equal Importance 2
Junior	Equal Importance 3
Senior	4

ID	High School Grade Level	Age
1	Freshman	14
2	Senior	17
3	Junior	16
4	Freshman	15
5	Sophomore	16



ID	High School Grade Level	Age
1	1	14
2	4	17
3	3	16
4	1	15
5	2	16



# Target Encoding

Encoding categorical variables: replacement by mean of target variable

## Procedure:

1. Compute the mean of the target variable for each category.
2. Substitute category instances with their mean values.

## Example:

City	Clicked Values	Mean
New York	[1,0]	0.5
Paris	[0,1]	0.5
London	[1]	1

ID	City	Clicked
1	New York	1
2	Paris	0
3	London	1
4	Paris	1
5	New York	0



ID	City	Clicked
1	0.5	1
2	0.5	0
3	1.0	1
4	0.5	1
5	0.5	0



# Target Encoding

Encoding categorical variables: replacement by mean of target variable

## Drawbacks:

1. Overfitting risk
2. Handling new categories
3. Dependency of target distribution
4. Limited applicability to small datasets
5. Risk of bias

## Example:

City	Clicked Values	Mean
New York	[1,0]	0.5
Paris	[0,1]	0.5
London	[1]	1

Can lead to overconfident prediction

ID	City	Clicked
1	New York	1
2	Paris	0
3	London	1
4	Paris	1
5	New York	0



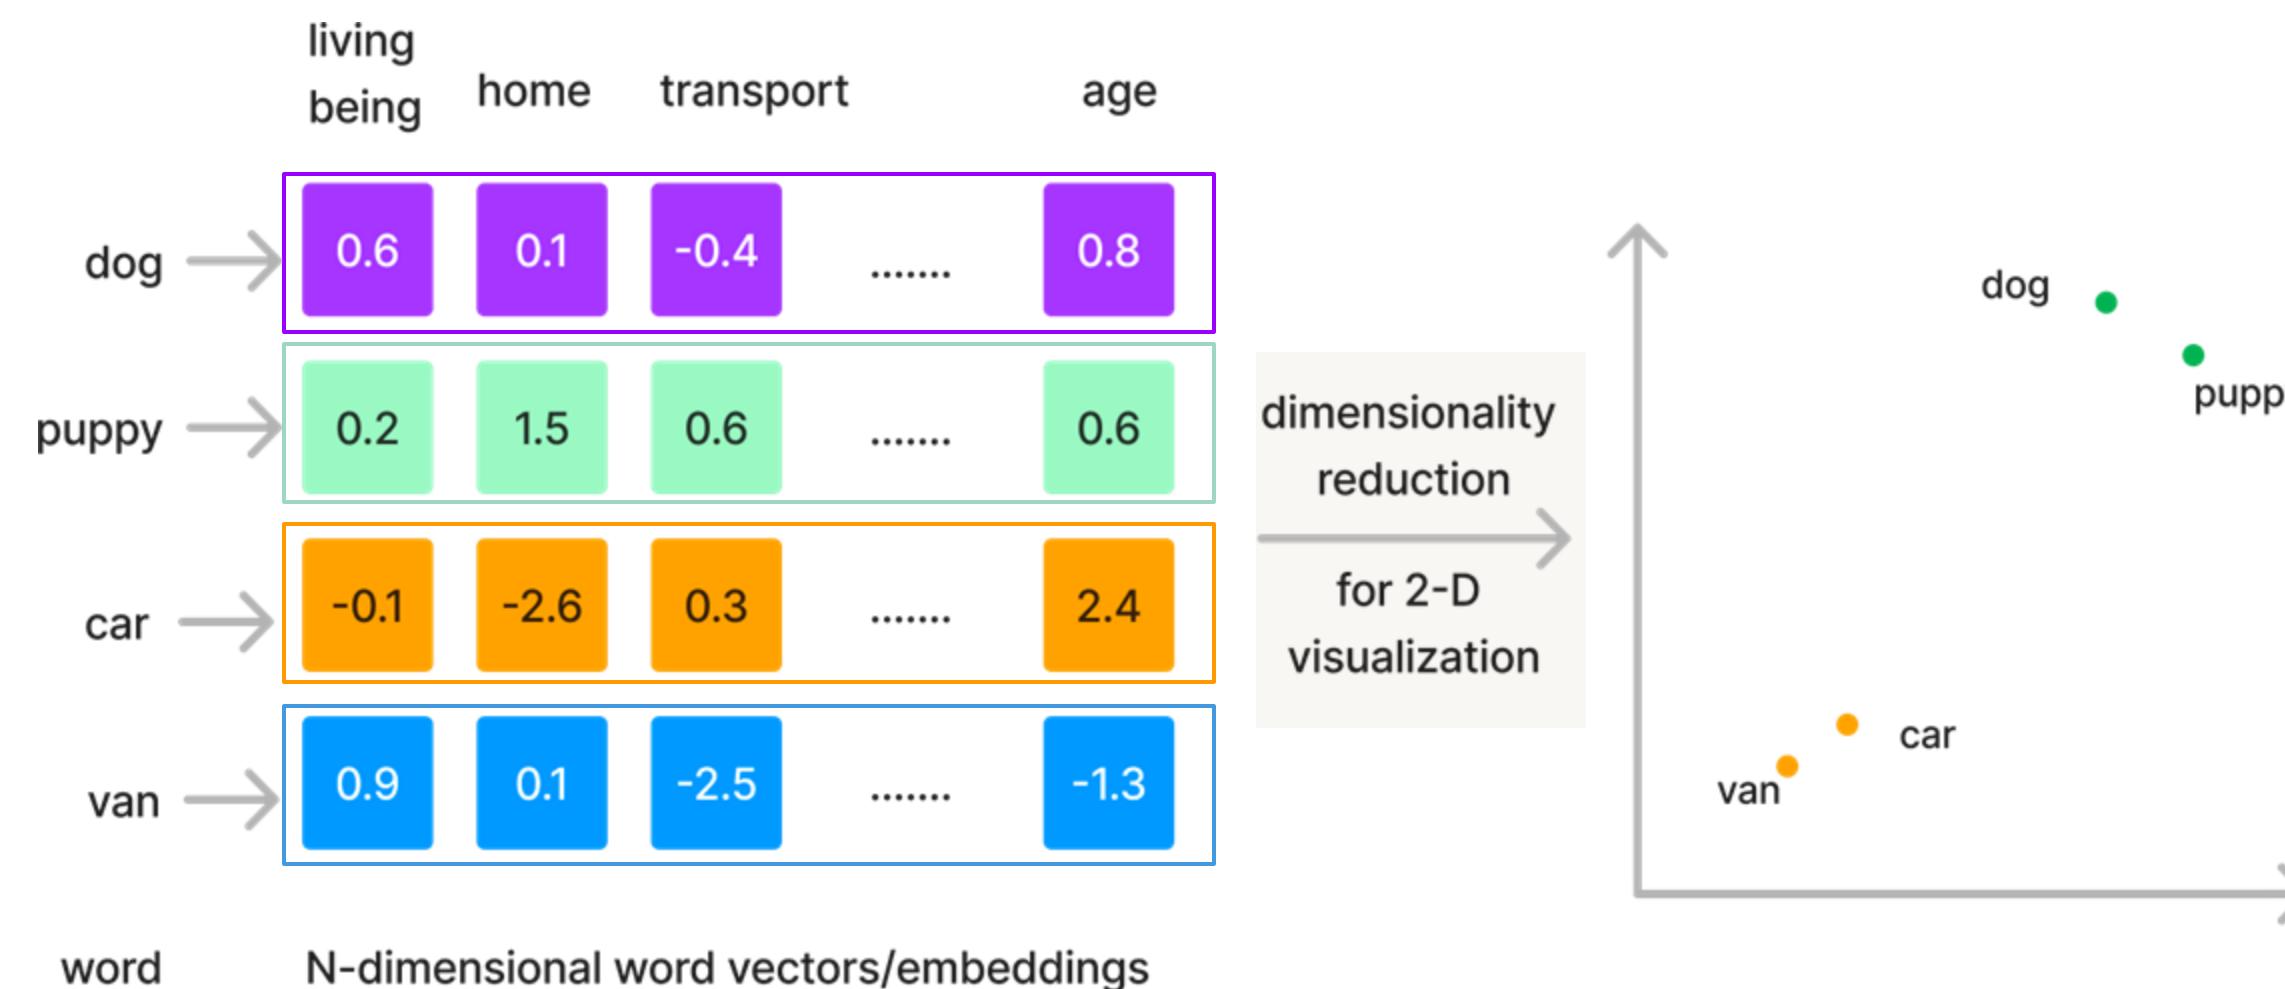
ID	City	Clicked
1	0.5	1
2	0.5	0
3	1.0	1
4	0.5	1
5	0.5	0



# What are Embeddings?

## Embeddings: Representing Data in a Meaningful Way

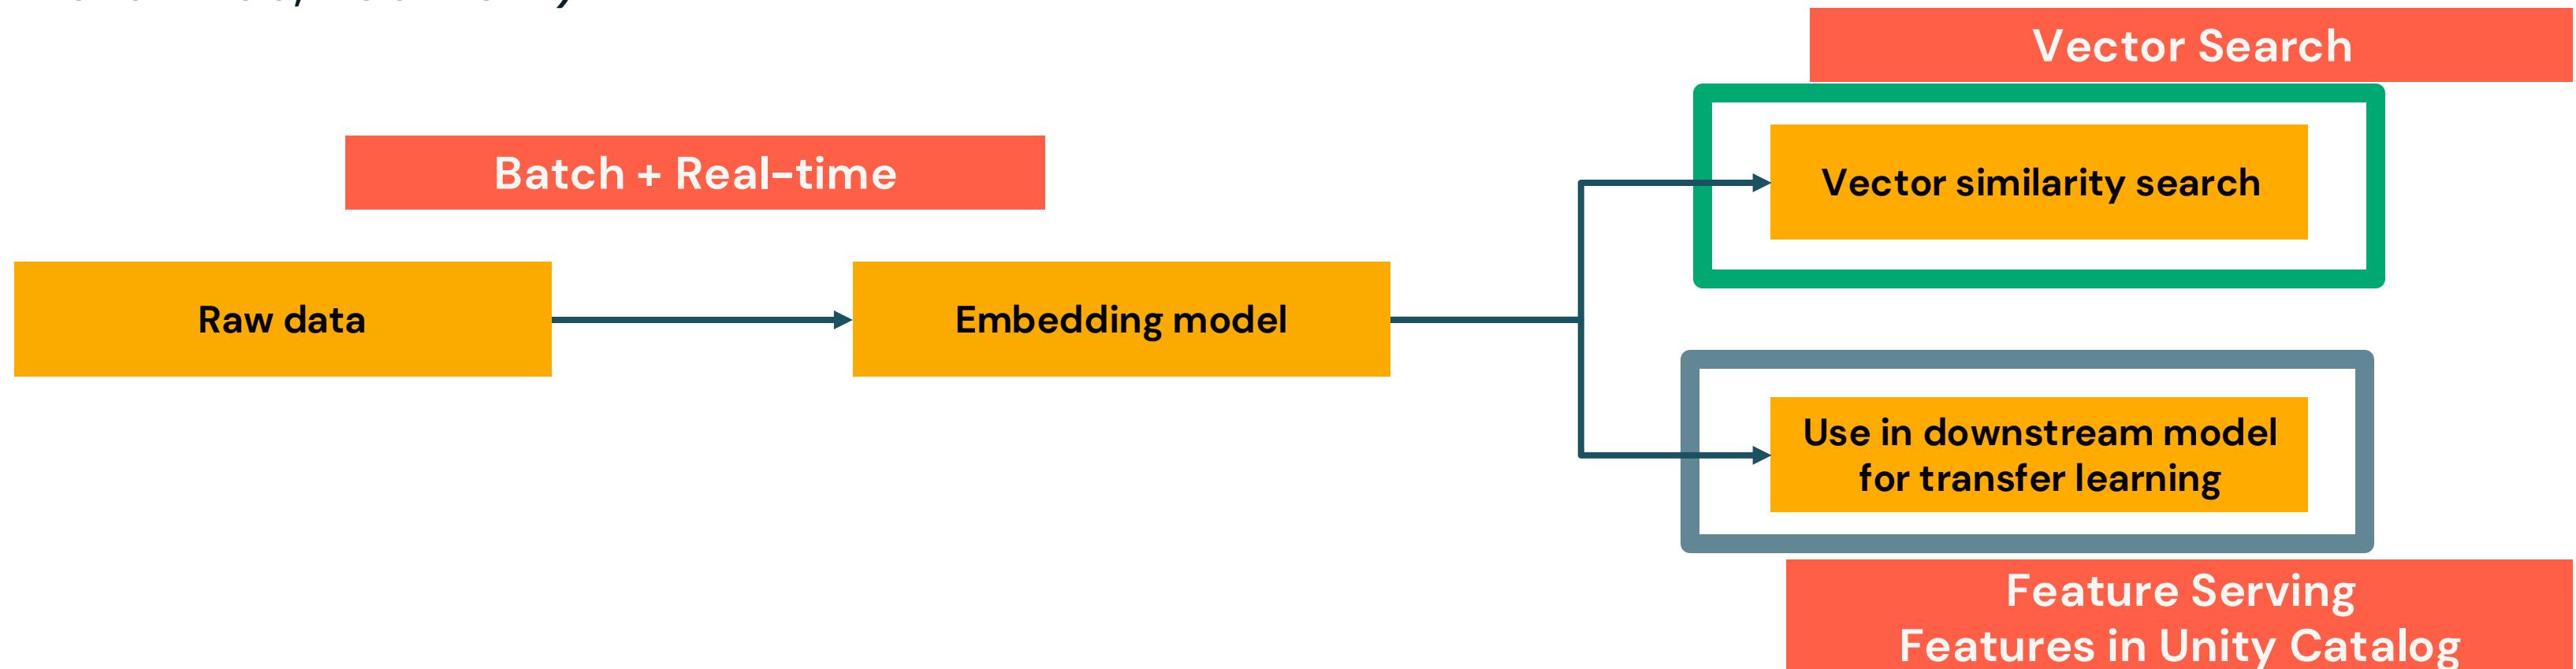
- Embeddings are **dense vector representations** of data, capturing semantic relationships.
- Unlike traditional encoding methods, embeddings **place similar items closer together** in a lower-dimensional space.
- Used widely in **text, categorical data, images, and audio processing**.



# How Embeddings Work?

## The Role of Embeddings in Machine Learning

- Embeddings transform high-dimensional categorical or textual data into a **compact, dense vector space**.
- These representations capture relationships and context among different entities.
- Used in **Recommendation Systems, NLP, Image Search, and more**.
- Can be learned from data using neural networks or retrieved from **pretrained models** (e.g., Word2Vec, FastText).



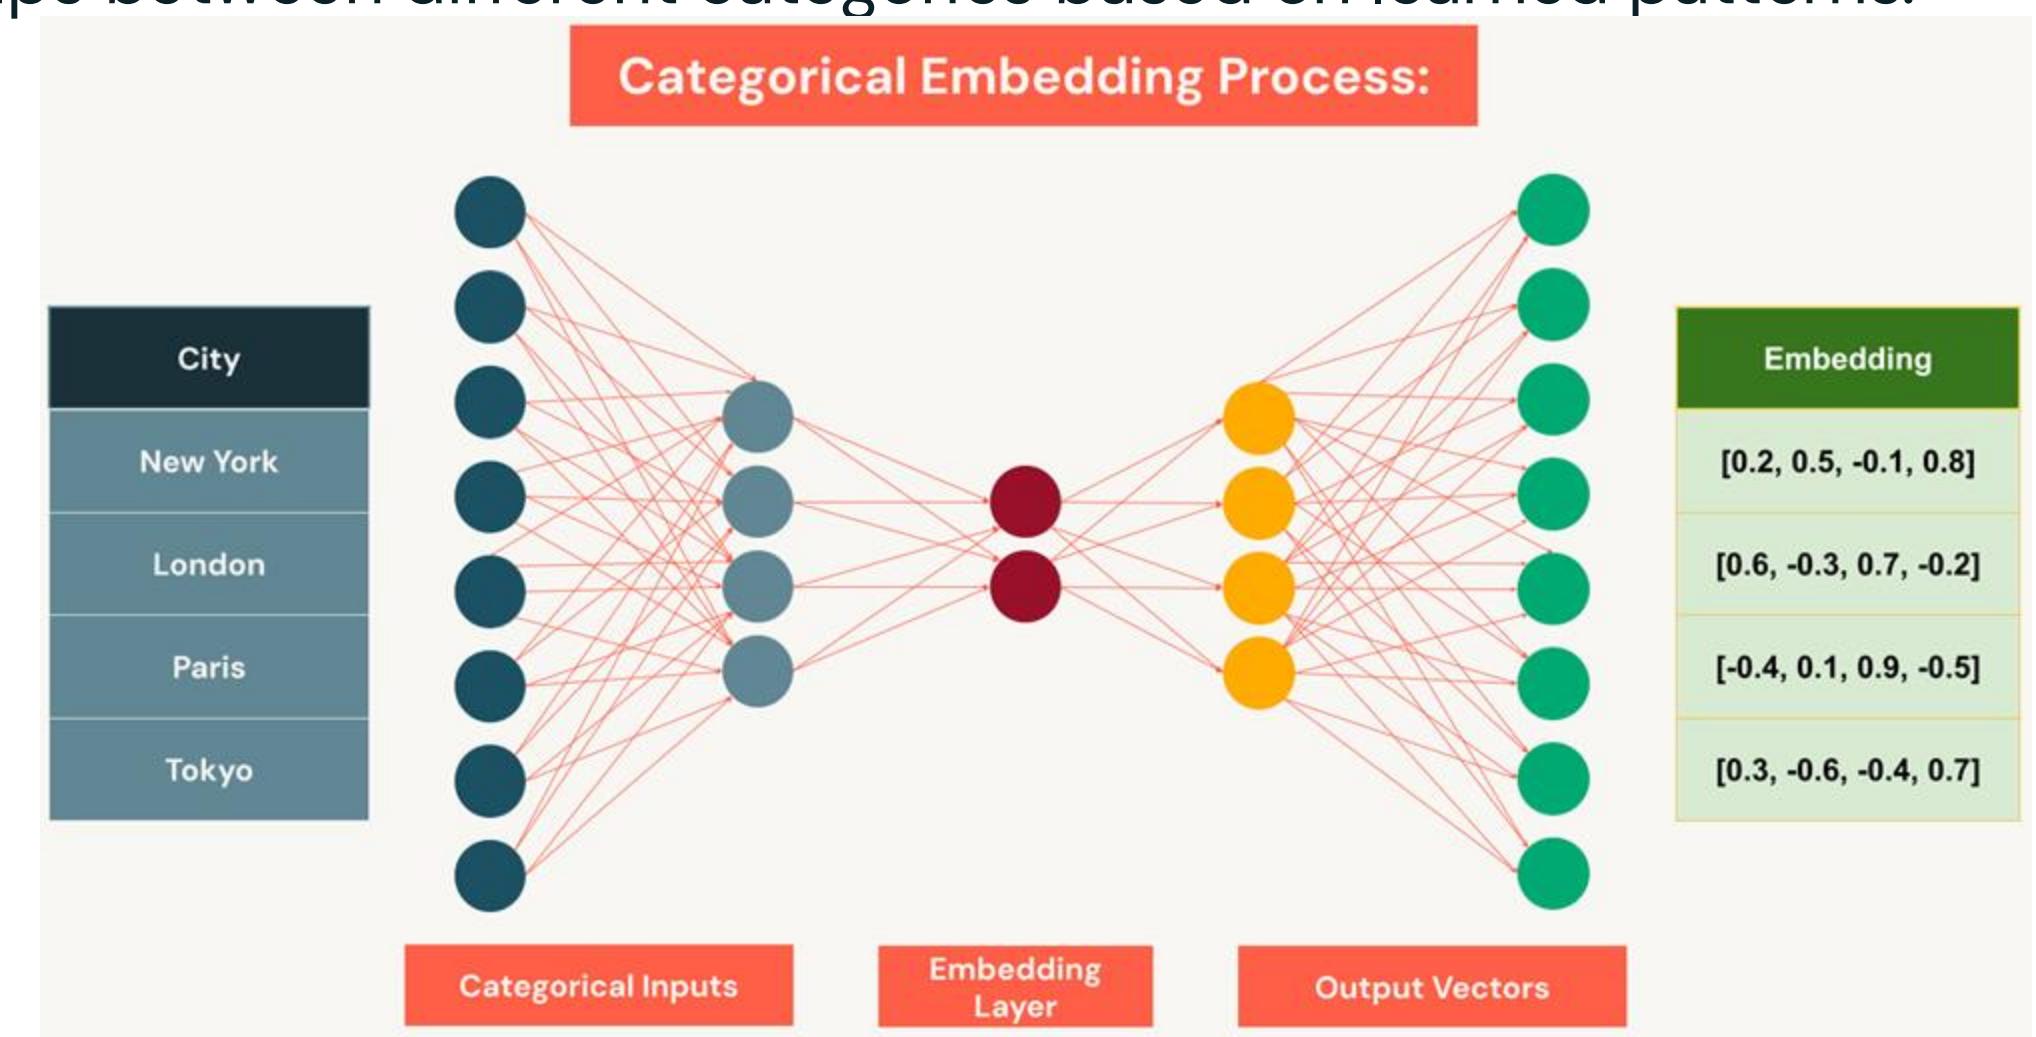
# Categorical Embedding

Transforming categorical variables into meaningful, continuous vectors

- Categorical embeddings replace sparse categorical encoding methods with lower-dimensional, dense vector representations.
- These embeddings capture relationships between different categories based on learned patterns.

## Process:

1. Convert categorical variables into unique indices.
2. Pass through an embedding layer (neural network or learned lookup table).
3. The model assigns dense vector representations to each category based on learned relationships.



# Embeddings vs. One-Hot Encoding

## Choosing the Right Encoding for Categorical Data

Feature	One-Hot Encoding	Embeddings
Representation	Sparse binary vectors (high-dimensional)	Dense numerical vectors (low-dimensional)
Semantic Relationships	Does not capture relationships between categories	Places similar categories closer in vector space
Scalability	Inefficient for high-cardinality data	Scales well with large category sets
Efficiency	Inefficient for high-cardinality features (sparse matrix)	Efficient for high-cardinality features
Model Suitability	Suitable for simple models like Decision Trees	Best for Neural Networks and deep learning models
Example:	TV → [1, 0, 0, 0] Laptop → [0, 1, 0, 0] Phone → [0, 0, 1, 0]	TV → [0.6, 1.2, -0.8] Laptop → [0.5, 1.1, -0.7] Phone → [0.2, -0.4, 1.5]





Data Preparation and Feature Engineering

LECTURE

# Data Standardization

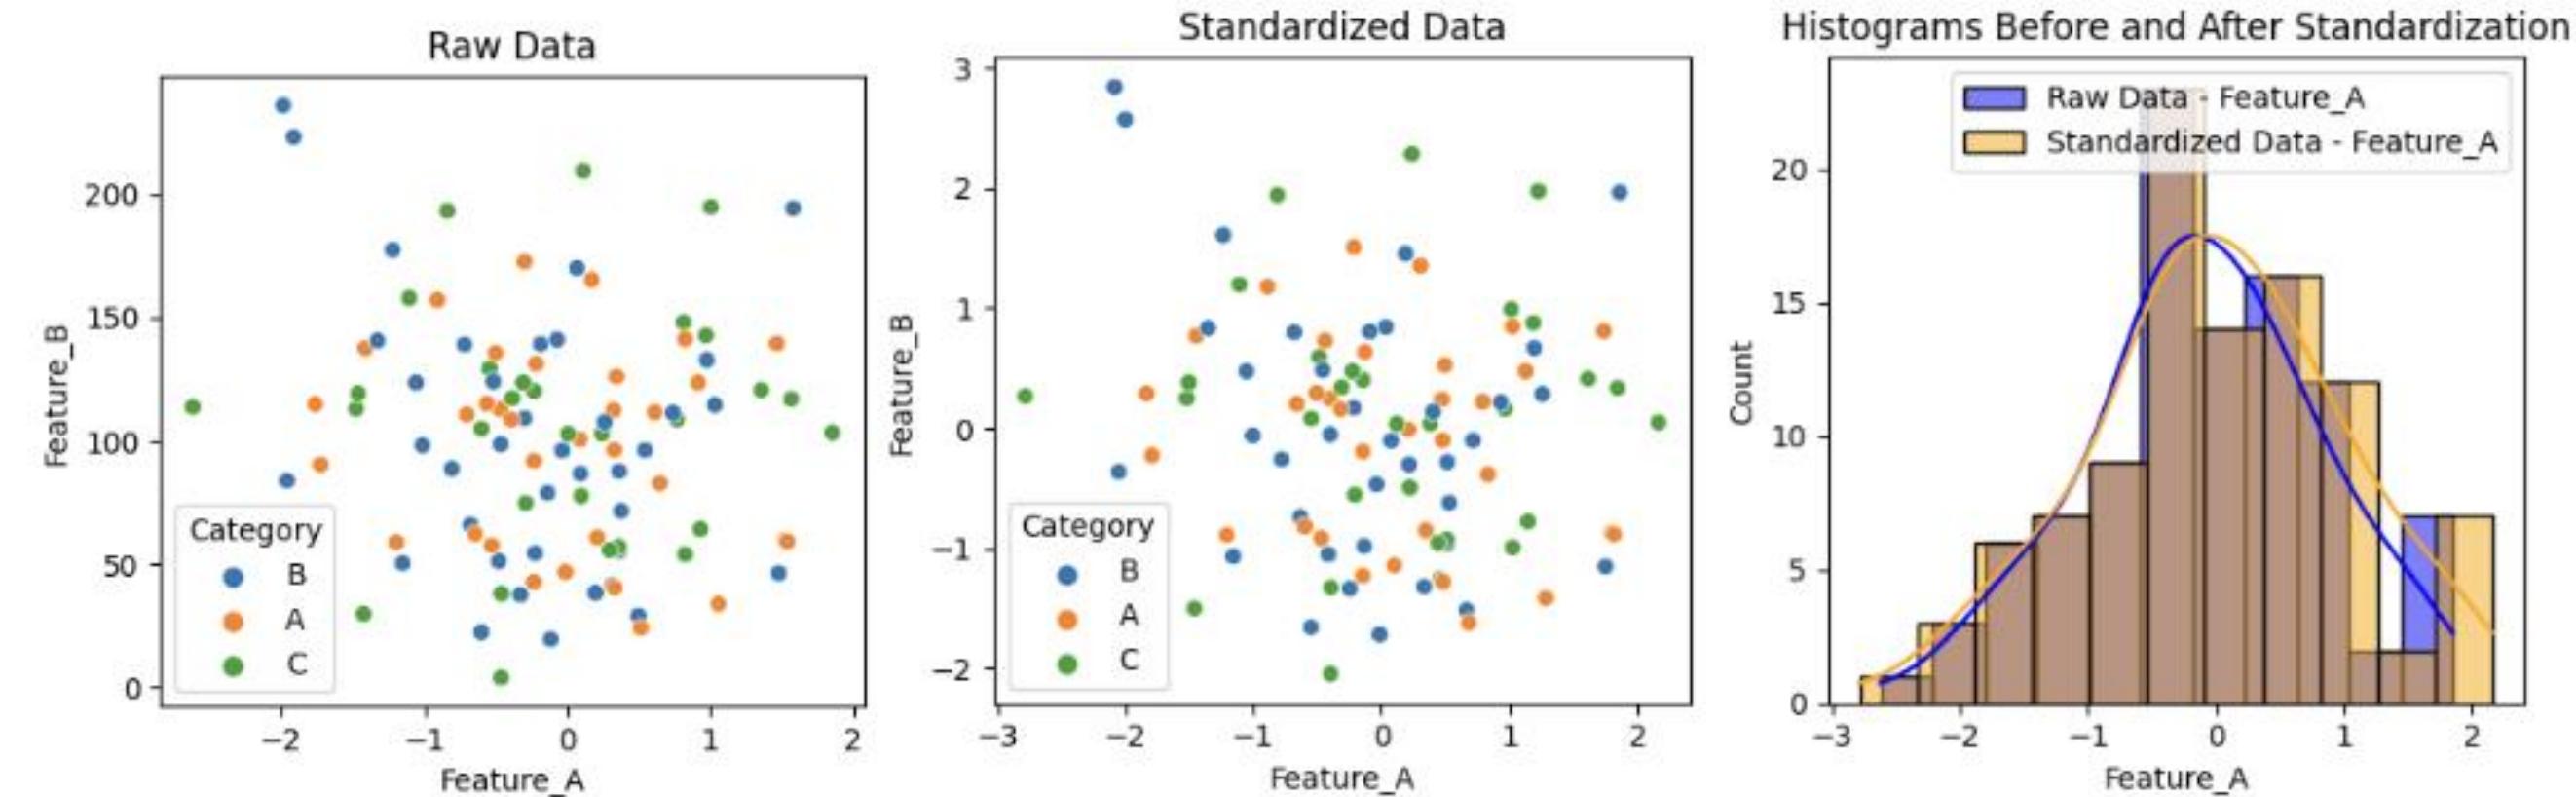


© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Data Standardization

Ensuring data consistency and optimal model performance.

Data standardization is the process of creating standards and transforming data taken from different sources into a consistent format that adheres to the standards.



# Benefits of Data Standardization

Standardizing features offers significant advantages



**Improved Interpretability**



**Prevention of Model Biases**



**Enhanced Convergence**



**Effective Regularization**



**Facilitates Comparison of Coefficients**



# Benefits of Data Standardization

Standardizing features offers significant advantages



Improved Interpretability



Prevention of Model Biases



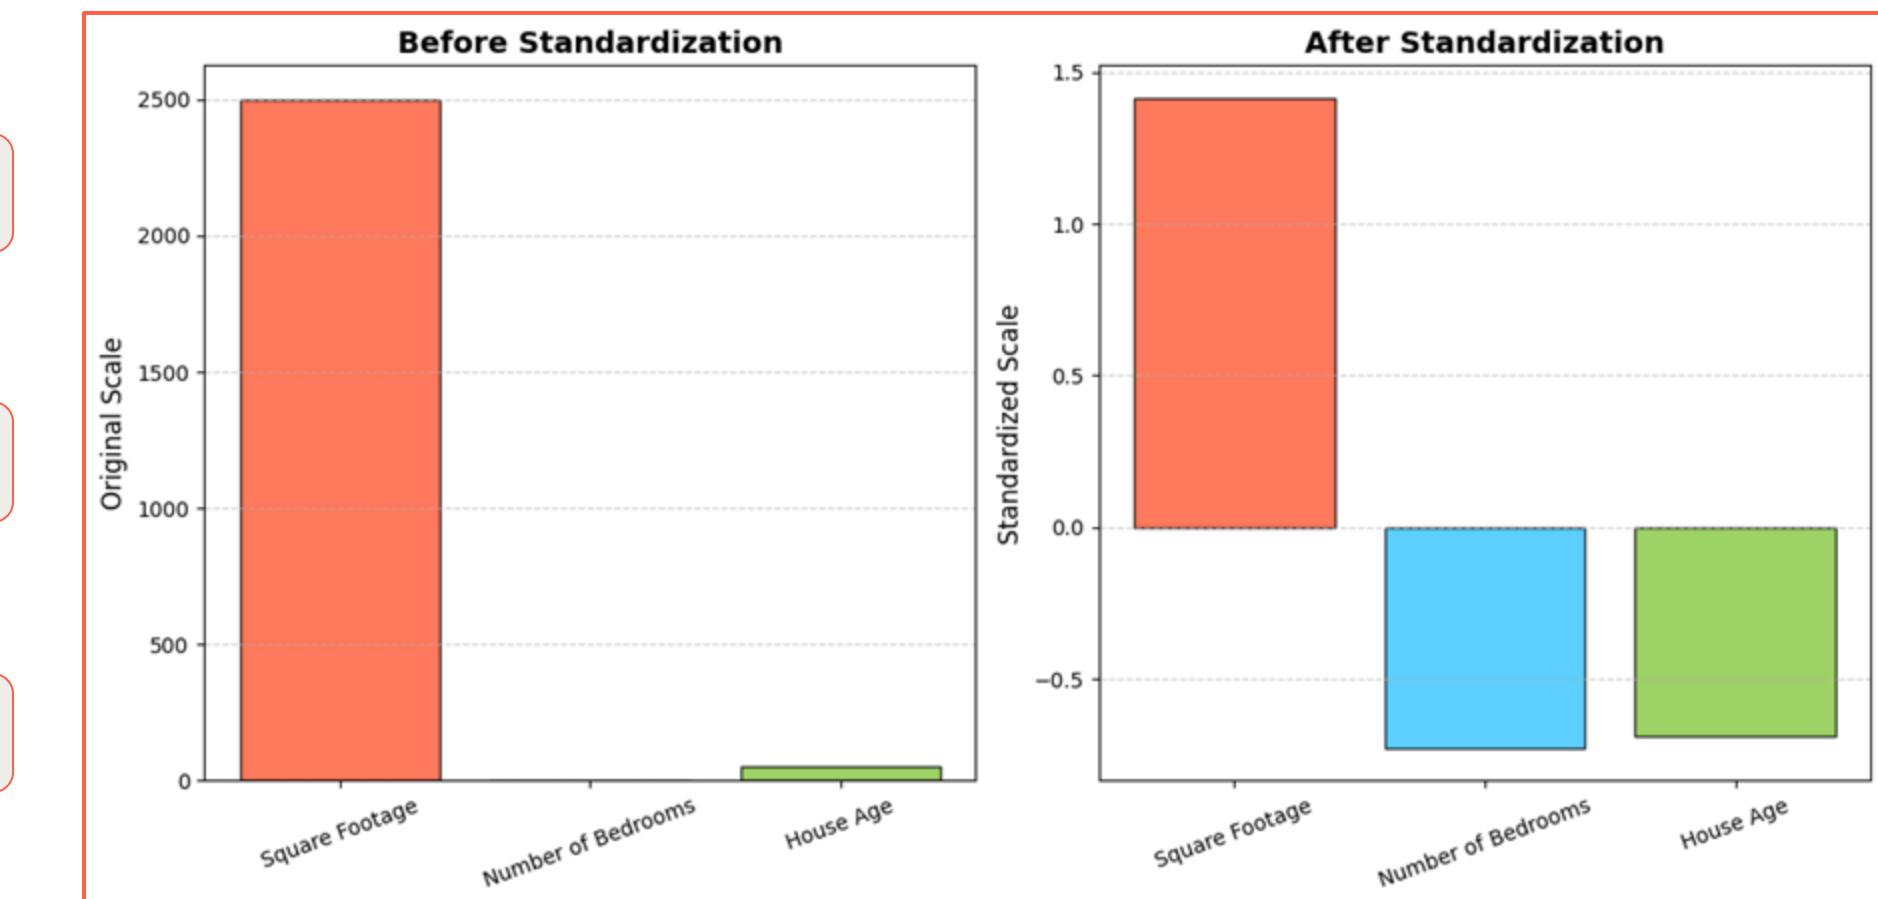
Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



# Benefits of Data Standardization

Standardizing features offers significant advantages



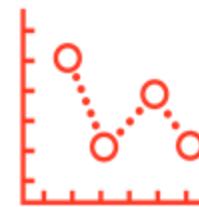
Improved Interpretability



Prevention of Model Biases



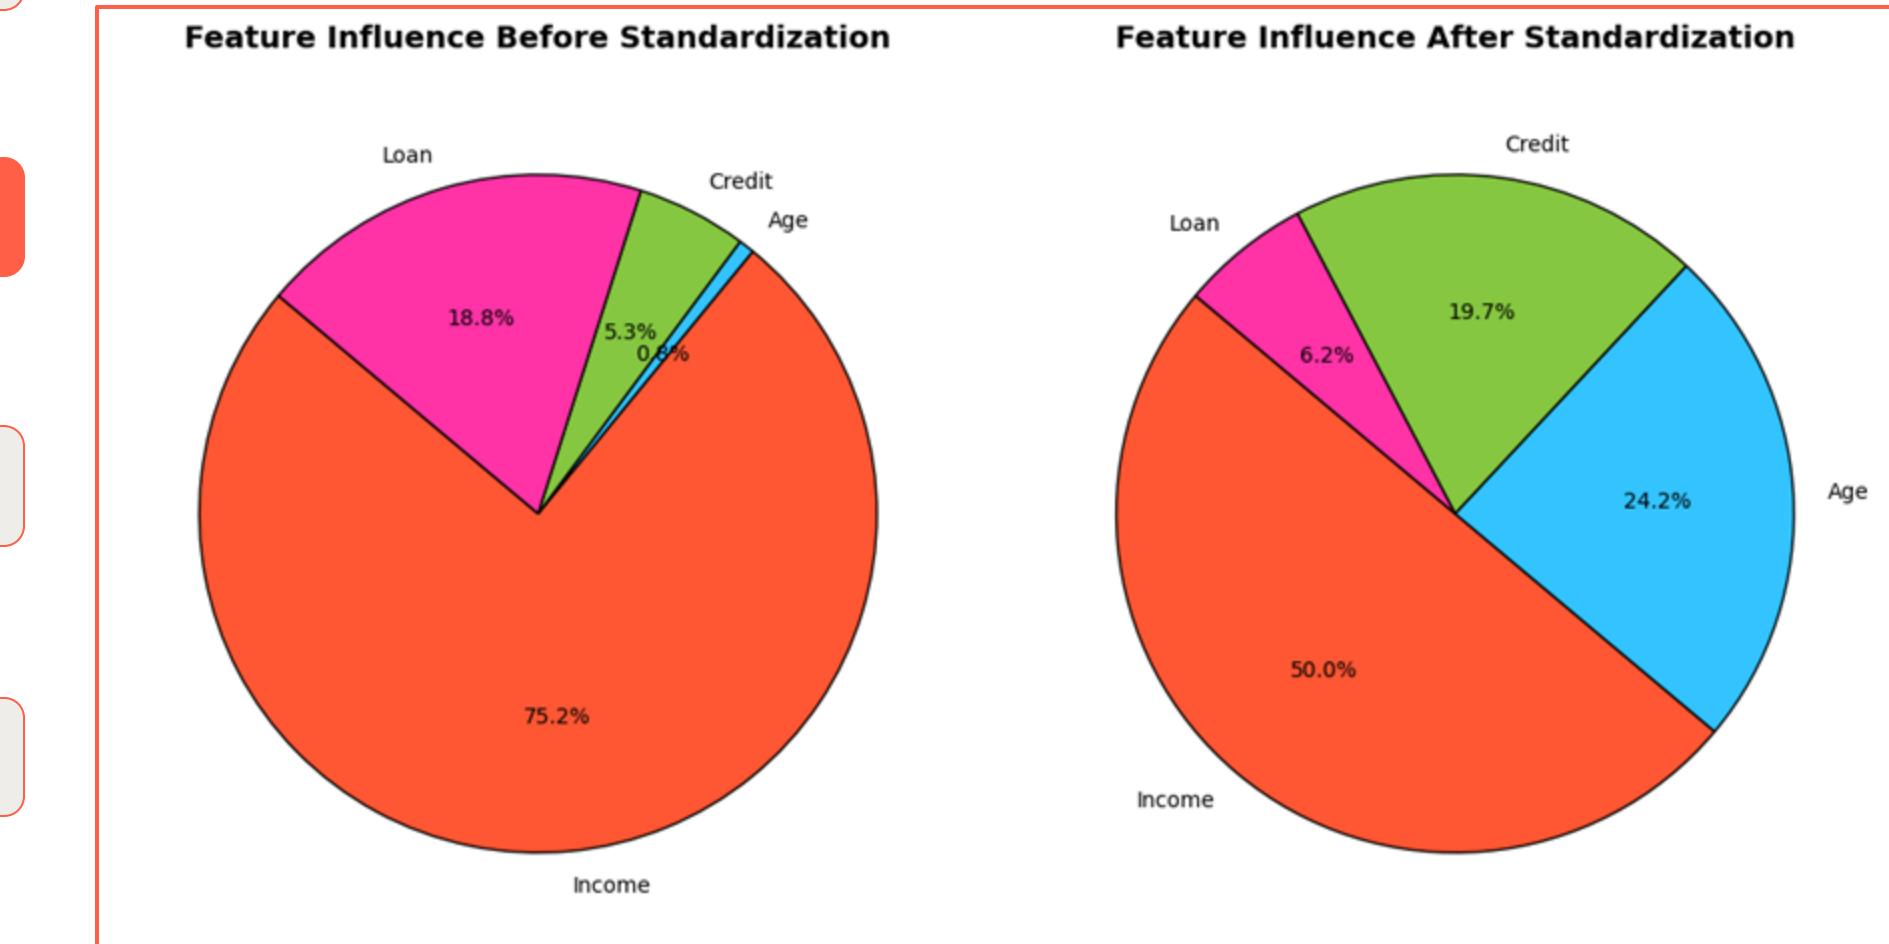
Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



# Benefits of Data Standardization

Standardizing features offers significant advantages



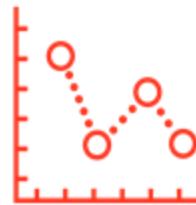
Improved Interpretability



Prevention of Model Biases



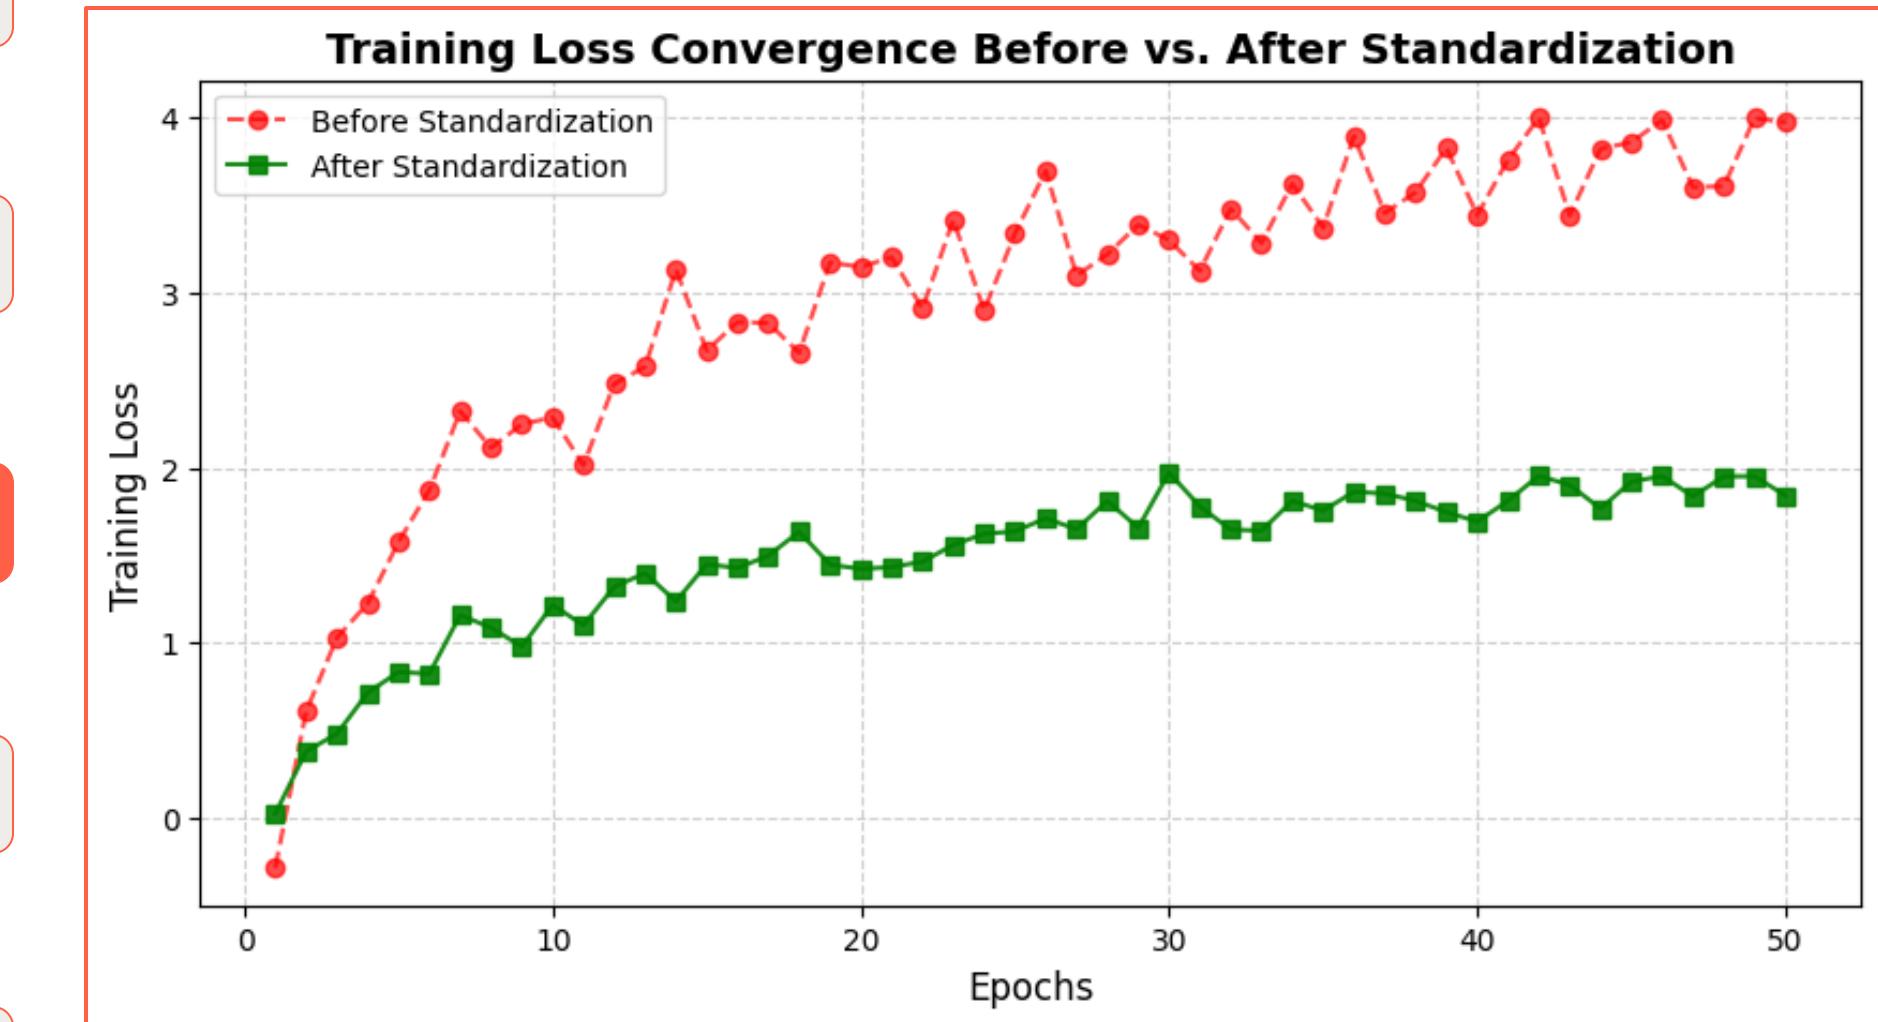
Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



# Benefits of Data Standardization

Standardizing features offers significant advantages



Improved Interpretability



Prevention of Model Biases



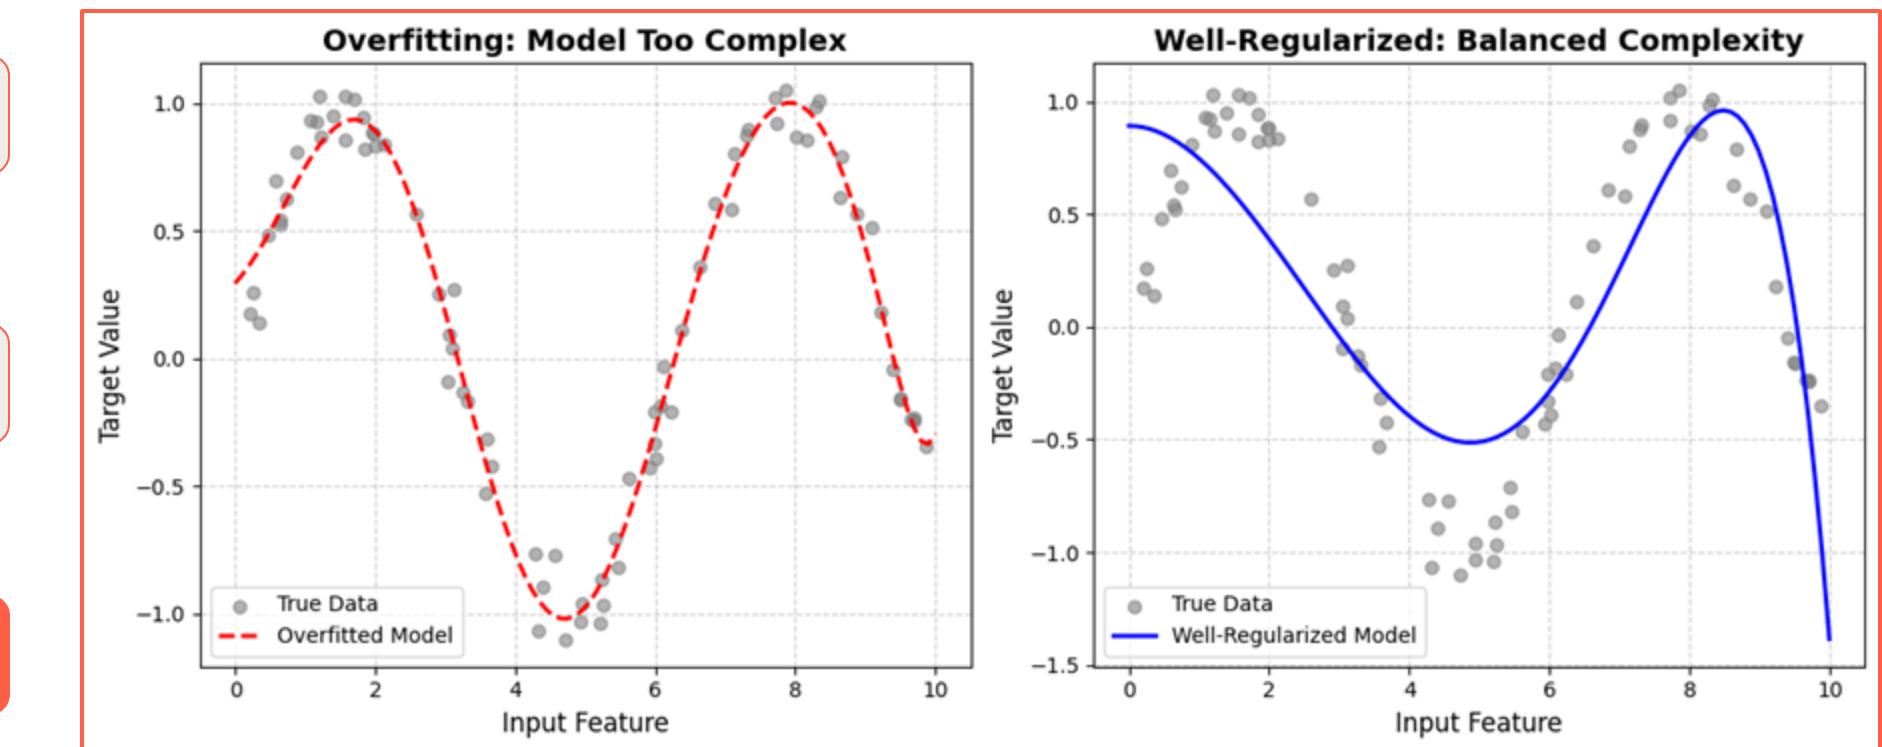
Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



# Benefits of Data Standardization

Standardizing features offers significant advantages



Improved Interpretability



Prevention of Model Biases



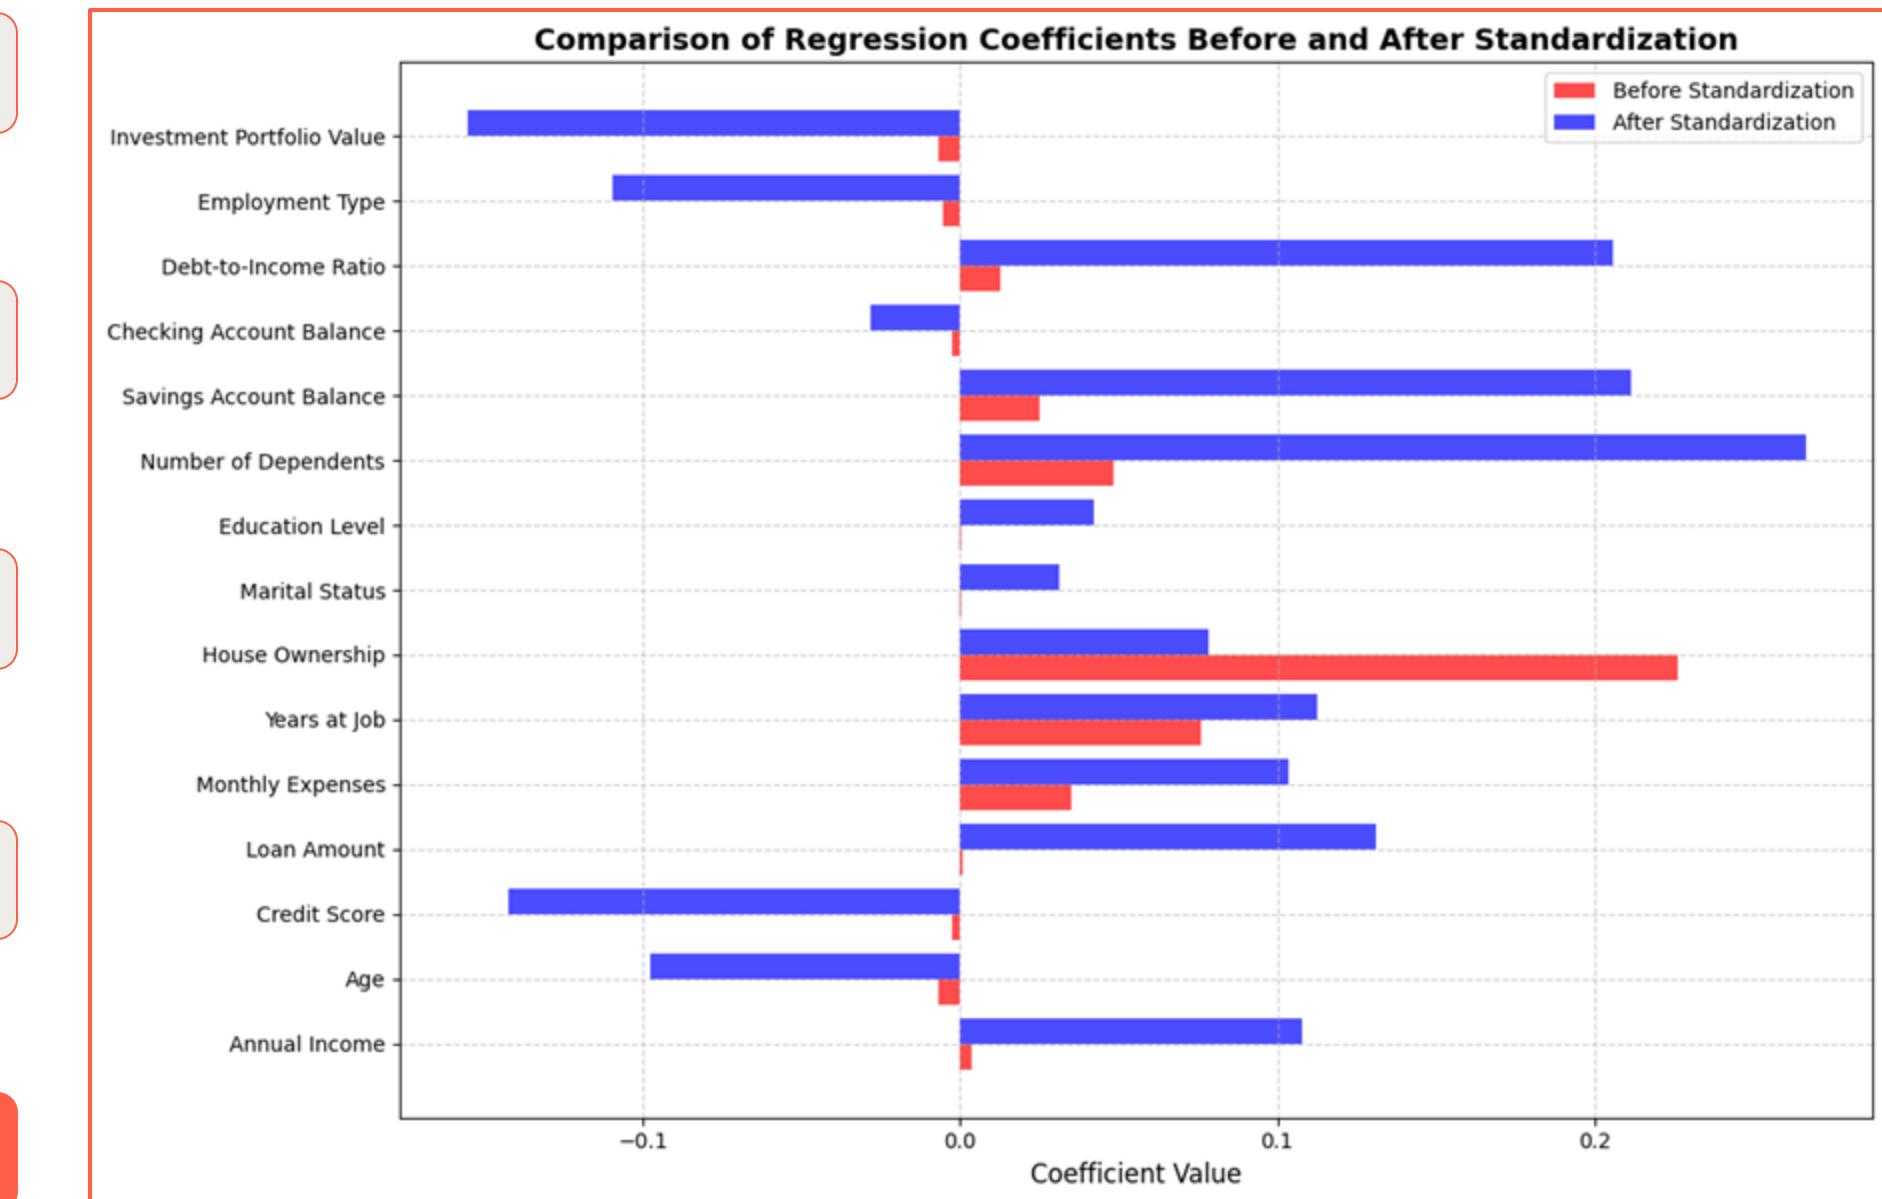
Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



# Which Models are Sensitive or Robust to the Scale of the Feature?

Sensitivity or robustness of ML models to the scale of features can vary

## Models Sensitive to Feature Scale:

- Linear Models
- Support Vector Machines (SVM)
- K-Nearest Neighbors (KNN)
- Neural Networks

## Models Robust to Feature Scale:

- Decision Trees and Random Forests
- Ensemble Models
- Tree-based Models
- Principal Component Analysis (PCA)



# Process of Data Standardization

Navigating the Steps to Standardize Data for Machine Learning

- 1 Assess feature distributions/scales and need to standardize
- 2 Prep-process + Hold-Out split (train/test then train/validation)
- 3 Create/Push-to Feature Store table(s)
- 4 Create standardization pipeline and fit to train/validation & test sets



# Drawbacks of Data Standardization

## Limitations in the Process of Data Standardization

- Loss of Original Interpretation
- Challenges in Interpreting Interaction Terms
- Challenges in Result Interpretation
- *Assumption of Normally Distributed Features\** (i.e. z-score)

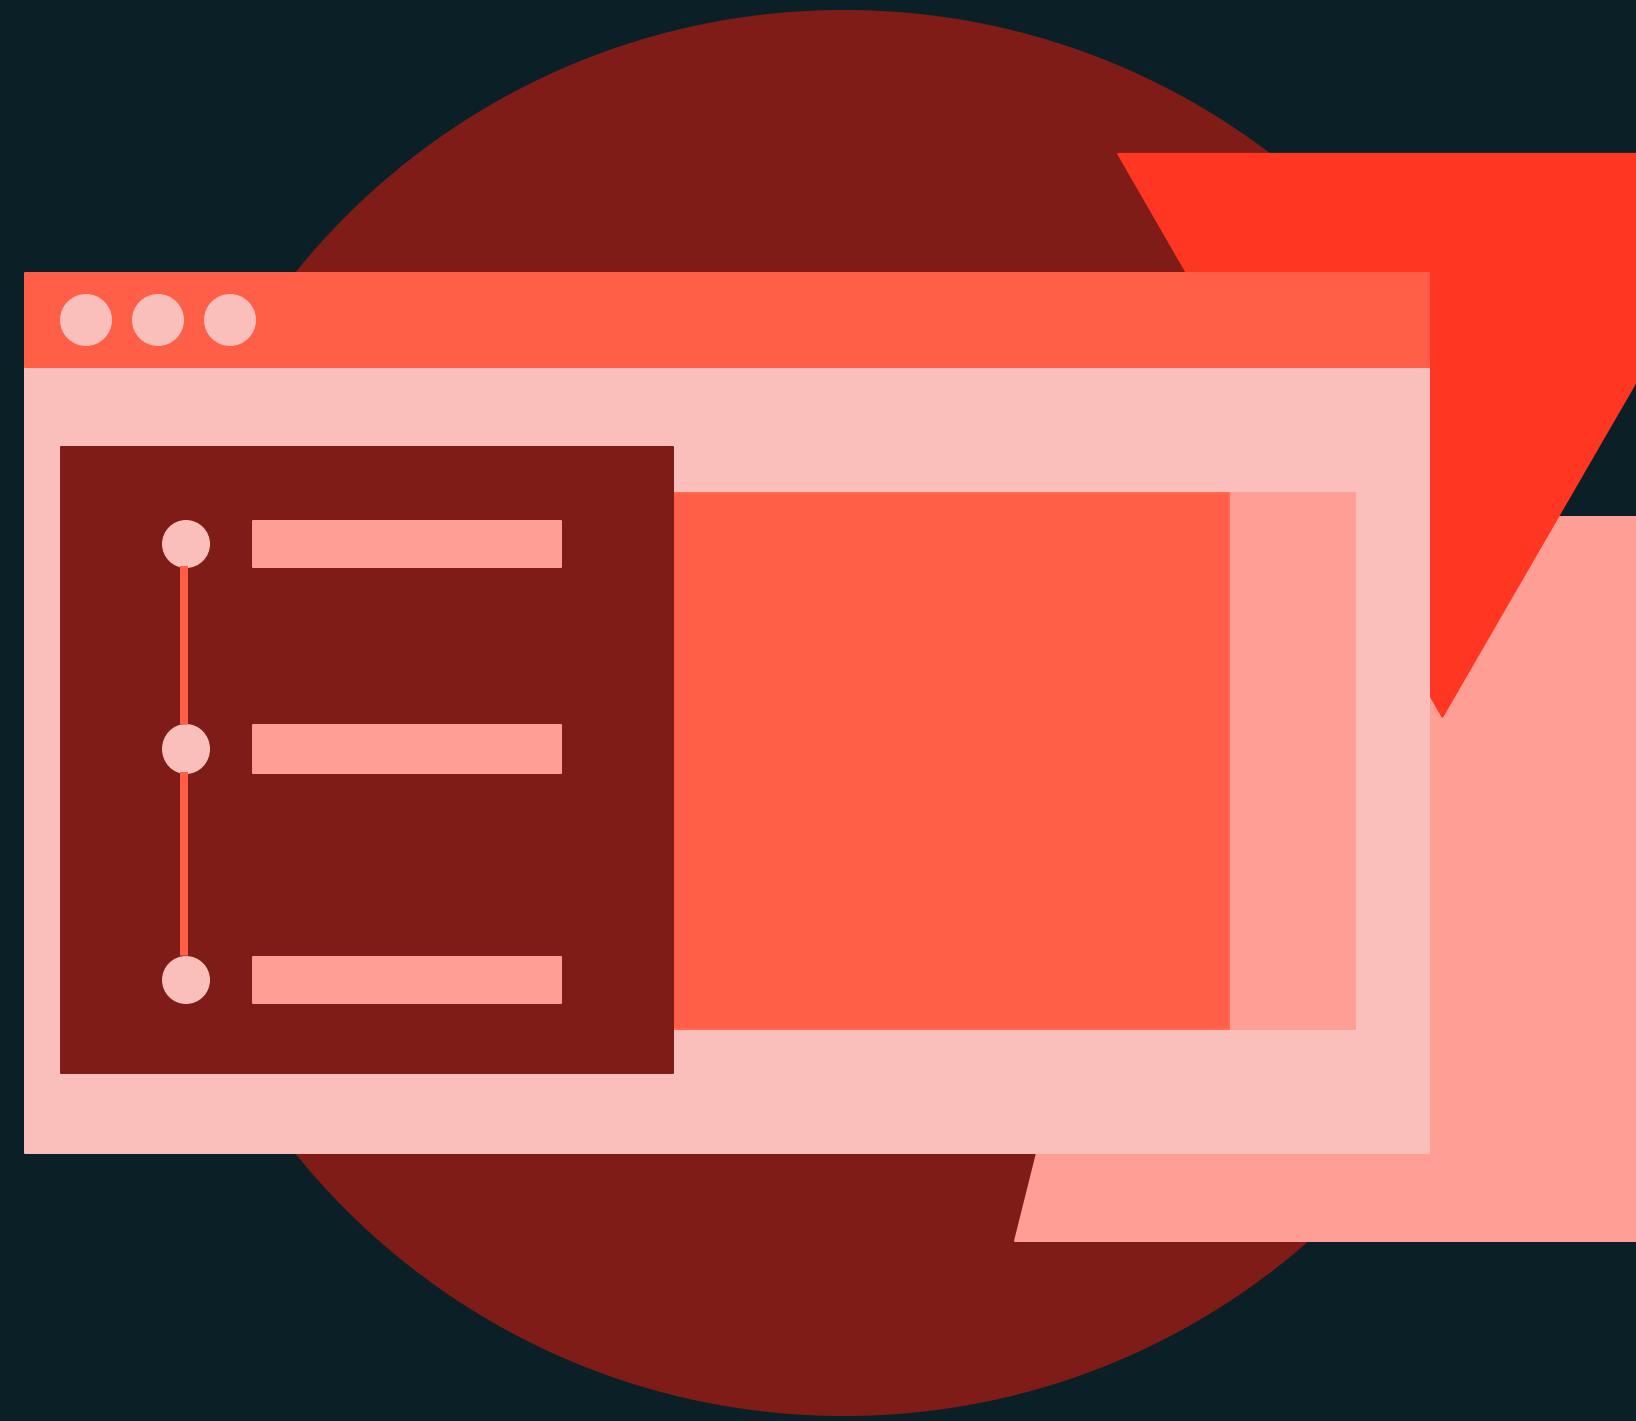




Data Preparation and Feature Engineering

**DEMONSTRATION**

# Data Imputation and Transformation Pipeline



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Demo

## Outline

### What we'll cover:

- Data cleaning and imputation
  - Coerce/fix data types
  - Handling outliers
  - Handling missing values
- Encoding categorical features
- Splitting data

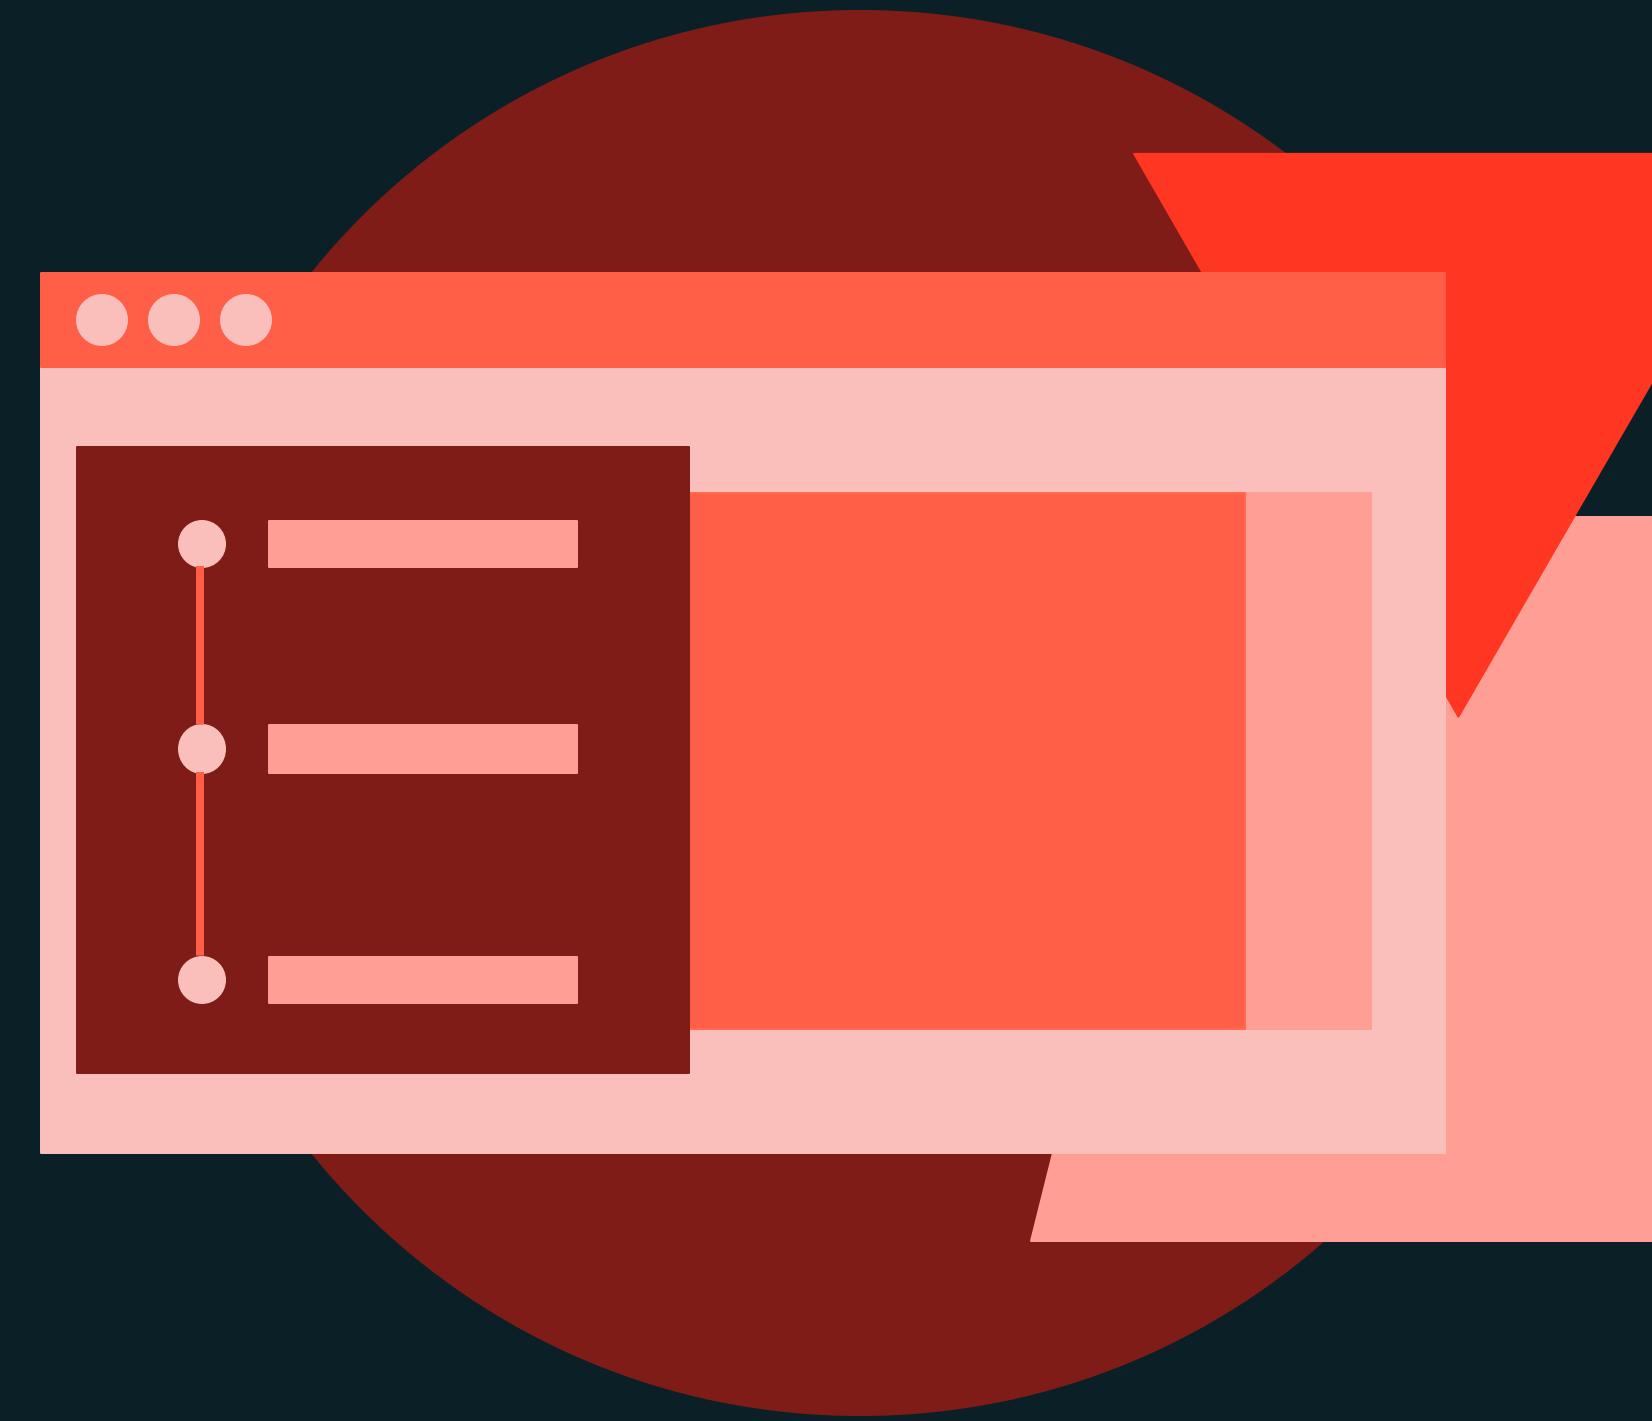




Data Preparation and Feature Engineering

**DEMONSTRATION**

# Build a Feature Engineering Pipeline with Embeddings



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Demo

## Outline

### What we'll cover:

- Build a structured feature engineering pipeline that includes multiple preprocessing steps.
- Create a pipeline with tasks for data imputation and numerical feature scaling.
- Generate embeddings for categorical features to represent categorical data effectively.
- Apply the feature engineering pipeline to both training and test datasets.
- Save a data preparation and feature engineering pipeline to Unity Catalog for potential future use.





Data Preparation and Feature Engineering

**LAB EXERCISE**

# Build a Feature Engineering Pipeline



# Lab

## Outline

### What you'll do:

- Data preparation
- Split dataset
- Create a pipeline using data imputation and transformation
- Fit the pipeline
- Show transformation results
- Save the pipeline

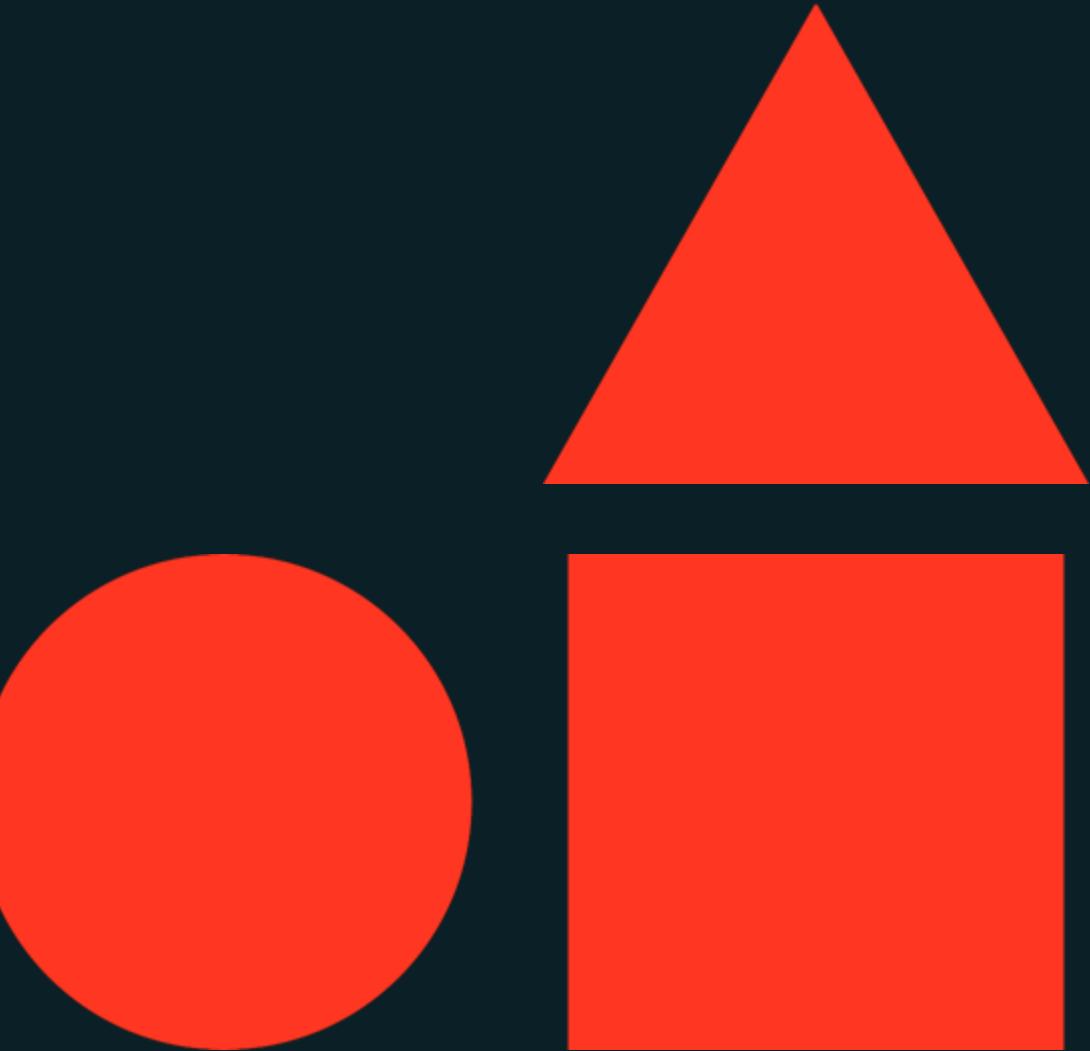




# Feature Store

---

Data Preparation for Machine Learning



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Learning objectives

Things you'll be able to do after completing this module

- Understand Feature Store concepts and its benefits/importance in machine learning (*including improved collaboration and data consistency*).
- Explain the relationship between Feature Store tables and Unity Catalog (*and compare legacy/workspace feature store and Unity Catalog feature store*)
- Using Feature Store for Feature Engineering by creating a feature store table





Feature Store

LECTURE

# Introduction to Feature Store



# What is a Feature Store?

A **feature store** manages features, or input data to a machine learning model.

In a model that predicts **customer churn**, for example, features could be:

- Aggregations of raw data over time windows, like **trailing 7-day purchases**
- Joined **combinations of data sets**, like customer demographic information joined to transaction features
- Complex functions of customer information, like **estimated customer lifetime value**

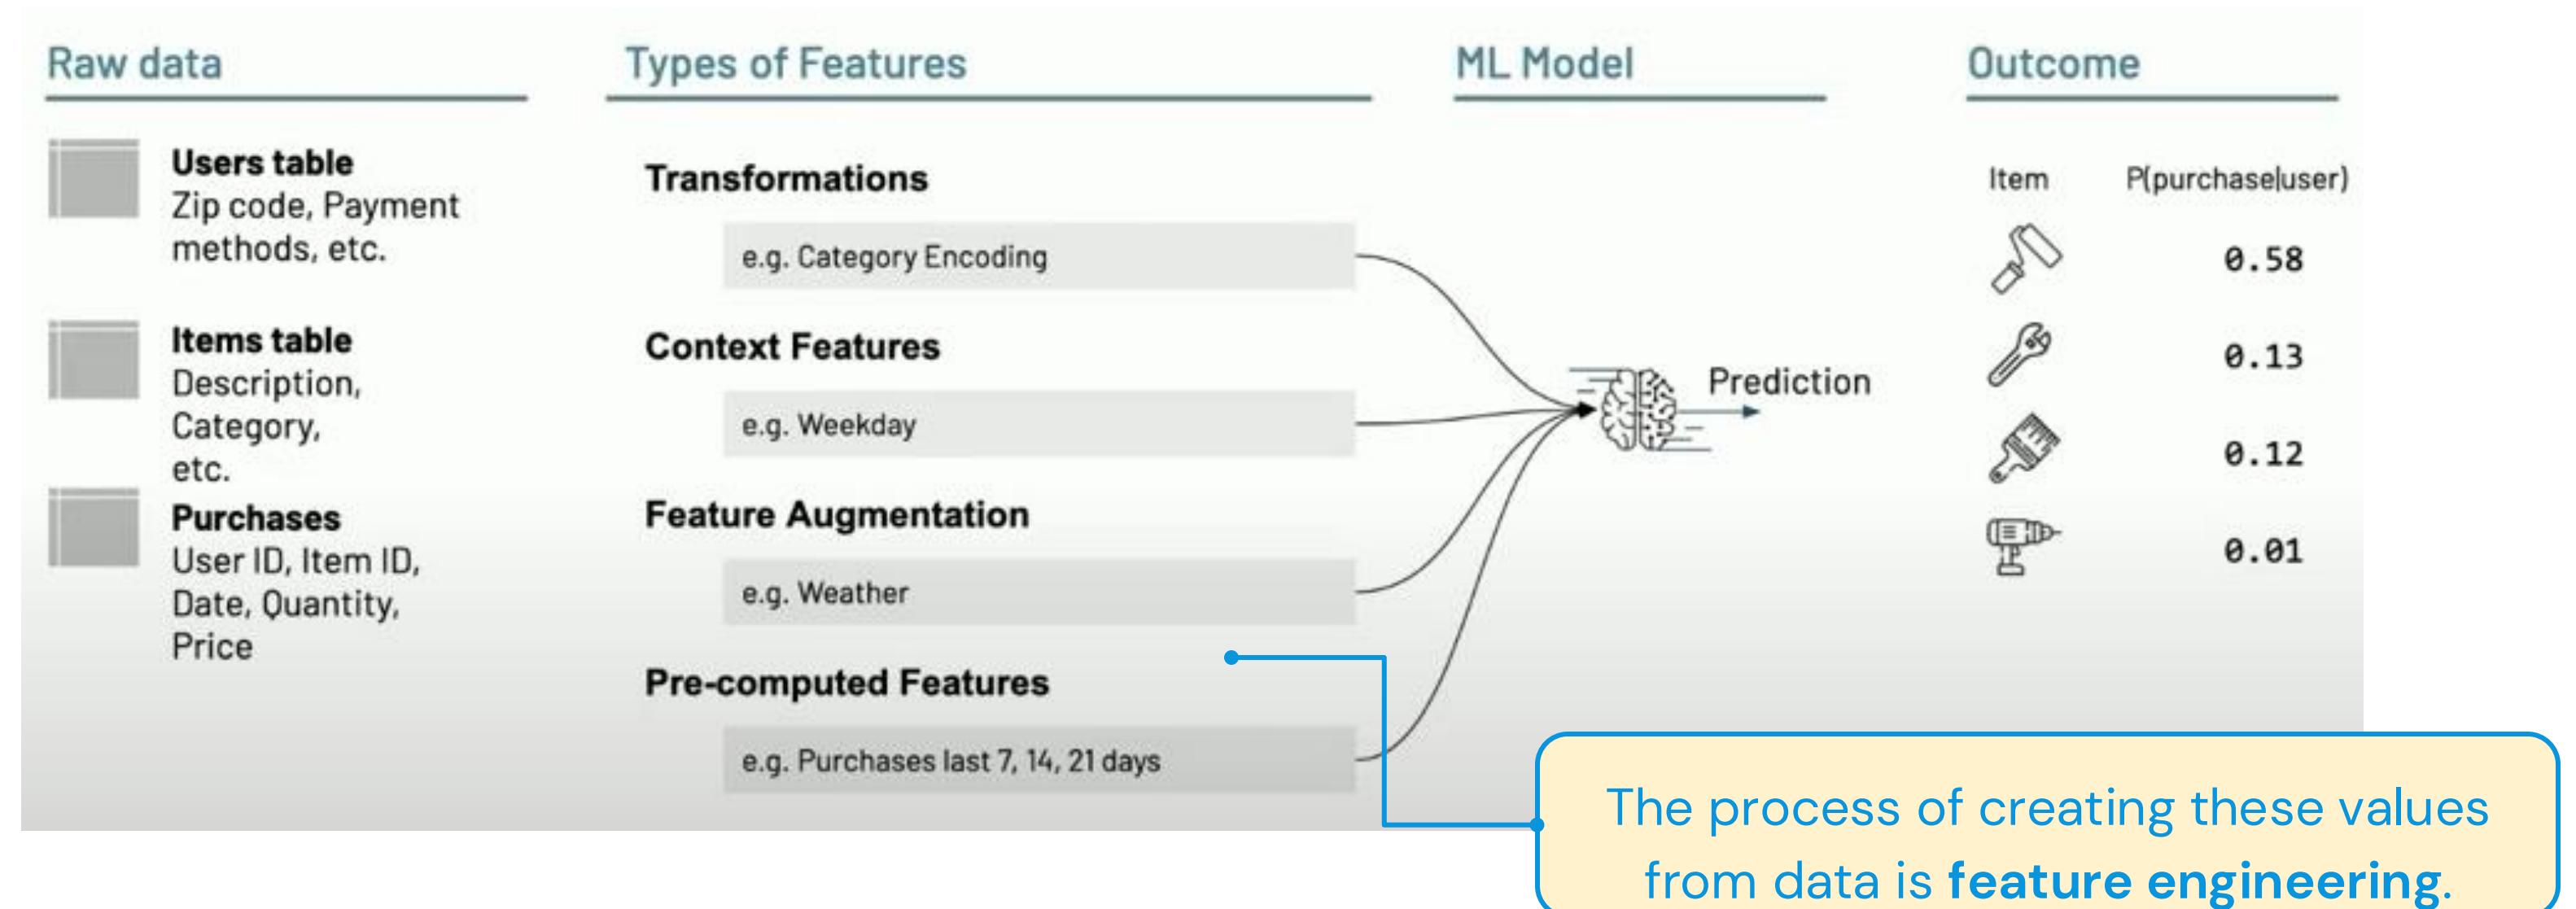
The process of creating these values from data is **feature engineering**.



# What is a Feature Store?

An example of a recommendation system.

A **feature store** manages features, or input data to a machine learning model.



# Why Would You Need a Feature Store?

## Basic Motivations

### Discovery



Multiple Data Scientists are trying to solve similar modeling tasks and come up with different definitions of the same features.  
**How can I find the features?**

### Lineage



Model governance requires documentation of the features used to train a model, as well as the **upstream lineage** of a feature to reliably use it. **How is it computed, and who owns it?**

### Skew



When multiple teams manage feature computation and ML models in production, minor yet significant **skew in upstream data** at the input of a feature pipeline can be very hard to detect and fix.

### Online Serving

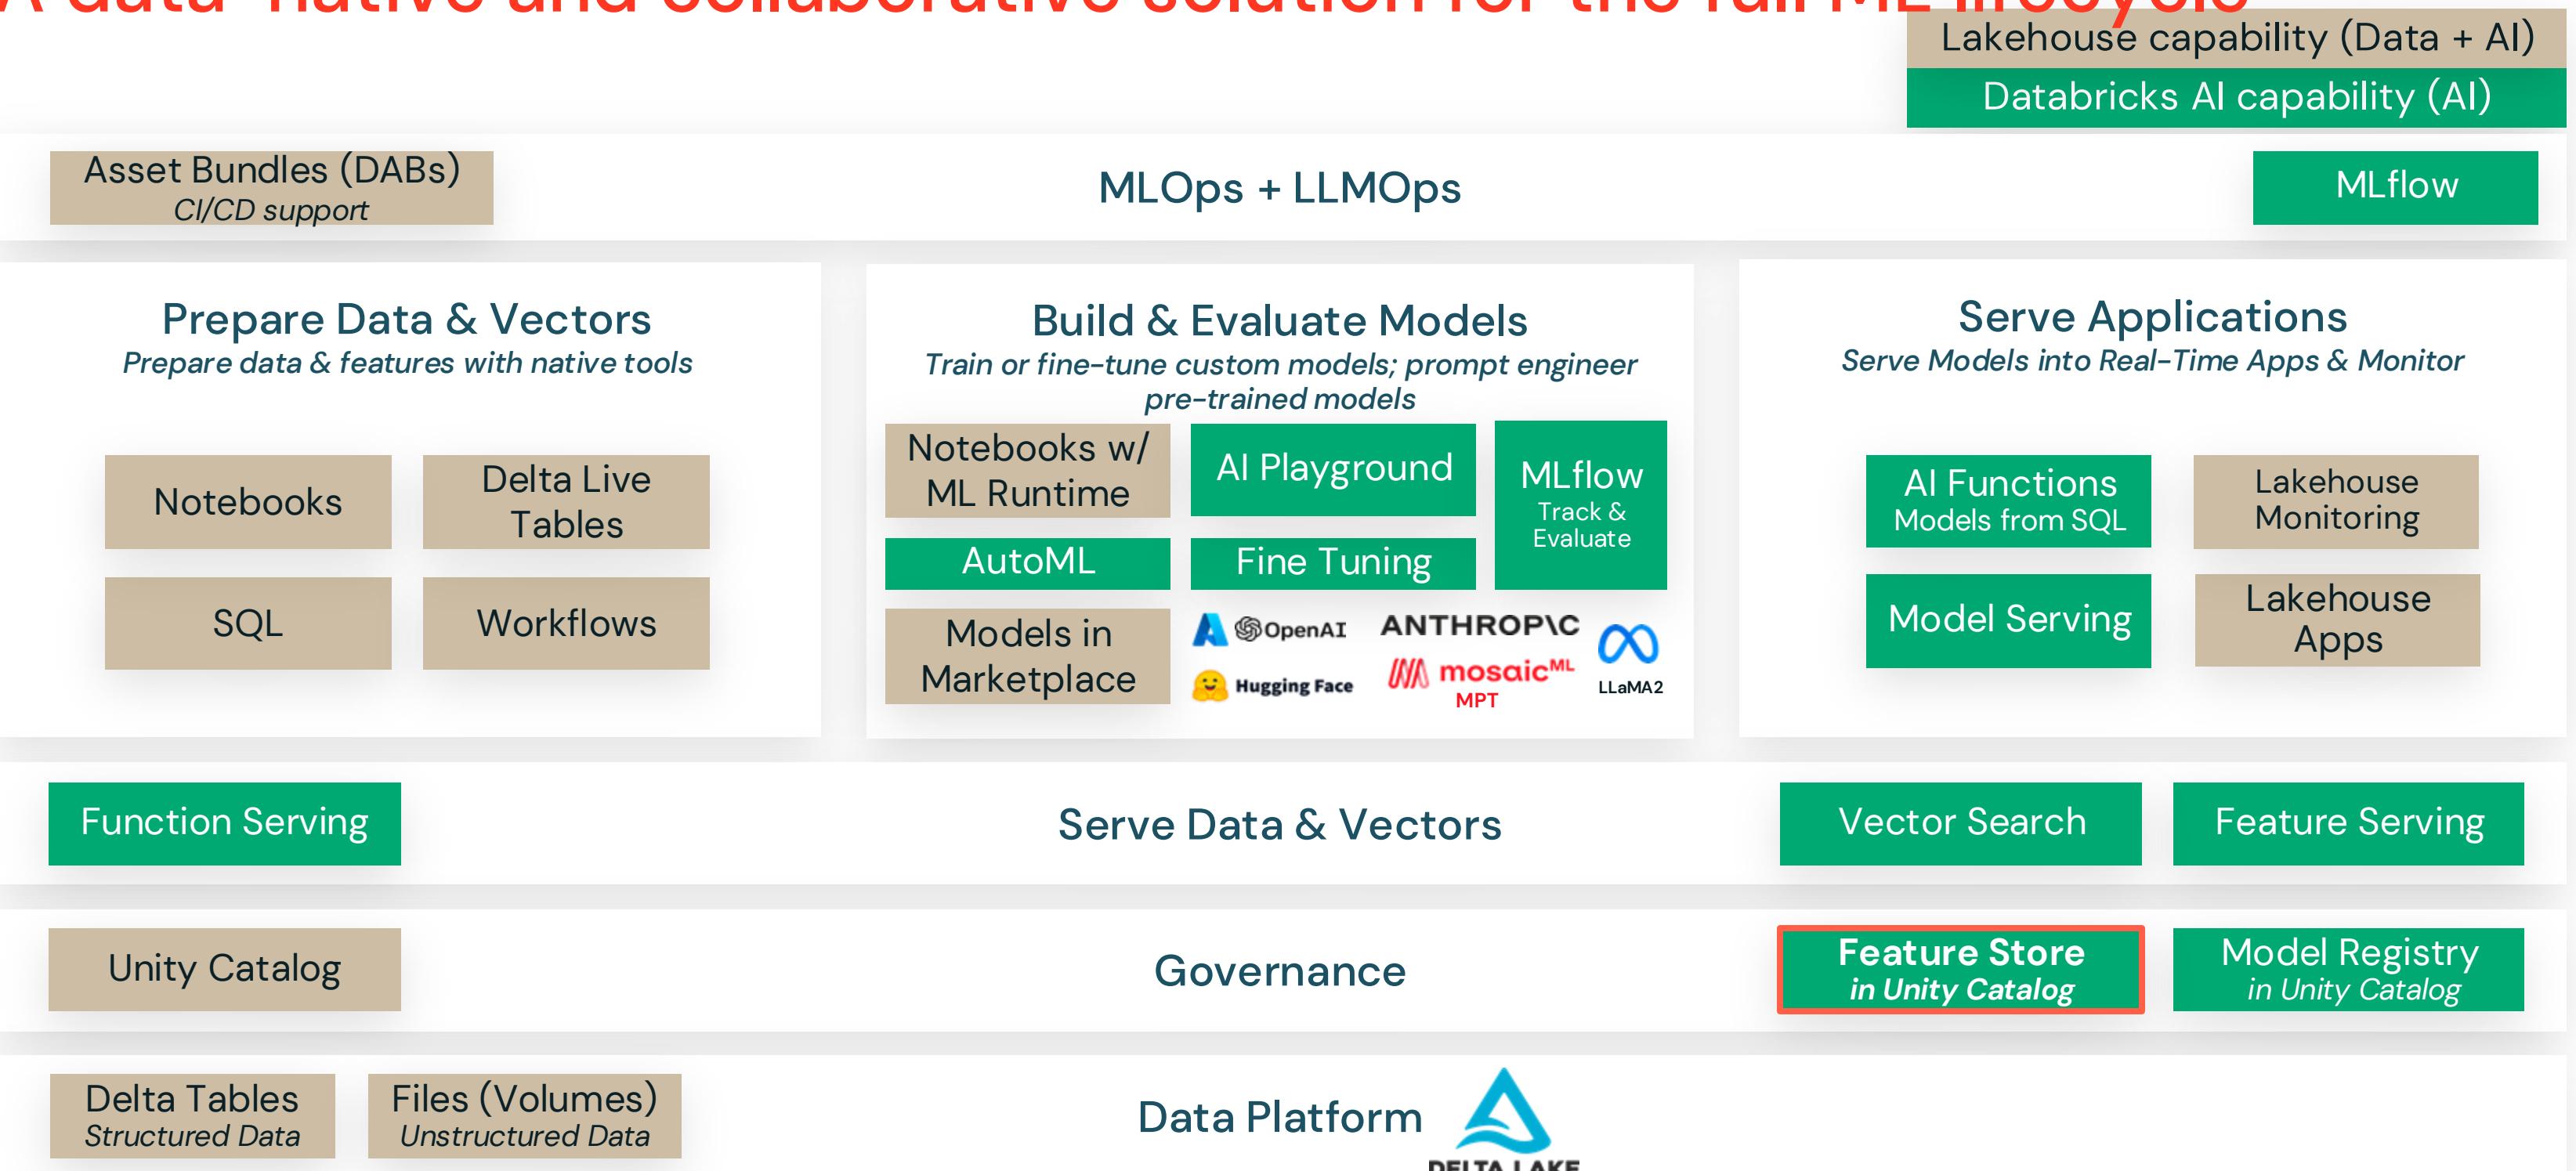


During exploration and model experimentation phases features are implemented in frameworks that do not scale to production.



# Databricks for Machine Learning

A data-native and collaborative solution for the full ML lifecycle

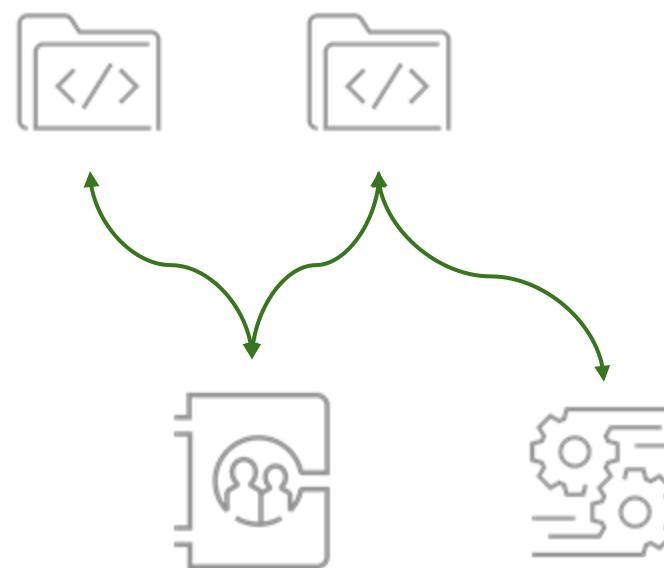


# Databricks Feature Store

## Featurization

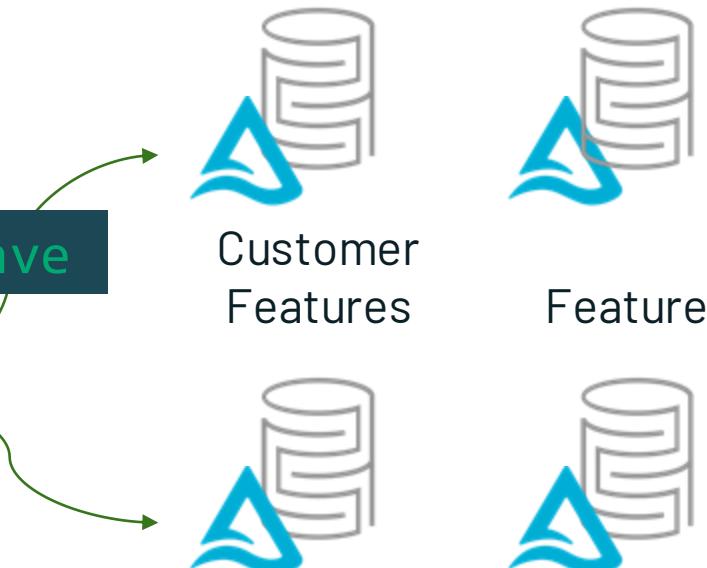
Define reusable, shareable featurization logic

Feature 1      Feature 2

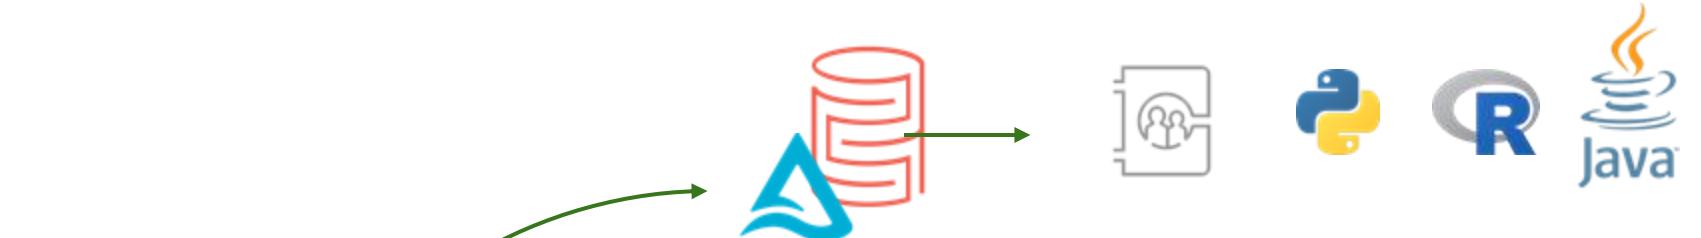


## Feature Tables

Delta Lake based: SQL, ACLs, versions, and performance optimizations



## Training Data Set Creation



## Batch Scoring



## Online Serving



Databricks  
Model Serving



# Complete Integration-FS with Unity Catalog

Any table can be a feature table

- Feature Tables become regular UC Tables with additional metadata.
- Shared properties are unified.
  - Feature table description == table comment.
  - Feature table schema == table schema
- Three-level namespace convention

The screenshot shows the Databricks Data Explorer interface. On the left, the sidebar navigation includes 'Data Explorer', 'Search for notebooks, folders and more', 'Production workspace', and 'user@acme.com'. The main area displays a tree view of databases and tables. Under 'billing\_prod\_catalog > users > Tables >', the 'user\_features' table is selected. The table details pane shows the following information:  
Owner: user@acme.com | Last updated: 8 days ago | Created: 7 weeks ago | Size: 544MB | Comment: Add comment  
ML Feature | Add tags  
Columns | Sample data | Details | Permissions | History | Lineage | Insights  
A table with columns: Name, Type, Tags, Comment. The columns listed are:

Name	Type	Tags	Comment
Column1	string	customer	🔗
Column2	date		🔗
mean_price_7d	float	department:rentals	🔗
Column4	int		🔗
Column5	int		🔗
Column6	string		🔗
Column7	string		🔗
Column8	string		🔗



# Unified Permission Model

Secure features with built-in governance

- Feature data and metadata are governed by the Unity Catalog permission model.
- Future improvements to data governance in Unity Catalog will apply to Feature data.

Catalogs > quickstart\_catalog > quickstart\_schema >

quickstart\_catalog.quickstart\_schema.quickstart\_table Delta

Comment:

Owner: haejoon.lee@databricks.com

Schema Sample Data Details Permissions History Lineage Preview

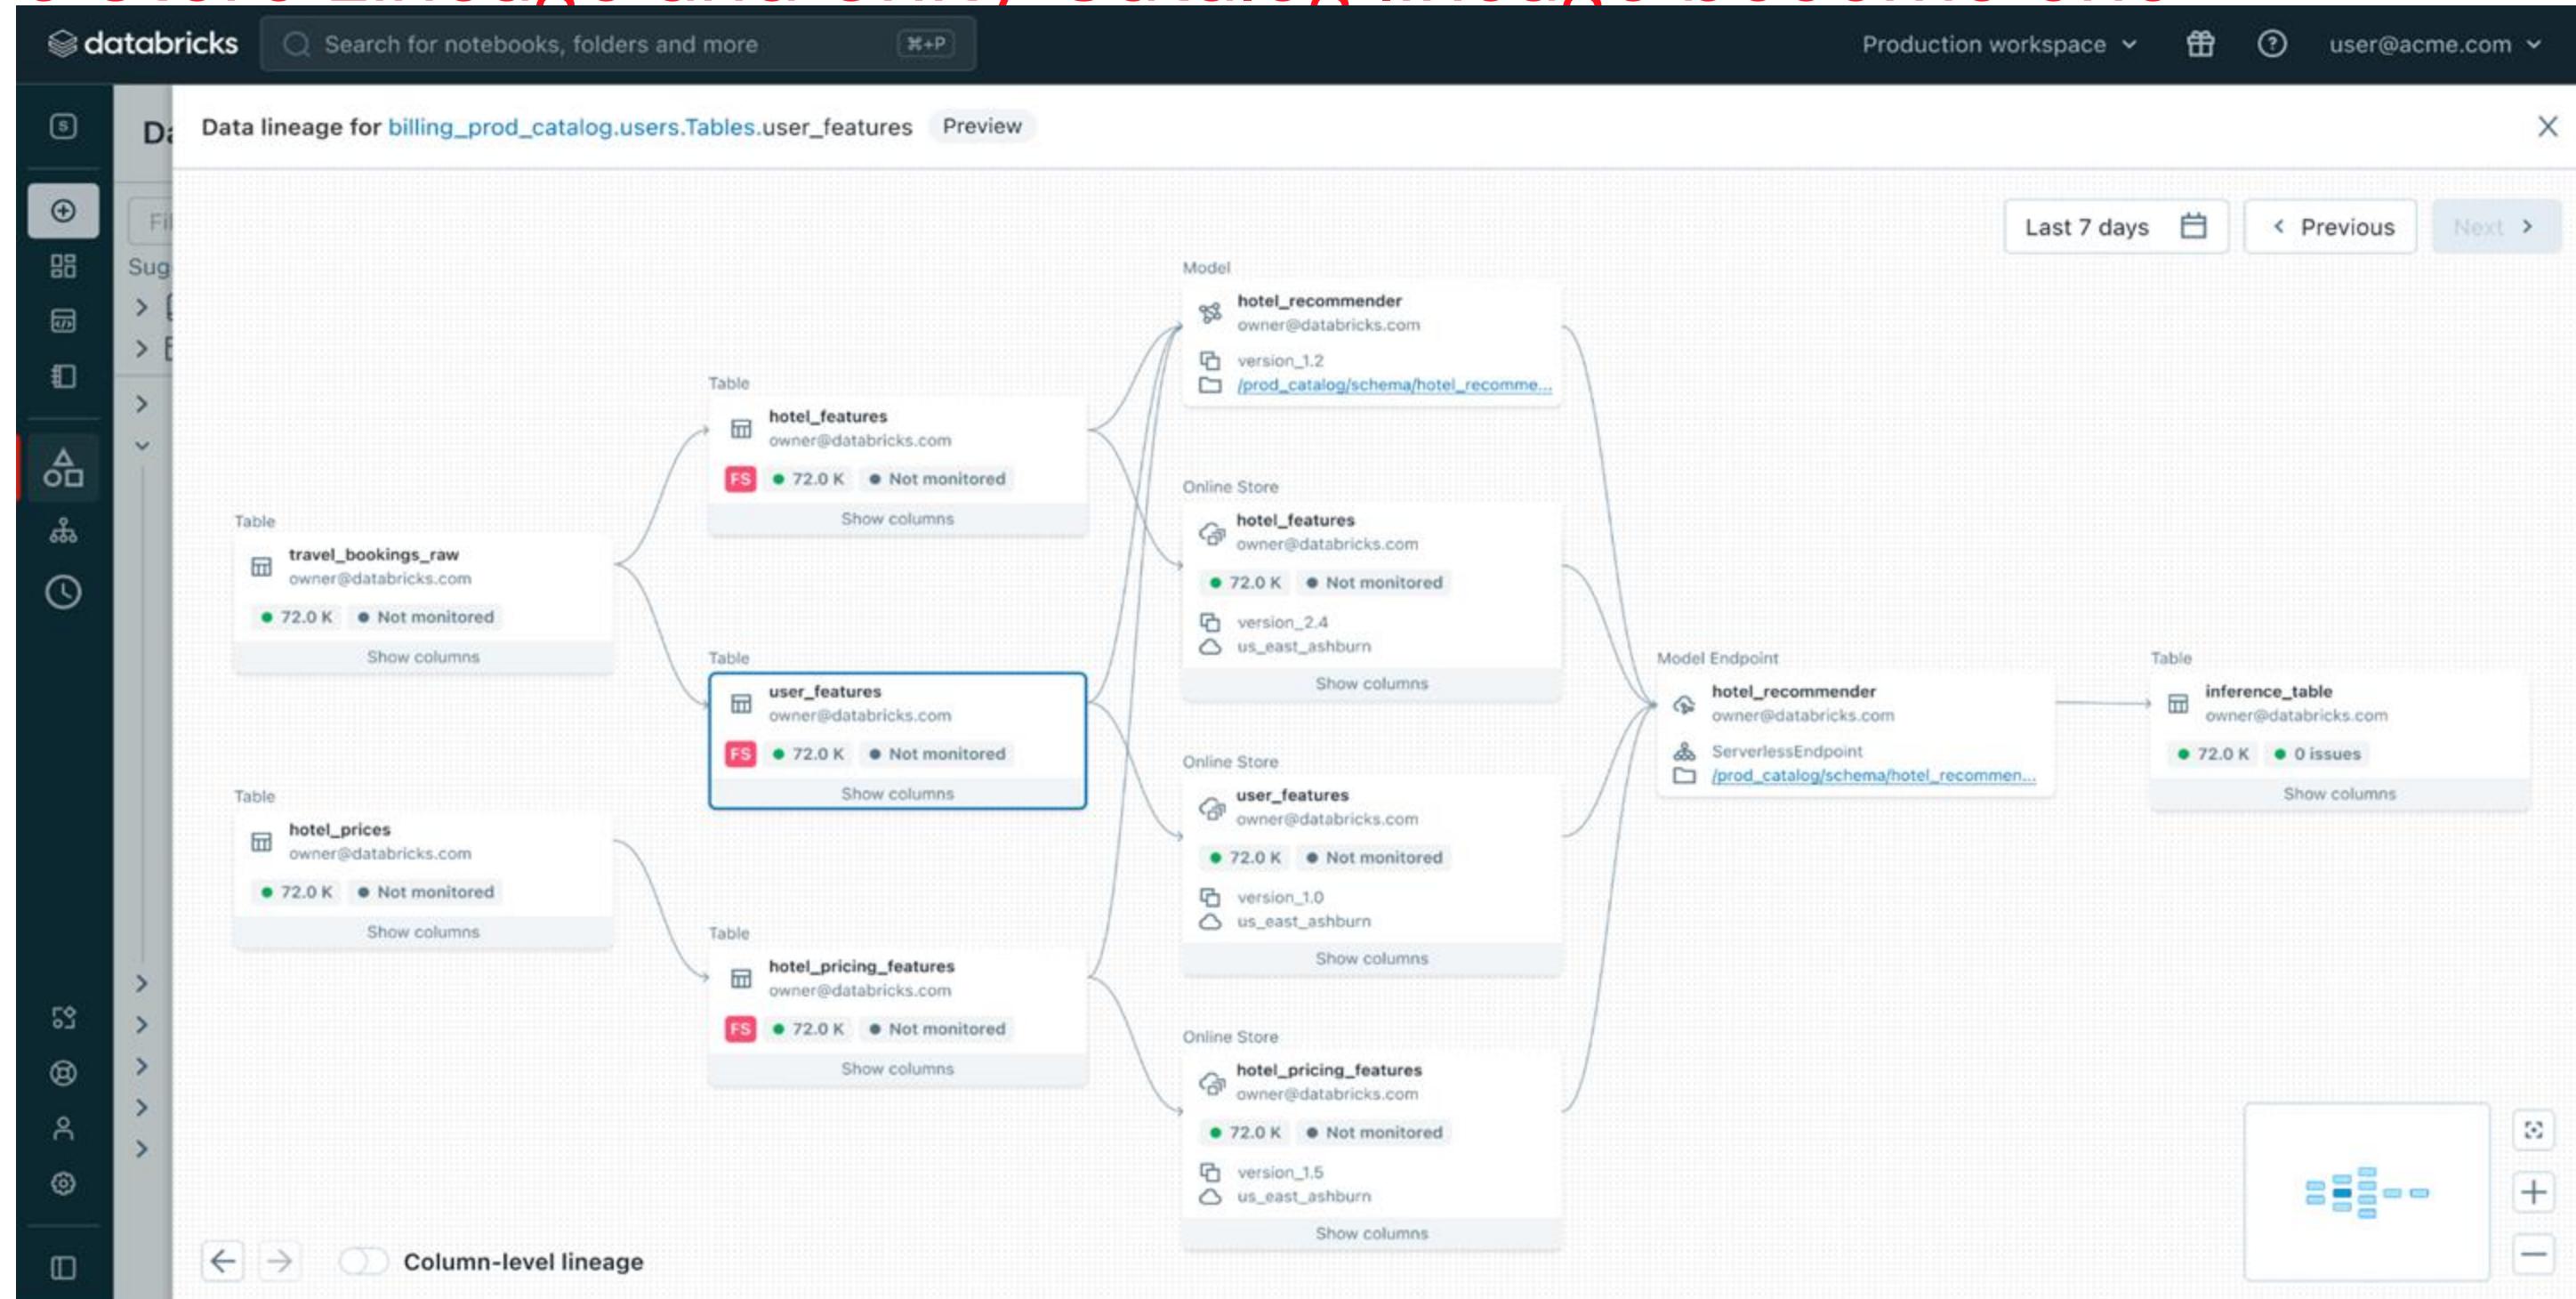
Grant  Revoke

Principal	Privilege	Object
account users	MODIFY	quickstart_catalog.quickstart_schema.quickstart_table
account users	SELECT	quickstart_catalog.quickstart_schema.quickstart_table



# Unified Data Lineage

Feature Store Lineage and Unity Catalog lineage become one



# Integration with MLflow and Model Serving

Makes model deployment easier

## Integration with MLflow

When features from FS are used for training, **the model is packaged with feature metadata.**

In inference time, the model can look up features at runtime.

## Integration with Model Serving

When the model is used for batch scoring or online inference, **it automatically retrieves features from Feature Store.**



# Workspace FS vs. Unity Catalog FS

## Workspace FS

- Specific to a single workspace
- Challenges in sharing feature tables across workspaces
- Uses `FeatureStoreClient` to create feature tables.

## Unity Catalog FS (**Recommended**)

- Any Delta table with a primary key can be a feature table (DBR 13.2+)
- All UC capabilities are available; discovery, governance, lineage and cross-workspace access.
- Uses `FeatureEngineeringClient` to create feature tables.

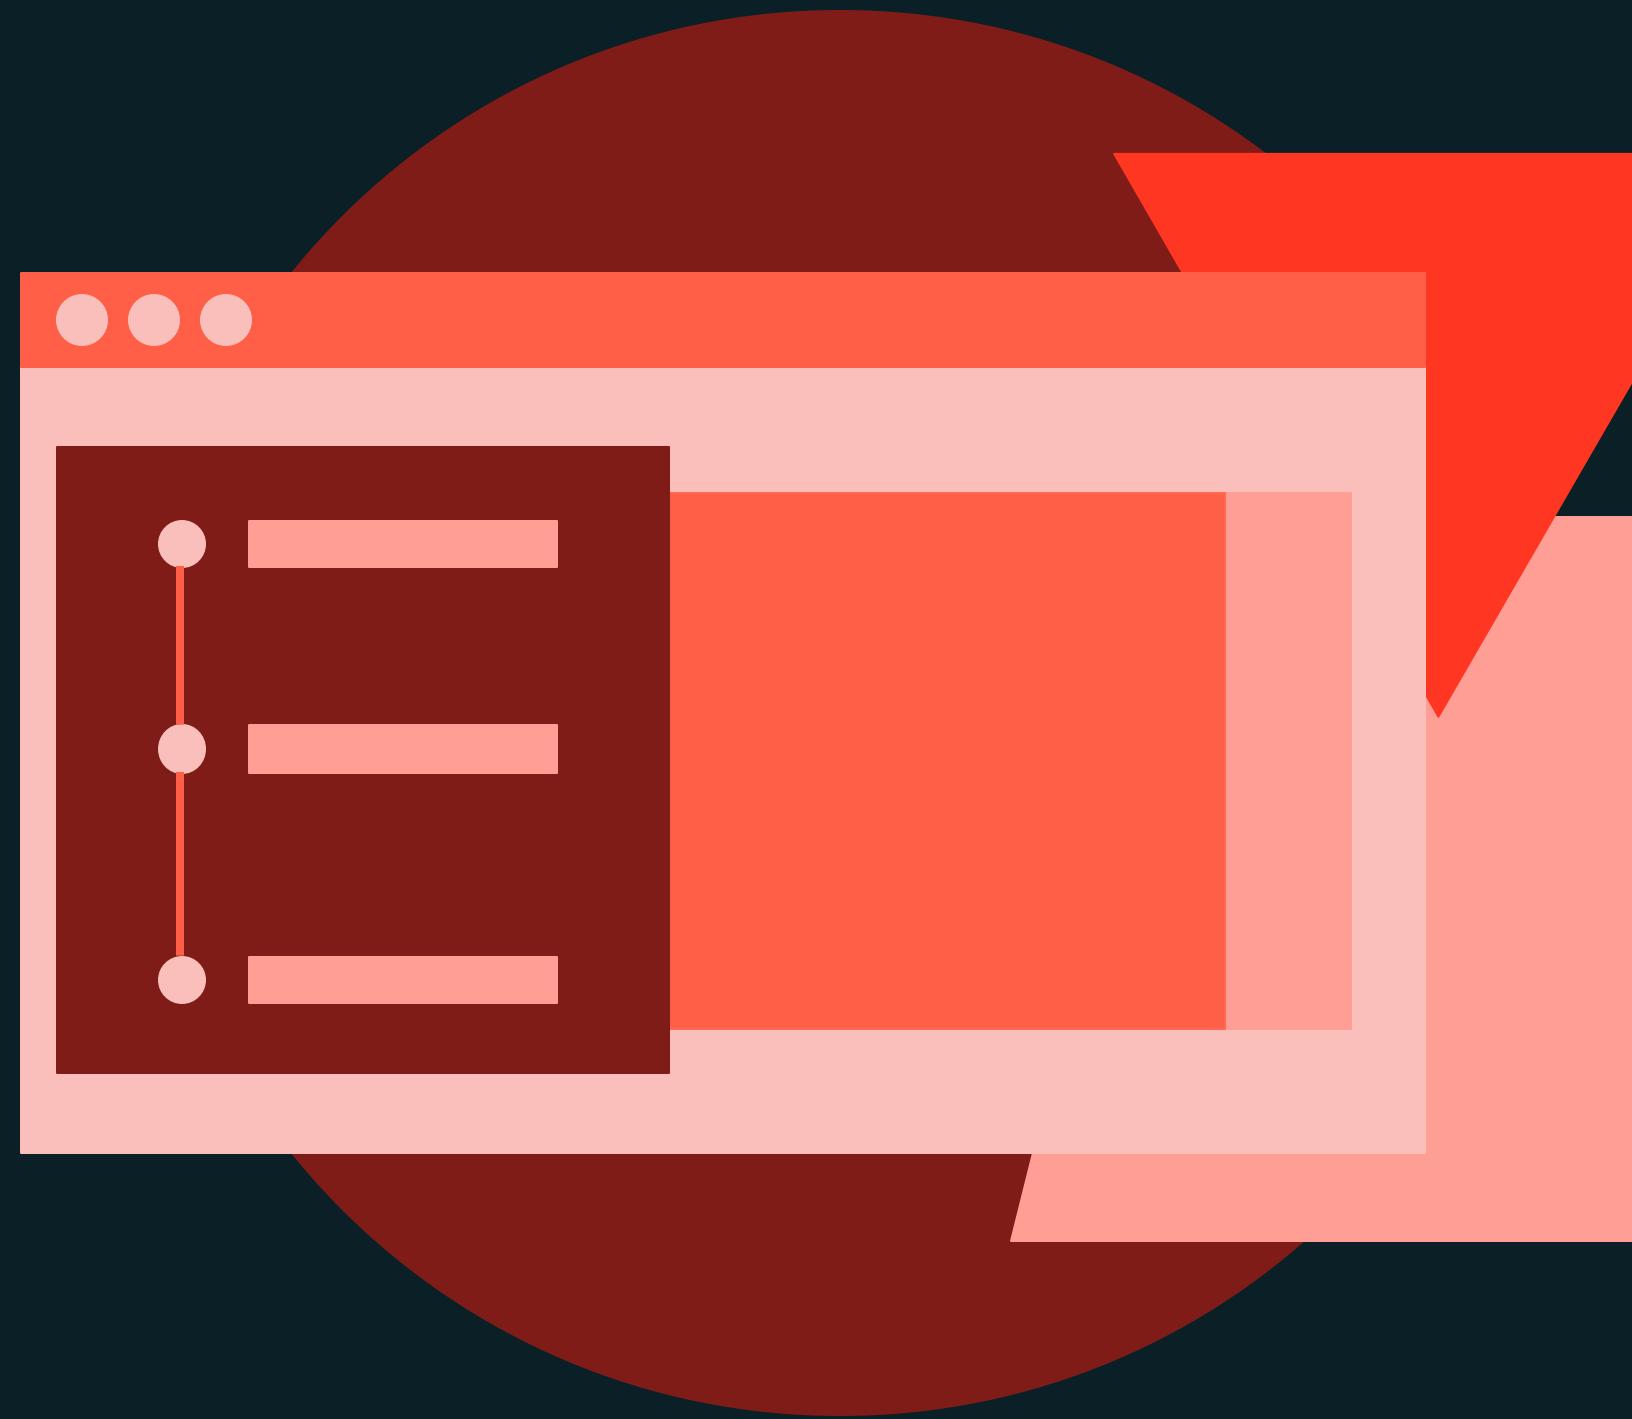




Feature Store

**DEMONSTRATION**

# Using Feature Store for Feature Engineering



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Demo

## Outline

### What we'll cover:

- Feature engineering
- Save features to feature table
- Create a feature table from existing UC table
- Update features in a feature table
- Optional: Migrate workspace feature table to UC





Feature Store

**LAB EXERCISE**

# Feature Engineering with Feature Store



# Lab

## Outline

### What you'll do:

- Data preparation and feature engineering
- Create a feature table
- Explore feature table with the UI
- Create a feature table from an existing UC table





# Course Summary and Next Steps

---

Data Preparation for Machine Learning





# databricks



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).